



Nuance Speech Recognition System
Version 8.0

Grammar Developer's Guide

Nuance Speech Recognition System, Version 8.0
Grammar Developer's Guide

Copyright © 1996-2001 Nuance Communications, Inc. All rights reserved.
1005 Hamilton Avenue, Menlo Park, California 94025 U.S.A.
Printed in the United States of America.

Information in this document is subject to change without notice and does not represent a commitment on the part of Nuance Communications, Inc. The software described in this document is furnished under a license agreement or nondisclosure agreement. The software may be used or copied only in accordance with the terms of the agreement. You may not copy, use, modify, or distribute the software except as specifically allowed in the license or nondisclosure agreement. No part of this document may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying and recording, for any purpose, without the express written permission of Nuance Communications, Inc.

Nuance and Nuance Communications are registered trademarks of Nuance Communications, Inc. Say Anything, SpeechObjects, and Verifier are trademarks of Nuance Communications, Inc. Any other trademarks belong to their respective owners.

Contents

About this guide	ix
Audience	ix
Organization	ix
Related documentation	x
Typographical conventions	xii
Where to get help	xiii
Chapter 1. The grammar development process	1
Principles of effective grammar design	1
Programming versus grammar writing	1
Predicting how callers will speak	2
Steps in the development process	3
Define the dialog	3
Identify the information items and define the slots	4
Design the prompts	5
Anticipate the caller responses	6
Identify the core of the grammar	7
Identify the filler portion of the grammar	8
Determine the appropriate type of grammar to use	8
Define your grammar in the Grammar Specification Language	9
Add natural language commands	9
Building a recognition package	10
Chapter 2. Defining GSL grammars	13
Creating a grammar file	13
Grammar names	14
Grammar descriptions	14
Grammar hierarchies	16

Top-level grammars and subgrammars	16
Including grammar files	18
Recursive grammars	19
Grammar probabilities	20
Probabilities in GSL	20
Identifying a text string as Nuance GSL	21
Language-specific grammar conventions	22
Chapter 3. Dynamic grammars: Just-in-time grammars and external rule references	23
Creating a dynamic grammar in a database	23
External rule references and redefinitions	25
Referring to external grammar rules	25
Fetching grammars	28
Referring to a specific external rule	29
Redefining an external rule	31
Specifying a grammar backoff	32
Failed references	33
Relative paths in external rule references and redefinitions	34
Specifying builtin: grammars	35
Configuring your system for external rule references and just-in-time grammars	35
Setting up a web server for external rule references	36
Configuring the compilation server	36
Parameters for external rule references	36
Just-in-time grammars	38
Writing a just-in-time grammar	38
Specifying just-in-time grammar options: Using the keyword RECOGNIZE	40
Redefining an external rule to a local rule	43
Enabling compilation of just-in-time grammars	44
Usage scenarios	44
Using the dynamic grammar gate technique	46
GSL syntax for enrollment	48

Chapter 4. Natural language understanding	51
Defining natural language interpretations	52
Natural language commands	52
The slot definitions file	58
Compiling a grammar with natural language support	58
Ambiguous grammars	58
Advanced features	59
Functions	59
Complex values	62
Semantic uniqueing	67
Robust natural language engine	67
Chapter 5. Say Anything: Statistical language models and robust interpretation	69
Overview	69
The challenge	69
A simple solution	70
A better solution: The SLM and the robust interpretation approach	71
Creating an SLM grammar	72
Create the training set	73
Create a vocabulary file (optional)	77
Determine the order of the model	79
Train the SLM grammar	79
Use the grammar in your application	81
Set appropriate parameters	83
Measure the perplexity of a model	84
Summary of the procedure for creating an SLM grammar	85
SLM and dynamic grammars	86
SLMs and just-in-time grammars	86
Robust natural language interpretation	86
Full and robust interpretation	87
Using robust interpretation	88
Ambiguity and robust interpretation	90
Statistical NL technology	92

Chapter 6. Compiling grammars	93
Compiling a grammar	93
Static compilation	93
Dynamic compilation in a database	96
Dynamic compilation to a file	96
Compilation of just-in-time grammars	98
Implicit compilation	99
Assembling grammars	100
Choosing a master package	101
Master package naming convention	101
Language-specific packages	102
Support for multiple languages	105
Adding missing pronunciations	105
Creating a dictionary file	106
Merging and overriding dictionaries	107
Automatic pronunciation generator	107
Phrase pronunciations	108
Optimizing recognition packages	110
Creating unflattened grammars	110
Speeding up compilation	110
Filler grammars	110
Tuning recognition performance	111
Chapter 7. Testing grammars	113
Grammar testing overview	113
Regression tests	114
Command-line tools for grammar testing	114
<i>parse-tool</i>	114
<i>generate</i>	115
<i>nl-tool</i>	116
<i>Xapp</i>	117
<i>batchrec</i>	118

Chapter 8. Testing recognition performance	119
Choosing <i>batchrec</i> test sets	120
Using <i>batchrec</i>	120
Creating the testset file	121
Additional options	124
Setting Nuance parameters	127
Output from <i>batchrec</i>	127
Using <i>batchrec</i> for SLM grammars	129
Example: Using <i>batchrec</i> to perform enrollment	130
Chapter 9. Creating application-specific dictionaries	133
The Nuance phoneme set	133
Multiple pronunciations	139
Sample dictionary file	140
Converting non-CPA dictionaries	142
Phoneme sets for other languages	142
Appendix A. GSL reference	143
GSL syntax	143
External grammar reference syntax	145
Summary of package files	146
Appendix B. Advanced SLM features	147
The <i>train-slm</i> tool	147
Usage	148
Editing discounting parameters	153
The <i>process-slm</i> tool	153
Usage	154
Appendix C. W3C grammar support	157
Working with W3C grammars in the Nuance System	157
Supported encodings	158
Tokens	158
Handling of the GARBAGE rule	158
DTMF mode	159

Nuance extensions to W3C grammars	160
The <rule> element	160
The <ruleref> element	161
Unsupported features	162
ABNF	162
Elements and attributes	162
W3C semantic interpretation tags	163
Language specification	163
Using W3C grammars	164
Sample grammars	164
Index	175

About this guide

This guide provides guidelines for developing and testing a grammar and the tools that Nuance offers for building efficient recognition packages.

In addition, this guide discusses the command-line utilities available for grammar development. You can use these if you are working in a Unix environment or if you prefer a command-line interface.

Audience

This guide is intended for developers wanting information about the major tasks involved in building a recognition package, including:

- Understanding the grammar development process
- Designing and writing a grammar
- Building a recognition package
- Supporting natural language interpretation
- Testing a grammar
- Defining application-specific word pronunciations

Organization

This guide is organized as follows:

Chapter 1 gives general guidelines for developing complex grammars and outlines the grammar development process.

Chapter 2 explains the basic features of the Grammar Specification Language used to describe a grammar, and describes grammar hierarchies and grammar probabilities.

Chapter 4 defines natural language interpretations and explains how to add them to a grammar. This chapter also introduces advanced interpretation features like functions and data structures.

Chapter 5 introduces the Say Anything™ feature and describes how to write statistical language models (SLM) grammars. It also explains how to use the robust natural language interpretation (robust NL) technology.

Chapter 6 describes how to compile a grammar into a recognition package and append additional word pronunciations. It also discusses some optimization and performance issues.

Chapter 7 presents a comprehensive way of testing all aspects of a grammar, including coverage, interpretation, ambiguity, and regression tests.

Chapter 8 describes how to use the Nuance *batchrec* utility to test the performance of recognition packages.

Chapter 9 explains how to make a dictionary file for an application, and how to specify word pronunciations using the Computer Phonetic Alphabet phoneme set.

Appendix A contains a formal definition of the Grammar Specification Language syntax and a summary of all the files used by the grammar compiler.

Appendix B provides advanced information on the Nuance SLM utilities.

Appendix C describes how Nuance supports W3C grammars. It also lists the Nuance extensions to the W3C format and the natural language features added to it.

Related documentation

The Nuance documentation contains a set of developer guides as well as comprehensive online API reference documentation.

In addition to this guide, the documentation set includes:

- *Introduction to the Nuance System*, which provides a comprehensive overview of the Nuance System architecture and features, the available tools and programming interfaces, and the speech application development process.
- *Nuance System Installation Guide*, which describes how to install and configure the Nuance Speech Recognition System.
- *Nuance Application Developer's Guide*, which describes how to develop, configure, and tune a Nuance speech application. This manual describes general application development and deployment issues such as setting Nuance parameters and launching recognition clients, servers, and other required processes.

- *Nuance System Administrator's Guide*, which describes topics related to monitoring the performance of a running speech application, including controlling Nuance process log output, using the Nuance Watcher to manage processes across a network, and using the runtime task adaptation feature.
- *Nuance Verifier Developer's Guide*, which describes how to use the Nuance Verifier™ to add security features to speech recognition and IVR applications.
- *Nuance Platform Integrator's Guide*, which describes the process and requirements for integrating the Nuance System with an existing IVR (interactive voice response) platform or toolkit. It provides detailed information on the primary Nuance integration APIs, the RCEngine and Java SpeechChannel™, and also discusses other specialized APIs for adding custom components such as audio providers to your integration.
- *Nuance System Glossary*, which defines terms and acronyms used in the Nuance documentation set.
- *Nuance API Reference*, which includes HTML-based documentation for Nuance APIs, parameters, command-line utilities, and GSL (Grammar Specification Language) syntax. To access this documentation, open the file `%NUANCE%\doc\api\index.html`.

You may also refer to the Foundation SpeechObjects™ documentation, including the *SpeechObjects Developer's Guide*, which describes the Nuance SpeechObjects framework and how to use it to build a speech recognition application. SpeechObjects documentation is shipped with the Foundation SpeechObjects product.

To view the entire set of documentation online, open the file `%NUANCE%\doc\index.html` in any HTML browser. Install the `%NUANCE%\doc` directory on an internal web server or file system to minimize space requirements. This directory includes all Nuance developer guides in both HTML and PDF format, and HTML reference documentation.

Typographical conventions

Nuance manuals use the following text conventions:

italic text Indicates variables, file and path names, program names, program arguments, web and email addresses, as well as terms introduced for the first time. For example:

Edit the *ag.cfg* configuration file.

Courier New Indicates method/member functions, parameter names and values, program constants, and onscreen program output. For example:

Nuance recommends that you set the parameter `audio.OutputVolume` to 255.

> Courier New Indicates commands or characters you type in and the responses that appear on the screen. The > character indicates the MS-DOS command prompt or Unix shell. Everything after this character is intended to be typed in. For example:

```
> resource-manager
```

In this example, the text that you actually type at the command line is “resource-manager”

Courier New Indicates a value that you replace. For example:

The usage for the *nlm* utility is:

```
> nlm license_key
```

In this example, *license_key* is a value that you replace, so the text you actually type could be, for example:

```
> nlm ncr8-16-100-a-b22-333c4d55eeff
```

Note: The Nuance System runs on both Windows and Unix platforms. Windows syntax is typically used throughout the documentation. If you are running on a Unix platform, substitute Unix syntax. For example, use *\$NUANCE* wherever *%NUANCE%* appears, and use “/” in place of “\” in path names. Differences in usage or functionality on Windows and Unix platforms are noted where relevant.

Where to get help

If you have questions or problems, Nuance provides technical support through the Nuance Developer Network™ (NDN), a web-based resource center that includes online forums, technical support guides on specific topics, access to software updates, as well as a support request form. If you are a member, go to *extranet.nuance.com* to log on, or see the Nuance website *www.nuance.com* for information on how to become a member.

To submit comments on the documentation, please send email directly to *techdoc@nuance.com*. Note that no technical support is provided through this email address. Technical support is provided through the NDN.

The grammar development process

1

An intuitive user interface that guides the user through a constrained but purposeful interaction is a crucial component of speech recognition application design. The user speaks naturally, yet, because of well-designed prompts and grammars, the user's utterances fall largely within an expected set of phrases, allowing the recognizer to achieve a high accuracy rate.

This chapter describes an approach to grammar development designed to help you create effective grammars. These guidelines are the result of experience gained by developers writing grammars for real-world applications. They provide a tested approach to the complex task of designing and writing an effective grammar—a critical component of a good speech recognition application.

Principles of effective grammar design

The following sections describe general principles that Nuance advocates for the design of quality grammars. Keep these guidelines in mind when writing grammars for your particular application. Chapter 2 describes the syntactic structure and elements of grammars in detail.

Programming versus grammar writing

The complexity of a grammar greatly affects the speed and accuracy of the recognizer. Complex grammars must be constructed with as much care as complex software programs.

Grammar writing is an unfamiliar task for most software developers, and creating a high-quality, error-free grammar requires somewhat different skills than programming in a language like Java or C++. Grammars are inherently

non-procedural and thus software programming and grammar writing cannot be approached in the same way.

Conceptually, a grammar is a set of phrases—possibly infinite—that a caller is expected to say during a dialog in response to a particular prompt. The grammar writer’s task is to *predict* that set of phrases and *encode* them in the grammar.

Predicting how callers will speak

Of the two tasks, predicting and encoding, predicting the set of responses is by far the more difficult. Even if you are just expecting a simple yes/no response, you will likely get a wide range of responses from real callers, such as “yeah,” “yup,” “no way,” “correct,” and others that you might not guess in advance.

Grammar writing is an iterative process: you make your best guess, collect some real data, refine the grammar, get some more data, refine further, and so on. As you refine the grammar by adding and removing phrases, it more closely approximates the way callers speak to the application.

In practice, you are not able to include *all* of the responses that can occur in your application because you cannot control how people speak. As a rule of thumb, a 5% out-of-grammar rate is considered acceptable. Even 10-20% out-of-grammar rates are not uncommon for certain types of grammars.

Note: A phrase is *out-of-grammar* if it cannot be parsed by a grammar.

A good piece of advice is to avoid putting too much effort into thinking of *every* alternative way of saying something right from the start—this will largely be a wasted effort. Instead, focus on guessing the most common ways that people will respond and build those into your first grammar release. After you have collected some data, expand your grammar based on that field data.

How do you guess the most common responses? Fortunately, it turns out that there are two types of responses that are by far the most common:

- The information item by itself
- The literal response to the question wording

So, if you ask “What is your departure city?,” most responses will be just a city name like “New York” with no other verbiage. A smaller group of responses will contain phrases like “My departure city is Miami” or “departure from Miami” in direct response to the prompt wording. If you change the question to “What city would you like?,” you’ll still get city-only responses, but you’ll also get “I’d like Miami” responses. In this second case, you probably won’t get many (or any) responses of the form “My departure city is Miami”.

Therefore it is important to:

- Word prompts carefully
- Coordinate grammars and prompts, making sure that your grammars correspond closely to the prompt wording

Furthermore, whenever you change the wording of a prompt, be sure to modify the corresponding grammar as well.

Steps in the development process

The following sections describe the generic process of developing a grammar. Details on specific concepts and terminology are included later in this guide and recommendations on ways to format the encoding of a grammar are found in Appendix A.

The tasks you should follow while developing a grammar are:

- 1 Define the dialog
- 2 Identify the information items and define the slots
- 3 Design the prompts
- 4 Anticipate the caller responses
- 5 Identify the *core* and *filler* portions of your grammars
- 6 Write the code for your grammars
- 7 Add natural language commands
- 8 Build a recognition package

Define the dialog

It is important to define the dialog before starting to write a grammar, because the dialog determines what grammars you have to write. For an important commercial application, the dialog is normally defined in a formal dialog specification document, though you can use whatever type of documentation makes sense for your project. For instance, a one-time demo would not require a large amount of detail specification. In any case, it is important to have a good understanding of the dialog *before* you start to write the grammars.

At a minimum, you should answer the following questions:

- 1 What pieces of information are required to complete the task?

- 2 In what order will the information be requested?
- 3 Will the dialog request one piece of information at a time in a particular order—a *directed dialog*—or will it allow several pieces of information at once, in any order, and prompt for missing items as necessary—a *mixed-initiative dialog*?

The answers to these questions determine the shape and content of the grammars you need to develop.

Identify the information items and define the slots

Once your dialog has been defined, it is relatively simple to determine what items your dialog should capture. Normally, you would use one *slot* for each piece of information. A slot is similar to an identifier in a data structure in that it holds a value of certain type. (See the introduction to Chapter 4 for more on slots.)

For example, if you're creating an air travel application, you might need to collect two cities (origin and destination), a date, and a time, and then confirm the validity of information assembled (a yes/no question). You might then ask the caller if they want to hear the return flight information, start over, or hang up (a small set of commands). That's six pieces of information in all. At this point, you may also want to determine the format and *type* in which the information will be returned.

You can summarize all that information in table like the following:

Item	Slot name	Value format	Value type
city #1	origin	3-letter code	string
city #2	destination	3-letter code	string
date	date	[<month> <day>]	NL structure
time	time	0-2359	integer
yes/no	confirm	"yes" or "no"	string
restart/hangup	command	"restart" or "hangup"	string

This information helps you set up your grammars to return the right values in the right format in the right slots.

Design the prompts

After defining the dialog and information items, you are ready to write the wording for your dialog’s prompts.

Prompt design is best done *before* writing the grammars because prompt wording can greatly affect the wording of the caller responses, as pointed out earlier. The grammar needs to capture those responses, so if the prompts are changing frequently while the grammars are developed, you will probably have to do a lot of rework.

Note: Core items, such as city and name lists, can typically be developed earlier in the process—that is, before completing the dialog design.

To request flight information, for example, consider the following possible prompts and slots they would fill:

Table 1: Sample prompts and natural language slots

Prompt	Slot
What city would you like to leave from?	origin
What city would you like to fly to?	dest
What date would you like to leave?	date
What time would you like to depart?	time
You're going from <origin> to <dest> on <date> at <time>. Is this correct?	confirm
Would you like to start over or hang up?	command

If you have additional error or help prompts that can immediately precede recognition, you should write these as well, and take them into consideration when you write the grammars.

Note: The preceding prompts are appropriate for a directed dialog. A mixed-initiative dialog might instead start by asking “Where would you like to travel?” or “How can I help you?” and then pose more specific questions to obtain the missing pieces of information. It is much more difficult to predict the range of responses to an open-ended question. This makes the grammar more difficult to write and tune, although it is doable. It is up to you to decide whether the dialog will elicit only simple responses or more complex ones.

Anticipate the caller responses

After designing your prompts, you can guess more accurately how callers will respond. Remember that the two most typical responses in a directed dialog will contain just one of the following:

- The information item by itself
- The literal response to the question wording

You should also consider that people tend to hesitate at the start and sometimes say “please” at the end.

Taking these points into account, here are some guesses as to how callers might respond to each of the prompts in Table 1 on page 5.

<i>What city would you like to leave from?</i>	
San Francisco	[the city name by itself]
I'd like to leave from San Francisco	[a literal response]
Uh, San Francisco	[initial hesitation]
San Francisco, please	[final “please”]
(I'm) leaving from San Francisco	[some additional possibilities]
(I'm) departing from San Francisco	
<i>What city would you like to fly to?</i>	
New York	[the city name by itself]
I'm flying to New York	[a literal response]
I'd like to fly to New York	[another literal response]
Uh, New York	[initial hesitation]
New York, please	[final “please”]
(I'm) going to New York	[some additional possibilities]
My destination is New York	
<i>What date would you like to leave?</i>	
May second	[the date by itself]
I'd like to leave on May second	[a literal response]
I'm leaving on May second	[a second literal response]
Leaving May second	[a third literal response]
Um, May second, please	[hesitation + final “please”]

<i>What time would you like to depart?</i>	
2 pm	[the time by itself]
I'd like to depart at 2 pm	[a literal response]
I'm departing at 2 pm	[a second literal response]
Departing 2 pm	[a third literal response]
2pm, please	[final "please"]

<i>You're going from <origin> to <dest> on <date> at <time>. Is this correct?</i>	
Yes	["yes" by itself]
No	["no" by itself]
Yes, that's correct	[a literal response]
Yes it is	[a second literal response]
No, that's not correct	[a third literal response]
No, it's not	[a fourth literal response]
Yeah (or yup, or you bet)	[casual alternatives]

<i>Would you like to start over or hang up?</i>	
Start over	[command by itself]
Hang up	[command by itself]
I'd like to start over	[a literal response]
Um, start over please	[hesitation + final "please"]

Identify the core of the grammar

A grammar typically consists of a *core* portion that contains the most important meaning-bearing words—like cities, dates, and times—and a *filler* portion that contains additional expressions such as “I’d like to...” or “please.”

The core portion is often highly reusable, so it makes sense to define a *subgrammar* or *grammar rule*—a smaller grammar used in building up hierarchies within larger grammars—describing just the core portion of a grammar. Information that pertains to a particular grammar can then be added in a higher-level more specific grammar.

In the flight information example, the core subgrammars should describe cities, dates, time, and confirmation. The “start over” and “hang up” commands are instead specific to this application, so no core grammar need be created for these.

Identify the *filler* portion of the grammar

The filler portion of a grammar depends largely on the prompt wording. If you have considered the caller responses, as described in “Anticipate the caller responses” on page 6, then you start by replacing the core portion of each utterance, in the list of anticipated phrases, with the name of a core grammar.

The portion of the original responses that remains after replacement is, very likely, the filler part of your grammar.

In the flight information example, you could use the grammar rules *CITY* and *DATE*, leading to the following types of transformed phrases:

- What city would you like to leave from?

CITY

I'd like to leave from *CITY*

Uh, *CITY*

CITY, please

(I'm) leaving from *CITY*

(I'm) departing from *CITY*

- What date would you like to leave?

DATE

I'd like to leave on *DATE*

I'm leaving on *DATE*

Leaving *DATE*

Um, *DATE*, please

At this point—once the core and filler portions have been clearly identified—you have nearly written the grammar. All you need to do is write the final grammar definitions.

Determine the appropriate type of grammar to use

Nuance provides two types of grammars: Grammar Specification Language (GSL) grammars and statistical language model (SLM) grammars. GSL grammars are useful when the application's prompts are sufficient to restrict the user's response. In general, GSL grammars are appropriate for most applications.

SLM grammars are appropriate for recognizing free-style speech, especially when the out-of-grammar rate is high. Consider an application that requires an

open-ended prompt such as “Please state the nature of your problem?”. Not only can the user’s response be highly variable and hard to predict, but it can also contain restarts, filled pauses (*um* and *uh*), and ungrammatical sentences. An SLM grammar would be very appropriate for such an application.

The rest of this chapter describes how to write a GSL grammar. See Chapter 5, “Say Anything: Statistical language models and robust interpretation” for more information about SLM grammars.

Note: Nuance also supports W3C grammars. See Appendix C, “W3C grammar support” for more information.

Define your grammar in the Grammar Specification Language

The Grammar Specification Language (GSL) is the language you use to formally specify a grammar for a Nuance System application.

The two grammars in the flight information example (departure city and date), are readily translated to GSL from the lists above. Assuming that you have the *CITY* and *DATE* subgrammars (for example, from the Grammar Library), the code looks like the following:

```
.DEPARTURE_CITY [CITY
                  (i'd like to leave from CITY)
                  (uh CITY)
                  (CITY please)
                  (?i'm leaving from CITY)
                  (?i'm departing from CITY)
                  ]
.DEPARTURE_DATE [DATE
                 (i'd like to leave on DATE)
                 (i'm leaving on DATE)
                 (leaving DATE)
                 (um, DATE please)
                 ]
```

CITY and *DATE* are subgrammars defined elsewhere.

Add natural language commands

The next step, adding natural language commands to the grammar, is straightforward, but you need to know how to do it using GSL. Note the following points in the code fragments below:

- *c* and *d* are *variables*
- The expressions *CITY:c* and *DATE:d* set the variables *c* and *d* with the values returned by the subgrammars *CITY* and *DATE*, respectively

- The expressions *\$c* and *\$d* are references to the values of the corresponding variables
- The expressions *{<origin \$c>}* and *{<date \$d>}* fill the slots *origin* and *date* with the values held in the variables *c* and *d*, respectively

```
.DEPARTURE_CITY [CITY:c
                (i'd like to leave from CITY:c)
                (uh CITY:c)
                (CITY:c please)
                (?i'm leaving from CITY:c)
                (?i'm departing from CITY:c)
                ] {<origin $c>}
.DEPARTURE_DATE [DATE:d
                (i'd like to leave on DATE:d)
                (i'm leaving on DATE:d)
                (leaving DATE:d)
                (um DATE:d please)
                ] {<date $d>}
```

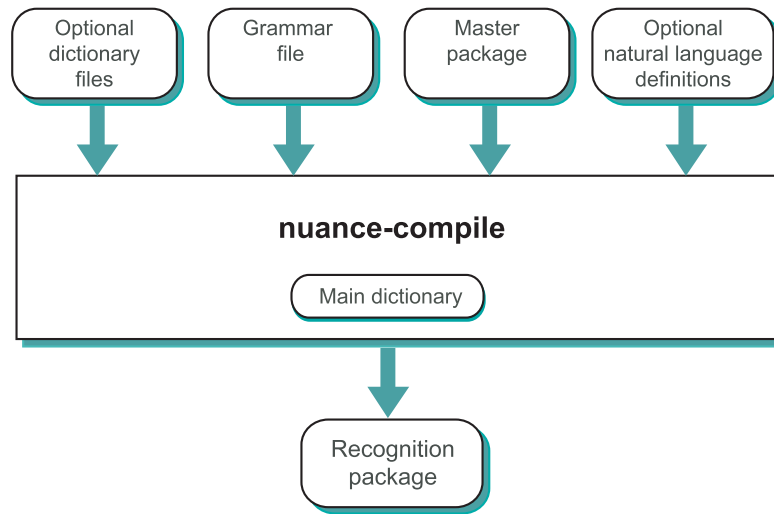
The grammars are now ready to be compiled and tested.

Building a recognition package

A *recognition package* contains the information to configure the Nuance Speech Recognition System for a specific application, including:

- Definitions of recognition grammars
- Pronunciations for the words in those grammars
- Pointers to the master packages selected for the application

It can also contain natural language processing information, to be used by the Nuance System's natural language processing engine (see Chapter 4). You generate recognition packages by creating the correct specification files and then compiling them with the command-line program *nuance-compile*. (The different compilation options are described in Chapter 6, "Compiling grammars.")



The two required inputs to *nuance-compile* are:

- A grammar file name
- A master package set name

The grammar file defines the recognition grammars used in an application. Each grammar describes a set of phrases that the Nuance System will consider during the recognition process. The structure and contents of a grammar are discussed in Chapter 2.

The master package defines, among other things, the languages that an application will use. Master packages differ in their accuracy, memory requirements, and computational requirements. Master packages are discussed in Chapter 6.

Optional inputs to *nuance-compile* such as natural language processing commands, a slot definitions file, and dictionary files containing phonetic pronunciations for words used in the grammars are also discussed in Chapter 6.

Defining GSL grammars

2

The grammars that an application uses for recognition are defined in grammar files. Each grammar describes a set of word sequences that the Nuance System weighs during the recognition process. A grammar can be as simple as “yes” versus “no,” as large as a list of all the names of people living in a city, or complex enough to support a dialog that registers a user in a course or traces a package for a delivery company.

This chapter describes how to write a GSL grammar.

Chapter 4 describes more complex grammar-writing techniques that include natural language interpretation.

Creating a grammar file

You write grammar definitions using a special language called the Grammar Specification Language, or GSL. You specify a GSL grammar in a *grammar file*. A grammar file is a text file—it can contain more than one grammar definition, and it has the file extension *.grammar*, for example, *myApp.grammar*. A grammar file is the basic component necessary to build a recognition package.

At the top level a grammar definition has the format:

GrammarName GrammarDescription

where *GrammarName* is the name of the grammar being defined and *GrammarDescription* defines the contents of the grammar. The simplest example of a grammar is one whose description is a single word:

```
.Account checking
```

The grammar file name has the format *package_name.grammar*, where *package_name* is the name of the recognition package you want to create. For

example, if you create a grammar file called *banking.grammar*, the resulting recognition package is named *banking*.

The grammar *package_name.grammar* is the *primary* grammar file since it is the file passed to the compiler. The primary grammar file, however, need not be the *only* grammar file in your project, as GSL provides a directive that lets you refer to other grammar files within a grammar file.

Note: A formal definition of all the constructs available in GSL is included in Appendix A.

Grammar names

GrammarName is the character string that other grammars or an application use to reference the named grammar. Grammar names must contain at least one uppercase character—typically the first alphabetic character—and can be up to 200 characters in length. All grammar names passed to the compiler must be distinct—that is, a grammar name can only have one grammar description associated with it.

The following characters are allowed in a grammar name:

- Uppercase and lowercase letters
- Digits
- The special characters:
 - - (dash)
 - _ (underscore)
 - ' (single quote)
 - @ ("at" sign)
 - . (period)

Other characters are not allowed.

Grammar descriptions

A *GrammarDescription* consists of a sequence of word names, grammar names, and operators that define a set of recognizable word sequences or phrases. Grammar and word names must be separated from one another by at least one white space character—space, tab, or newline.

Word names are the terminal symbols in a grammar description. Word names are lowercase character strings that correspond directly to the actual words

spoken for recognition. For example, the word name *dog* corresponds directly to the spoken word “dog.”

Word names cannot contain any uppercase letters, but can contain the other special characters that are allowed in grammar names (digits, “-”, “_”, and so on, as listed in “Grammar names” on page 14). You can include other special characters (except for white space and double quotes) if you enclose the word in double quotes. For example, “foo*bar” defines a legal word, but “foo bar” does not.

You can add comments in a grammar description by using a semicolon (;)—all text in a line after a semicolon is ignored by the compiler. Comments can be included anywhere in a grammar file (or in any of the other package files mentioned in this manual).

You construct a grammar description by using a set of five basic grammar operators: (), [], ?, +, and *, described in Table 2. White space is optional between operators and operands (grammar or word names).

The symbols A, B, C, and D in the following table denote a grammar or word name.

Table 2: Grammar operators

Operator	Expression	Meaning
() <i>concatenation</i>	(A B C ... D)	A and B and C and ... D (in that order)
[] <i>disjunction</i>	[A B C ... D]	One of A or B or C or ... D
? <i>optional</i>	?A	A is optional
+ <i>positive closure</i>	+A	One or more repetitions of A
* <i>kleene closure</i>	*A	Zero or more repetitions of A

Here are some simple GSL expressions and some of the phrases they describe:

[morning afternoon evening]

“morning”, “afternoon”, “evening”

(good [morning afternoon evening])

“good morning”, “good afternoon”, “good evening”

(?good [morning afternoon evening])

“good morning”, “good afternoon”, “good evening”, “morning”,
“afternoon”, “evening”

(thanks +very much)

“thanks very much”, “thanks very very much”, and so on

(thanks *very much)

“thanks much”, “thanks very much”, “thanks very very much”, and so on

Caution: The following strings are reserved for internal use and cannot be used to denote words or grammar names (*n* designates any integer):

AND-*n*, OR-*n*, OP-*n*, KC-*n*, PC-*n*

Grammar hierarchies

You can build a hierarchy of grammars using subgrammars, also called grammar rules. By breaking a grammar into smaller units, you can create components that are reusable by multiple grammars or applications.

The use of subgrammars:

- Simplifies grammar creation and revision
- Helps focus the grammar development to the task at hand
- Hides unnecessary details and promotes modularity

Top-level grammars and subgrammars


A grammar is either a *top-level* grammar or a subgrammar. A top-level grammar is, by definition, one whose name has a period (.) as its first character, for example, *.Account*.

Only top-level grammars can be referenced by an application at runtime. Any other grammar—one whose name does not begin with a period—is a subgrammar that can only be referenced by other grammars. A grammar named *.SENTENCE*, for example, can be referenced by an application or by other grammars, while a grammar named *SENTENCE* can only be referenced by other grammars.

Caution: The distinction between top grammars and subgrammars does not apply to grammars used dynamically, including just-in-time, VoiceXML, and SpeechObjects grammars.

Subgrammars let you define complex grammars as a hierarchy of smaller grammars. This simplifies your specifications and provides a more efficient way to specify grammars with shared internal structure. For example, a grammar defining how to say a date might have subgrammars for day, month, and year,

and those subgrammars might reference subgrammars for different ways of saying numbers, and so on:



The top-level grammar *.Date* can then be referenced by applications for recognition, without needing to know anything about the subgrammars used by the *.Date* grammar.

To include the contents of a subgrammar in a grammar description, you just refer to the subgrammar by its unique name as you do with a word name. The description of the subgrammar referred to is included in exactly the location you specify it in.

For example, the following grammars describe *naturally phrased* numbers from 0 to 99. Four subgrammars are defined and then used to create the top-level grammar *.N0-99*, that can be referenced by applications.

```

; Sample GSL code for natural numbers from 0 to 99
; subgrammar for one of the non-zero digits
NZDIGIT [one two three four five six seven eight nine]

; subgrammar for one of the digits including zero
DIGIT [one two three four five six seven eight nine zero oh]

; subgrammar for one of the teen words and ten
TEEN [ten eleven twelve thirteen fourteen fifteen sixteen
      seventeen eighteen nineteen]

; subgrammar for one of the decade words except ten
DECADE [twenty thirty forty fifty sixty seventy eighty
        ninety]

; the top-level grammar, which references the subgrammars
; defined above
.N0-99 [(?NZDIGIT DIGIT) TEEN (DECADE ?NZDIGIT)]
  
```

Phrases defined by the top-level grammar *.N0-99* include: “two,” “eight two,” “zero,” “seventeen,” “fifty,” and “thirty nine.”

Including grammar files

In many cases the grammars for a particular application are defined in a single main grammar specification file. However, you can use the `#include` directive to include a grammar file in a static grammar. This lets you create modular subgrammars that you can later include in the grammar file for one or more applications, keeping your application grammar files smaller and simpler. Included grammar files can contain both subgrammars and top-level grammars.

When the grammar file is compiled, any GSL line of the form

```
#include "filename.grammar"
```

is replaced by the contents of the file *filename.grammar*. The grammar compiler searches for your include files locally—that is, in the directory where the compiler is run. If you want to include a Nuance sample grammar, copy this grammar from the *grammars* directory (for example, `%NUANCE%\data\lang\<language>\grammars\`) to the directory where the compiler is run.

An included file may itself contain `#include` lines. Using the `#include` directive is necessary when your grammars are defined in more than one file, since you can pass one file only as an argument to the command-line compiler.

Default grammar location

The default directory location for American English grammars is `%NUANCE%\data\lang\English.America\grammars`.

Grammars for languages other languages are installed in the directories `%NUANCE%\data\lang\<language>\grammars`, where `<language>` indicates one of the supported natural languages—for example, `%NUANCE%\data\lang\Spanish.America\grammars`.

The grammar files installed in the directories `%NUANCE%\data\lang\<language>\grammars` cover common application vocabularies such as numbers, dates, money amounts, and confirmation (yes/no) responses. You can copy these files into your package directory and include them in your grammars, using them verbatim or editing them as needed.

The name format of the grammar files in the *English.America* directory is `<name>.grammar-v<vnumber>`, where *vnumber* denotes the version of the grammar. For example, to use the latest version the *money. grammar*, make a copy of that file into your working area, and then use the following directive in other files, as needed:

```
#include "money.grammar"
```


Caution: Be careful with the use of the include directive. Multiple inclusions of a grammar file result in compilation errors, as all grammar names passed to the compiler must have exactly one definition.

Recursive grammars

Recursive grammars—grammars that reference themselves—are only allowed in certain cases. Here is a simple example of a recursive grammar:

```
SENT [(see spot run) (SENT and SENT)]
```

Recursion can also occur indirectly, as in the following grammar description:

```
NounPhrase (NounPhrase1 *PrepositionalPhrase)
NounPhrase1 (Determiner Noun)
Determiner [the a]
Noun [cat dog]
Preposition [from with]
PrepositionalPhrase (Preposition NounPhrase)
```

The NounPhrase grammar is recursive because it references *PrepositionalPhrase*, whose description references back to *NounPhrase*.

GSL does not support *left-recursive* grammars. Left recursion occurs when the self reference (direct or indirect) is located in the leftmost position. Any other type of recursion is valid in GSL.

For example, the following grammar will not compile because it is left recursive:

```
Digits (Digits Digit)
```

But the following one is right recursive, and will therefore compile:

```
Digits (Digit Digits)
```

This next one is middle recursive, and will also compile:

```
Digits (Digit Digits end)
```

Indirect left recursion is also prohibited, for example:

```
NounPhrase (Determiner Noun)
Determiner [ the
             a
             NounPhrase ]
```

The previous grammar is prohibited because *Determiner* appears in the initial position within *NounPhrase*, and *NounPhrase* appears in the initial position within *Determiner*. Note that while *NounPhrase* does not literally appear in the initial (leftmost) position in the definition of *Determiner*, it is treated as such

because it appears in an OR construction where the order is not meaningful. An equivalent definition for the *Determiner* subgrammar could also be:

```
Determiner [ NounPhrase
              a
              the ]
```

where the *NounPhrase* subgrammar more obviously appears in the leftmost position.

There is an important requirement to remember when you are compiling a valid recursive grammar (that is a non-left recursive grammar): you *must* use the *-dont_flatten* compiler option. If *-dont_flatten* is omitted from the compilation options, you get a fatal error (stating that a grammar is recursively calling itself) even when you are using valid recursion in that grammar. See “Creating unflattened grammars” on page 110 for details on this compiler option.

Grammar probabilities

The recognition engine uses probabilities to weight constructs while searching for matches in its data space. These weights force the recognizer to favor certain phrases over others. Typically you assign higher probabilities to phrases expected to be spoken more frequently.

Adding probabilities to your grammar can potentially increase both recognition accuracy and speed—however, assigning bad probability values can actually hurt recognition performance. Grammar probabilities are recommended only for large vocabulary grammars (over 1000 words) where probabilities can be accurately estimated based on real usage.

Probabilities in GSL

You can assign probabilities to GSL constructs by using the tilde character (~). The GSL syntax to specify a probability for a construct C is:

```
C~prob
```

where *prob* is a non-negative number.

Probabilities are mostly used in disjunct (OR) constructs. For example, the following *City* grammar assigns different probabilities to each of the four city names:

```
City [ boston~.2
       (new_york)~.4
       dallas~.3
```

```

    topeka~.1
]

```

Specifying no probabilities is the same as specifying a probability of 1.0 for each item—effectively stating that any given phrase is as likely as any other.

You can also assign probabilities to the operand of optional (?), kleene closure (*), and positive closure (+) operator. The meaning of such an expression is illustrated in the following table:

Table 3: Grammar probability expressions

Expression	Meaning
? A~.6	A is 60% likely
+ A~.6	The probability of any additional occurrence of A (after the first one) is 60%
* A~.6	The probability of each occurrence of A (including the first one) is 60%

In regards to assigning overall and individual probabilities, note that the following two probability specifications are equivalent:

```

[A~.3 B~.3 C~.1]~.5
[A~.15 B~.15 C~.05]

```

Note: Nuance recommends that, if you choose to use grammar probabilities, you test your application’s performance both with and without probabilities. The *nuance-compile* program has the option *-dont_use_grammar_probs* that allows you to compile a package ignoring any probability specifications found in your grammars. This feature is useful when you want to compare recognition performance of a package with and without the use of probabilities.

Identifying a text string as Nuance GSL

To specify that a text string is actually a set of grammar rules in Nuance GSL, you can use the `;GSL2.0` header, as follows:

```

;GSL2.0
TEST_SENTENCE:public (the quick COLOR fox jumped over the TYPE
dog)
COLOR [red blue brown]
TYPE [lazy diligent]

```

This header specifies that the content that follows is a set of Nuance GSL rules. For example, the sample code above identifies `TEST_SENTENCE`, `COLOR`, and `TYPE` as GSL rules.

The header can be specified inside a Nuance GSL file, as shown above, or to an API function directly when using a just-in-time grammar, as shown in the following example:

```
Recognize(" ;GSL2.0
          PEOPLE:public <http://server/people.gsl>");
```

Nuance strongly recommends that you begin to use it in all your GSL grammars; it will become required in a future version.

Language-specific grammar conventions

Through deployment experience, Nuance has developed some general recommendations and language-specific conventions that you should conform to when writing grammars. For the most up-to-date information for the language you are working with, see the documentation section of the Nuance Developer Network website at extranet.nuance.com.

Dynamic grammars: Just-in-time grammars and external rule references

3

A *dynamic grammar* is a grammar that can be dynamically created and modified while an application is running. The Nuance dynamic grammar mechanism lets you create and update grammars at runtime and use them for recognition immediately, without needing to recompile the recognition package and start a new *recserver* process.

A dynamic grammar can be a file referenced using external rule references, or it can be created directly in a database using API functions. This chapter first introduces dynamic grammars created using API functions. It then describes how to refer to dynamic grammars using external rule references and how to use just-in-time grammars. It also provides information on configuring your system to use external rule references and just-in-time grammars.

Creating a dynamic grammar in a database

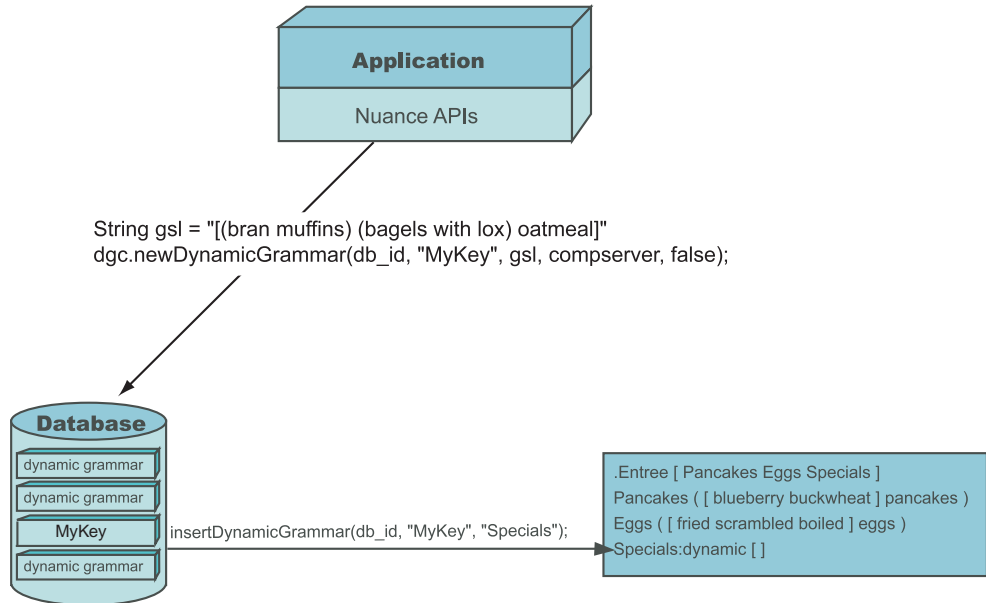
You can create a GSL or W3C XML (GrXML) grammar directly in a database using API functions. To create a dynamic grammar directly in a database, you first add the grammar as a component in a top-level grammar using the following GSL syntax:

```
GrammarName:dynamic []
```

This specification acts as a placeholder for the dynamic grammar that your application inserts at runtime. For example, this shows a simple GSL file that specifies a grammar for a breakfast menu that includes daily specials:

```
.Entree [ Pancakes Eggs Specials ]
Pancakes ( [ blueberry buckwheat ] pancakes )
Eggs ( [ fried scrambled boiled ] eggs )
Specials:dynamic []
```

When you compile the grammar *.Entree*, it contains the phrases “blueberry pancakes,” “buckwheat pancakes,” “fried eggs,” “scrambled eggs,” and “boiled eggs,” but the *Specials* subgrammar is empty. To fill the *Specials* grammar at runtime, your application finds (or creates, as shown in the example below) the correct dynamic grammar in the appropriate database and inserts it into the *.Entree* grammar at the location *Specials*:



After inserting this dynamic grammar, the grammar *.Entree* contains the phrases “bran muffins,” “bagels with lox,” and “oatmeal” in addition to its previous contents. You can specify whether these phrases should remain in the grammar indefinitely or only for the duration of the current call.

A dynamic grammar can also fill *slots*. For information about slots, see Chapter 4.

Note: You can have up to 255 `:dynamic` labels in a package. Also, a dynamic grammar must be a subgrammar in another top-level grammar—it cannot be a top-level grammar itself. It can, however, be as simple as:

```
.DynamicGrammar MyDynamicGrammar
MyDynamicGrammar:dynamic [ ]
```

In this case you can perform recognition using the top-level grammar *.DynamicGrammar*, whose contents are completely dynamic.

You must compile a dynamic grammar and a static grammar using the same master package. For complete details on creating GSL and GrXML dynamic grammars, see the *Nuance Application Developer's Guide*.

You cannot have a dynamic grammar in one package and a top-level grammar—of which that dynamic grammar is a subgrammar—in another package, and use both grammars at once. Both dynamic and static grammars must be compiled into a single package, if they are to be used together.

Note: As of Nuance 8.0, information kept in a dynamic grammar database can instead be kept in a grammar on file and loaded on demand, as provided by the just-in-time and external rule references grammar features. If your application requires the reliability, redundancy, and scaling associated with a commercial database, Nuance recommends that you keep your grammars in a dynamic grammar database. Otherwise, a just-in-time grammar or a dynamic grammar referenced using an external rule reference may be a more flexible approach.

External rule references and redefinitions

External rule references let you store dynamic grammars not only in a database but also in a file system or on a web server.

This section first describes how to reference and redefine external rules. It then presents the external grammar reference syntax.

Note: Your system configuration must include a web server for the external rule references to work. See “Configuring your system for external rule references and just-in-time grammars” on page 35 for more information.

Referring to external grammar rules

From any GSL grammar, whether static or just-in-time, you can refer to grammar rules in another grammar file. Nuance provides the following access modes to refer to external grammar rules:

- `http:`—points to a grammar on a web server
- `file:`—points to a grammar on disk
- `dgdb:`—points to a grammar in a dynamic grammar database
- `static:`—points to a precompiled static grammar
- `local:`—points to a public rule in the same grammar document
- `special:`—points to a grammar that modifies the search probabilities

- `fail:`—specifies a failure condition

The access modes are described below.

Using `http:` or `file:` URIs

To refer to a grammar residing on a web server, specify the grammar URI inside angle brackets, as shown in the following example:

```
;GSL2.0
TEST_SENTENCE:public (the quick COLOR fox jumped over the
    <http://types.com/todayslist.gsl> dog)
COLOR [red blue brown]
```

This GSL code tells the recognition server to get the list of types, at recognition time, from the specified `http:` URI and insert it in the specified grammar. Each time the `TEST_SENTENCE` grammar is used, the latest version of *todayslist.gsl* is fetched.

To refer to an external grammar rule on disk, specify the keyword `file:` followed by the location of the grammar, and place this reference inside angle brackets, as shown in the following example:

```
;GSL2.0
TEST_SENTENCE:public (the quick COLOR fox jumped over the
    <file:/c:/user/local/types.gsl> dog)
COLOR [red blue brown]
```

This GSL code gets the list of types from the specified file location and inserts it in the current grammar.

The `http:` and `file:` URIs can refer to the following types of files:

- GSL files, with the *.gsl* extension

```
http://gramserv.com/groceries.gsl
file:/c:/tmp/groceries.gsl
```

- W3C XML (GrXML) grammar files, with the *.grxml* extension

```
http://gramserv.com/groceries.grxml
file:/c:/tmp/groceries.grxml
```

- Nuance Grammar Object (NGO) files, with the *.ngo* extension

```
http://gramserv.com/groceries.ngo
file:/c:/tmp/groceries.ngo
```

Nuance provides the Nuance Grammar Object (NGO) grammar format. A Nuance Grammar Object is a grammar that is precompiled using the *nuance-compile-ngo* utility. This utility can compile a GSL, GrXML (W3C format), or SLM grammar file into a binary grammar file with the *.ngo* extension. The

compiled file can then be loaded on demand by a just-in-time grammar. Since an NGO file is precompiled, it is faster to load at runtime.

See “Dynamic compilation to a file” on page 96 for more information about the *nuance-compile-ngo* utility.

Using dgdb: URIs

You can refer directly to compiled dynamic grammars in a dynamic grammar database by specifying the keyword `dgdb:` followed by the database key and database descriptor. The syntax for the `dgdb:` URI is:

```
dgdb:?key=val1&dbdesc=val2
```

where *val1* is the database key and *val2* is the database descriptor for the grammar. Both *val1* and *val2* must be URI-encoded, which means that no illegal characters are used. The list of characters that cannot be used include all characters other than:

- a-z
- A-Z
- 0-9
- ; / ? : @ & = + \$, - _ . ! ~ * ' ()

The illegal characters must be converted to hexadecimal and written as `%xx`. For example, the space character, ASCII 32, becomes `"%20"`. For more information about URI encoding, see RFC 2396 (www.ietf.org/rfc/rfc2396.txt, section 2.4)

For example, in the following reference, the portion in bold must be URI-encoded:

```
dgdb:?key=PeterGrammar&dbdesc=provider=fs,root=/usr/dbs,name=db1,  
class=dgdb
```

To convert a `DBDescriptor` object to the URI-encoded string required by the `dgdb:` URI scheme, use the Nuance `DBDescriptor` function

`DBDescriptorToQueryString()`. See the *Nuance API Reference* for more information about this function.

Using static: URIs

A grammar can explicitly refer to a `:dynaref` or `:public` rule in the host package by using the `static:` keyword. For example, in the following code, the grammar refers to the static grammar rule *OtherInterestingSentences*:

file:/c:/usr/grammars/other.gsl

```
;GSL2.0
TEST_SENTENCE:public
[(the quick brown fox)
  <static:OtherInterestingSentences>
]
```

The static grammar rule to which the grammar refers must be identified by the `:dynaref` or the `:public` keyword in the static grammar file. For example, the static grammar rule *OtherInterestingSentences* in the example above would have been defined as following in a grammar file, and then compiled with *nuance-compile*:

```
OtherInterestingSentences:dynaref [...]
```

To compile these grammars into your application's recognition package, you must use the *nuance-compile* option *-dont_flatten*.

Note that, when compiling the grammar that refers to a static grammar rule, the compiler does not check that the static grammar exists. For example, when compiling the above grammar *other.gsl*, the compiler does not check that grammar *OtherInterestingSentences* exists. This grammar must be present only when it is used at recognition time.

Fetching grammars

An external reference is not evaluated until recognition time. If the contents of a grammar change over time, the latest version of the grammar is fetched for each recognition. The Nuance implementation of grammar fetching is very efficient and is designed to minimize network traffic. Grammar fetching is configured with the parameter `config.EGRCacheMinFreshSecs`. This parameter specifies the time (in seconds) that a grammar referred with an `http:` or `file:` URI is assumed to be *fresh*. The default value is 10 seconds.

This parameter works as follows. When an application refers to the grammar, the compilation server loads it, compiles it, and stores the compiled version in its cache. For the next 10 seconds (or as specified by `config.EGRCacheMinFreshSecs`), all other requests for that grammar are served out of the cache without testing whether the source file has changed. The first request for that grammar after the 10 seconds have expired forces the compilation server to check the eTag of the grammar, via an HTTP request, or the timestamp of the grammar file. If the grammar has not changed, the cached copy is used for another 10 seconds. If the grammar has changed, it is fetched and recompiled.

This feature provides a nice way to minimize network usage. No matter how many times a grammar is requested (across any number of ports simultaneously), the grammar is examined only once every 10 seconds.

Another parameter, `config.EGRTIMEOUTMS`, is also used to configure grammar fetching. This parameter specifies the maximum number of milliseconds that are allotted for fetching all the external grammar components for a just-in-time recognition. Any requests that are still outstanding once the timeout is reached are reported as load failures with a “Timed out waiting for HTTP response” annotation in the recognition server log. These references are replaced with the backoff reference. See “Specifying a grammar backoff” on page 32 for more information.

Referring to a specific external rule

A grammar file often contains multiple grammar rules. When referring to an external grammar file in a grammar, if you do not specify a rule, the URI refers to the *root* rule by default. For GSL files (and NGO files compiled from GSL files), the root rule is defined as the first `:public` rule in the file.

Note: For W3C grammars (and NGO files compiled from W3C grammars), the root rule is determined according to the W3C grammar specification, described in *Speech Recognition Grammar Specification for the W3C Speech Interface Framework*. www.w3.org/TR/2001/WD-speech-grammar-20010820/ for more information. Nuance is currently working from the draft dated **August 20, 2001**.

You can also refer to a specific rule in a grammar file by adding the `#` symbol followed by the rule name, as shown below:

http: URI	http://URI_of_grammar#RuleName
For example	
	http://gramserv.com/groceries.gsl#Nuts
file: URI	file:/location_of_grammar#RuleName
For example	
	file:/c:/usr/grammars/groceries.gsl#Nuts
dgdb: URI	dgdb:DBDescriptor_and_key#RuleName
For example	
	dgdb:?key=PeterGrammar&dbdesc=provider=fs,root=/usr/dbs,name=db1,class=dgdb#Nuts

Defining a rule as public

When making a reference to a specific external rule, you can specify a public rule only. A rule is defined as a *public* rule if it is followed by the `:public` or `:dynaref` keyword. For example, all of the rules in this GSL file are public:

```
;GSL2.0
MyGrammar1:public [...]
MyGrammar2:dynaref [...]
```

In a GSL file, if none of the rules in a grammar file are declared public, then all rules are automatically made public, and the first rule becomes the root rule. In this case, a warning message is displayed. For consistency and clarity, Nuance recommends that you explicitly identify the rules that you want to make public with the `:public` keyword.

Examples

For example, in the following grammar files, both *Fruits* and *Nuts* are public rules, and the *Fruits* rule is the root rule.

<http://server/groceries.gsl>

```
;GSL2.0
Fruits [apple orange]
Nuts [peanut cashew]
```

<file:/c:/usr/grammars/groceries.gsl>

```
;GSL2.0
Fruits:public [apple orange]
Nuts:public [peanut cashew]
```

Given the files above, the following URI refers to the *Fruits* rule:

`http://server/groceries.gsl`

While the following URI refers to the *Nuts* rule:

`file:/c:/usr/grammars/groceries.gsl#Nuts`

Finally, consider the following example:

<http://server/groceries.gsl>

```
;GSL2.0
Fruits [apple orange]
Nuts:public [peanut cashew]
```

In this example, only the *Nuts* rule is public. Therefore, you cannot refer to the *Fruits* rule directly from another grammar. Given this file, the following URI refers to the *Nuts* rule:

```
http://server/groceries.gsl
```

Redefining an external rule

You can dynamically redefine an external dynamic rule in a grammar by using the `!` symbol, as follows:

```
< <Grammar_File_URI> ! RuleName = <New_URI> & ... >
```

where `RuleName` is tagged with the `:dynamic []` keyword in the grammar file specified by `Grammar_File_URI`. Redefining an external rule is very useful in personalized applications, where the grammars need to be customized per users.

For example, consider the following grammar file for a personal dialing application:

file:/c:/usr/local/template.gsl

```
;GSL2.0
MAIN_MENU:public [ (please call MyFriends)
    help
    cancel
    Other_Cmds ]
Other_Cmds:dynamic []
MyFriends:dynamic []
```

A grammar can refer to this file and redefine the *Other_Cmds* and *MyFriends* rules dynamically as follows:

```
;GSL2.0
My_Main_Menu:public < <file:/c:/usr/local/template.gsl> !
    Other_Cmds = <file:/c:/usr/local/cmd.gsl> &
    MyFriends = <http://listserver/user18257_contactlist.ngo>>
```

In this example, the *Other_Cmds* rule is redefined with a grammar available on a file system, while the *MyFriends* rule is redefined with the grammar for the caller's contact's list. Note that you can also redefine an external rule to one that is defined locally using the `local:` keyword. See "Redefining an external rule to a local rule" on page 43 for more information.

An external rule redefinition is very similar to a dynamic grammar insertion, but it has many advantages:

- The inserted grammar can come from any type of URI, not only from a dynamic grammar in a database
- It is specified in the grammar itself; using a dynamic grammar function call is not required
- It is not persistent—it lasts for one utterance only

Note that any grammar rule previously inserted with the `InsertDynamicGrammar()` function is cleared after an external rule redefinition. Nuance does not recommend that you mix external rule redefinitions with dynamic grammar insertion functions in your application.

Specifying a grammar backoff

You can specify a *grammar backoff*—the alternate grammar to use if a grammar referenced using a `file:`, `http:`, or `dgdb:` URI cannot be found.

To specify the grammar backoff, you use the `%` symbol, as follows:

```
< <URI_1> % <backoff_URI> >
```

For example:

```
;GSL2.0
MAIN_MENU:public (<<http://polite.com/nice.gsl > %
                  <file:/c:/user/local/MyNice.gsl>>
                  get me a quote for STOCK)
STOCK:dynamic []
```

In this example, the *nice.gsl* grammar will be replaced by the local grammar file *MyNice.gsl* if the HTTP request fails. If the local file is not found, it is substituted with the global grammar backoff defined with parameter `egr.Backoff`. This parameter specifies a default backoff grammar that is used when any URI reference fails, even if you do not use the `%` symbol. To do so, you set parameter `egr.Backoff` to the backoff grammar.

You can also use the following special references as grammar backoffs:

- `special:passthrough`—defines a rule that is automatically matched, without the user speaking any word
- `special:roadblock`—defines a rule that can never be spoken
- `special:resistor`—lets you change the probability that any given rule will be taken instead of simply disabling the rule

The following example shows how special references can be used as grammar backoffs.

```

;GSL2.0
MAIN_MENU:public
    (?<<http://polite.com/nice.gsl> % <special:passthrough>>
    get me a quote for
    <<http://broker.com/stock.gsl> % <special:roadblock>>)

```

In this example, if the first `http:` request fails, the *nice.gsl* grammar is replaced by a `passthrough` and is therefore simply ignored. However, if the *stock.gsl* grammar cannot be found, it is replaced by a `roadblock`, which means that recognition will never happen for the *MAIN_MENU* rule.

For more information about `passthroughs` and `roadblocks` and about the `special:resistor` keyword, see “Using the dynamic grammar gate technique” on page 46.

Failed references

To cause a recognition to fail immediately and return an `EXCEPTION RecResult`, you can use the `fail:` reference. The text of the recognition result will include all the grammar references that generated a failure as well as their source location.

For example, consider the following *http://up/foo.gsl* grammar:

```

;GSL2.0
Foo:public <http://down/foo.gsl>

```

Then consider the following just-in-time grammar:

```

;GSL2.0
ROOT:public (<http://down/foo.gsl>
    <<http://down/bar.gsl> % <fail:died>>
    <http://up/foo.gsl>)

```

Suppose that the `http://down` server is down. Since the just-in-time grammar includes a `fail:` reference, trying to perform recognition with the just-in-time grammar will create an `EXCEPTION RecResult` similar to the following:

```

<http://down/foo.gsl> failed: couldn't load
<<http://down/foo.gsl> % <fail:died>> failed
In <http://up/foo.gsl>, <http://down/foo.gsl> failed: couldn't
load

```

The first reference failed because the `http://down` server is down. The second reference also failed, but because we used the `fail:` keyword, the `<fail:..>` reference is also included in the `RecResult`, providing more information as to why the reference failed. The third reference failed because the *http://up/foo.gsl* grammar contains a reference to a grammar on the server that is down.

Note that if you do not specify a grammar backoff as described in “Specifying a grammar backoff” on page 32, the default `fail: reference` “`<fail:couldn't load>`” is used, as specified by parameter `egr.Backoff`.

Relative paths in external rule references and redefinitions

When using the `file:` and `http:` URIs, a *relative path* is relative to the grammar file that contains the reference.

For example, consider the file `file:/c:/tmp/head.gsl` which contains the following string:

```
file:/c:/tmp/head.gsl
```

```
...<tail.gsl>...
```

An external rule reference to the file `file:/c:/tmp/head.gsl` will try to load the file `file:/c:/tmp/tail.gsl`.

Similarly, consider the file `http://megacorp.com/dave/test.gsl` which contains the following string:

```
http://megacorp.com/dave/test.gsl
```

```
...<../top.gsl>...
```

An external rule reference to `http://megacorp.com/dave/test.gsl` will try to load the file `http://megacorp.com/top.gsl`.

Note that for grammar backoffs and external rule redefinitions, a relative path is relative to the grammar file containing the backoffs and external rule redefinitions, not to the rules they are referencing or redefining.

For example, consider the file `file:/c:/tmp/dave/top.gsl` which contains the following string:

```
...<<file:/c:/nuance/car.gsl> ! Mine=<hiscars.gsl>>  
%<allcars.gsl>>...
```

An external rule reference to `C:\tmp\dave\top.gsl` will resolve references to files `hiscars.gsl` and `allcars.gsl` in the `C:\tmp\dave` directory, and not in the `C:\nuance` directory.

If there is no parent insertion, a relative path is resolved as follows:

- If the relative reference has an `http:` scheme, it is not resolved. Nuance does not recommend using a relative `http:` reference.

- If the reference has a `file:` scheme, and parameter `egr.file.context` is not set, the relative path is relative to the directory in which the compilation server was started. If parameter `egr.file.context` is set, all file references (whether relative or not) are resolved relative to that path. See “Parameters for external rule references” on page 36.
- If the reference does not have a scheme, then it is resolved only when a `Content-Base` key is specified in the `RECOGNIZE` header of the JIT grammar, as described in “Specifying just-in-time grammar options: Using the keyword `RECOGNIZE`” on page 40. For example, consider the following grammar:

```
RECOGNIZE
Content-Base: http://up/here/

;GSL
ROOT: <foo.gsl>
```

In this example, the `<foo.gsl>` reference is resolved as `<http://up/here/foo.gsl>`.

Specifying builtin: grammars

You can create built-in grammars and reference them using the `builtin:` URI scheme. To use this scheme, you set the `egr.builtin.context` parameter to the location of the built-in grammars, for example “`http://tango:8080/`”. You then use the `builtin:` URI as follows:

```
Rule:public (i want <builtin:grammar/digits> scoops of ice cream)
```

To find this grammar, the recognition server takes the string that follows the `builtin:` keyword and prepends the value of the parameter `egr.builtin.context`. For example, if parameter `egr.builtin.context` is set to “`http://host/`”, the expression `<builtin:grammar/digits>` is expanded to `<http://host/grammar/digits>`.

Note: This feature is required by VoiceXML. The Nuance System does not provide built-in grammars.

Configuring your system for external rule references and just-in-time grammars

This section provides information on configuring your system for external rule references and just-in-time grammars.

Setting up a web server for external rule references

For the URI-based external rule references to work, your system configuration *must* include a web server.

You need to define the MIME content types for the grammar files as follows. Refer to your web server documentation for information on setting MIME types.

Grammar file extension	MIME type
.gsl	application/x-nuance-gsl
.ngo	application/x-nuance-dynagram-binary
.grxml	application/grammar+xml

The Nuance System handles downloaded grammar documents as follows. It first looks at the web server content type. If it's one of the content types defined above, it handles the file appropriately. If it's an unknown content type, it looks at the file extension to determine how to handle the file. If that fails, it examines the document header: a header of `;GSL2.0` indicates a GSL file, while a header of `<?xml ... >` indicates a W3C XML file. If that fails, the Nuance System resolves the reference in a platform-specific manner. If the grammar document type cannot be resolved, a failure is reported.

Configuring the compilation server

Since the compilation server performs caching of dynamic grammars, Nuance recommends that you configure your system to share as few compilation servers as required for redundancy. While the system works with multiple compilation servers, grammars may get compiled by many compilation services and reside in many caches.

If you have a large system, run the compilation servers on a powerful computer with a lot of RAM, and increase the value of parameter `config.EGRCacheMB`.

Parameters for external rule references

The following table describes the parameters that apply to external rule references.

Note: The *compilation-server* executable only uses parameter `config.EGRProxy` and ignores all other parameters. To specify these options when starting the *compilation-server*, use the appropriate *compilation-server* options as described in the *Nuance API Reference*.

Parameter name	Description
<code>egr.Backoff</code>	Specifies the global grammar backoff used when an URI reference in a grammar fails.
<code>config.EGRCacheMB</code>	Specifies the maximum amount of dynamic grammars to be cached in memory. A larger value will avoid recompiling some <code>http:</code> and <code>file:</code> grammars but will increase the process size accordingly. Note that the process DOES NOT preallocate the space specified by this parameter. It just accumulates grammars until it hits the maximum amount specified and then starts flushing the least-recently-used ones. Applies to the compilation server and recognition server.
<code>config.EGRCacheMinFreshSecs</code>	Specifies the time (in seconds) an <code>http:</code> or <code>file:</code> grammar is assumed to be “fresh”. See “Fetching grammars” on page 28 for more information.
<code>config.EGRProxy</code>	If this parameter is set, the compilation server and recognition server do not fetch <code>http:</code> grammars directly, but instead fetch them via the specified HTTP proxy server. Valid values are <i>machinename</i> or <i>machinename:port</i> , for example, <i>mycomputer:1234</i> . This parameter is useful if you want all the HTTP requests leaving your site to go through a proxy. Applies to the compilation server and recognition server.
<code>config.EGRTimeoutMS</code>	Specifies the maximum number of milliseconds that are allotted for fetching all the external grammar components for a just-in-time recognition. Any requests that are still outstanding once the timeout is reached are reported as load failures with a “Timed out waiting for HTTP response” annotation in the recognition server log. These references are replaced with the backoff reference. See “Specifying a grammar backoff” on page 32 for more information.
<code>egr.builtin.context</code>	Specifies the location of the built-in grammars in the system.
<code>egr.file.context</code>	Specifies a base <code>file:</code> URI; all <code>file:</code> URIs in the grammar are relative to this base URI.

Parameter name	Description
<code>egr.dgdb.context</code>	<p>Specifies the DBDescriptor for all the <code>dgdb:</code> references. This is useful when your application uses the same DBDescriptor for most of the <code>dgdb:</code> requests. Set this parameter to the DBDescriptor content, excluding the <code>dbdesc</code> keyword. For example:</p> <pre>"provider=fs,root=c:/Temp,class=dgdb,name=Nuance"</pre> <p>The following <code>dgdb:</code> references would then fetch the same item:</p> <ul style="list-style-type: none"> ▪ <code><dgdb:?key=foo&dbdesc=provider=oci,server=nuance,name=Nuance,class=dgdb></code> ▪ <code><dgdb:?key=foo></code> ▪ <code><dgdb:foo></code> ▪ <code><dgdb:?key=foo&dbdesc=provider=fs,root=c:/Temp,class=dgdb,name=Nuance></code>

Just-in-time grammars

The Nuance System provides *just-in-time* grammars—grammars that let you specify grammar expressions directly as the starting point of a recognition. This section first describes how to write a just-in-time grammar. It then describes how to specify the compilation options required by the recognizer to perform just-in-time compilation.

Writing a just-in-time grammar

A just-in-time grammar is a grammar that lets you specify GSL expressions as well as GSL and GrXML content directly as the starting point of a recognition. A just-in-time grammar does not need to be precompiled. You can define, compile, and use a just-in-time grammar in only one step. You simply pass content directly to the recognizer at runtime through an API call, as shown in the following example:

```
String gsl = "[hello (hi there)]";
jsc.playAndRecognize(gsl);
```

where `jsc` is a `nuance.core.sc.NuanceSpeechChannel`.

The recognition server compiles the grammar “just in time” for the recognition and then discards the grammar after the recognition.

Note that the string passed to the method in the example above is an *extended right-hand side*: it does not contain a name for the rule and does not constitute by itself a complete valid GSL definition. Specifying a GSL fragment only is useful when you do not need to specify a complete set of rules.

You can specify an actual set of GSL or GrXML rules as a just-in-time grammar, as shown in the following example:

```
String gsl = ";GSL2.0 \n" +
    "Rule1:public (please call <http://myhost/people.gsl>)" ;
jsc.playAndRecognize(gsl);
```

You can also specify a grammar document through an external rule reference. For example, consider the following file:

<http://host/grammars/MyGrammar.grxml>

```
<?xml version="1.0" ?>
<grammar xmlns="http://www.w3.org/2001/06/grammar" version="1.0"
    xml:lang="en-US">
  <rule id="main">
    <one-of>
      <item>hi</item>
      <item>bye</item>
    </one-of>
  </rule>
</grammar>
```

You use this grammar just-in-time by specifying a reference to the file through a recognition API function, as follows:

```
String grxml = "<http://host/grammars/MyGrammar.grxml>";
jsc.playAndRecognize(grxml);
```

Here are more examples:

```
;external reference to a static grammar rule
String gsl = "<static:Names>";
jsc.playAndRecognize(gsl);
```

```
;external reference to a dynamic grammar rule
String gsl = "<dgdb:?key=MyGram&dbdesc=provider=fs," +
    "root=/usr/dbs,name=db1, class=dgdb>";
jsc.playAndRecognize(gsl);
```

```
;external reference to a GSL file
String gsl = "<file:/c:/nuance/names.gsl>";
jsc.playAndRecognize(gsl);
```

Specifying just-in-time grammar options: Using the keyword RECOGNIZE

You can specify a number of options related to just-in-time grammar compilation and usage with the keyword `RECOGNIZE`. This keyword takes a set of `<key>:<value>` pairs, followed by a blank line, followed by a GSL or GrXML document. For example:

```
RECOGNIZE
Nuance-Package-Name: French.1
Nuance-Config-Name: comp-server-french

;GSL2.0
Rule1:public (please call <http://myhost/people.gsl>)
```

Here is an example of a GrXML grammar:

```
RECOGNIZE
Nuance-Package-Name: French.1
Nuance-Config-Name: comp-server-french

<?xml version="1.0" ?>
<grammar xmlns="http://www.w3.org/2001/06/grammar" version="1.0"
  xml:lang="en-US" root="main">
  <rule id="main" scope="public">
    please call <ruleref uri="http://myhost/people.gsl"/>
  </rule>
</grammar>
```

Note: Any whitespace on either side of the colon in the `<key>:<value>` pair is ignored.

Nuance supports three keys:

Key	Description
Nuance-Config-Name	Specifies the compilation server to use for a just-in-time recognition request. See “Choosing between multiple compilation servers” on page 41.
Nuance-Package-Name	Specifies the master package to use for a just-in-time recognition request. See “Multi-package support” on page 42.
Content-Base	Specifies a base URI; the relative URIs in the grammar are relative to this base URI. See “Specifying the base for relative URIs” on page 43.

Choosing between multiple compilation servers

By default, when you use a just-in-time grammar, each external grammar reference is sent to any available compilation server. However, some applications may need to work with multiple, differently configured compilation servers. For example:

- You need to turn optimization off for some grammars (using parameter `comp.Optimize`) because they would take too long to compile, but you want to keep it on for other grammars.
- You want to set the number of pronunciations generated (using parameter `pron.MaxProns`) to a small number for large grammars, but to a large number for small grammars.

In cases like this, you can run separate compilation servers, each with the appropriate compilation options.

To run a system with multiple compilation servers:

- 1 Specify a name for each server you start using the `-config_name` argument. For example:

```
> compilation-server -package English.America.1
pron.MaxProns=10 -config_name SmallGramServer
> compilation-server -package English.America.1 pron.MaxProns=3
-config_name LargeGramServer
```

These commands start two compilation servers, *SmallGramServer* and *LargeGramServer*.

- 2 When you call a recognition function with a just-in-time grammar, specify the name of the compilation server using the `Nuance-Config-Name` key-value pair in the `RECOGNIZE` header of your just-in-time grammar specification, as follows:

```
String gsl = "RECOGNIZE\n" +
    "Nuance-Config-Name: LargeGramServer\n" +
    " \n" +
    ";GSL2.0" +
    "PEOPLE:public <http://server/people.gsl>";
jsc.playAndRecognize(gsl);
```

In this example, grammar *people.gsl* will be compiled using the *LargeGramServer* compilation server, which only generates three pronunciations.

Note: The words in the just-in-time grammar itself are compiled by the recognition server, not the compilation server. For example, in the following grammar, the words "hello hi there" are compiled by the recognition server. Therefore, in such a case, the GSL option `Nuance-Config-Name` key-value pair is ignored.

Multi-package support

```
String gsl = ";GSL2.0 \n" +  
            "[hello (hi there)]";  
jsc.playAndRecognize(gsl);
```

You can use just-in-time grammars with multiple Nuance master packages. To do so, you compile your grammar with the different master packages and you specify which master package to use for a specific recognition.

To use a just-in-time grammar with multiple packages:

- 1 Use *nuance-compile* to compile a grammar with the different master packages. For example, consider an application that needs to perform recognition both in French and English using the same grammar. You compile the grammar as follows:

```
> nuance-compile MyEnglish.grammar English.America.1  
-enable_jit  
> nuance-compile MyFrench.grammar French.1 -enable_jit
```

These commands will create two recognition packages. Each package will be given a default name, which is the name of the master package (*English.America.1* and *French.1*). If necessary, you can override the default package name with the *-package_name* option, as follows:

```
> nuance-compile MyEnglish.grammar English.America.1  
-enable_jit -package_name MyEnglishPackage  
> nuance-compile MyFrench.grammar French.1 -enable_jit  
-package_name MyFrenchPackage
```

- 2 When you perform recognition, specify the package name to use by setting the *Nuance-Package-Name* key-value pair in the *RECOGNIZE* header of the just-in-time specification. For example:

```
String gsl = "RECOGNIZE\n" +  
            "Nuance-Package-Name: French.1 \n" +  
            " \n" +  
            ";GSL2.0" +  
            "Greetings:public [salut bonjour]";  
jsc.playAndRecognize(gsl);
```

If you have overwritten the default package name, you must specify the new package name, as follows:

```
String gsl = "RECOGNIZE\n" +  
            "Nuance-Package-Name: MyFrenchPackage \n" +  
            " \n" +  
            ";GSL2.0" +  
            "Greetings:public [salut bonjour]";  
jsc.playAndRecognize(gsl);
```


Specifying the base for relative URIs

The `Content-Base` key lets you use relative URIs in your grammar document. The relative URIs in the grammar are relative to the base URI specified in the `Content-Base` key.

Consider the following grammar:

C:\tmp\head.gsl

```
RECOGNIZE
Content-Base: http://yahoo.com/grammars/stocks/

;GSL2.0
Rule1:public (please call <../people/work.gsl>)
```

The `<../people/work.gsl>` URI will be expanded to `<http://yahoo.com/people/work.gsl>`.

Note that any character following the final “/” character in the `Content-Base` value is ignored. For example, the following specification is the same as the specification in the example above:

```
Content-Base : http://yahoo.com/grammars/stocks/abcdef
```

Redefining an external rule to a local rule

The standard GSL syntax described in the *Nuance Grammar Developer's Guide* also applies to just-in-time grammars. You can also use the keyword `local:`, available in just-in-time grammars only. This keyword lets you redefine an external rule to one that is defined locally.

For example, consider the following file:

file:/c:/usr/local/template.gsl

```
;GSL2.0
MAIN_MENU:public [ (get me a quote for STOCK)
                    help
                    cancel
                    OTHER_CMDS ]
OTHER_CMDS:dynamic []
STOCK:dynamic []
```

The following just-in-time grammar redefines the `OTHER_CMDS` rule to a local one:

```
String gsl = ";GSL2.0 \n" +
    "MY_MAIN_MENU:public [<<file:/usr/local/template.gsl> ! \n" +
    "    OTHER_CMDS = <local:CMDS> & \n" +
```

```

        " STOCK = <http://stocks.com/list.gsl>>] \n" +
        " CMDS:public [up down left right] \n" +
        " ])" ;
jsc.playAndRecognize(gsl);

```

Note that, as with other references, the rule that you are redefining locally must be a public rule.

Enabling compilation of just-in-time grammars

To compile a recognition package to use with a just-in-time grammar, run *nuance-compile* with the *-enable_jit* option, as follows:

```
> nuance-compile my.grammar English.America.1 -enable_jit
```

In this example, the *English.America.1* package can be used to perform recognition of a just-in-time grammar. If you do not specify the *-enable_jit* option, your application can only perform recognition using the package's static grammars and contexts.

See "Compilation of just-in-time grammars" on page 98 for more information.

Note that when referring to an NGO grammar in a JIT grammar, the NGO grammar and the JIT grammar must have been compiled using the same master package. Otherwise, the reference to the NGO grammar will be treated as a *fail:reference*.

Usage scenarios

This section presents typical situations in which using just-in-time grammars allows a more efficient and simpler implementation of an otherwise complex piece of code.

Case 1: A personalized bill pay application

Consider a bank that lets its customers pay bills using a voice-recognition application. Each customer has a personalized list of a few dozen names. These lists are stored in a set of files on an internal web server; for example:

```

http://megabank.com/payees/customer1.gsl
http://megabank.com/payees/customer2.gsl
...

```

The contents of these files might look like this:

```

;GSL2.0
MyPayees:public [  (?my visa card)
                   (?the mortgage)
                   .....
                   ]

```

When a customer calls, for example customer #1, the application invokes recognition with a just-in-time grammar, as follows:

```
;GSL2.0
ROOT:public [(i ?really want to pay)
              (?please pay)]
          <http://megabank.com/payees/customer1.gsl>)
```

Variant 1: Using a precompiled grammar

In the example above, each payee list was stored as a GSL file on a web server. To load the files to the recognition server more quickly, you could also run *nuance-compile-ngo* on each file and refer to the *.ngo* file instead. The GSL and NGO files can also be referenced using `file:` URIs.

Variant 2: Using a dynamic grammar database

Using the `newDynamicGrammar` functions or the *dgdb-edit* utility, you could also store the compiled grammars in a dynamic grammar database, and point to the grammars using `dgdb:` URIs. Nuance recommends this approach if the company deploying the application requires the reliability, redundancy, and scaling associated with a commercial database.

Variant 3: Using speaker-trained grammars

If customers can add names to their list using voice enrollment—by speaking the names to add—you could store these grammars in a dynamic grammar database. At runtime, you point to the appropriate grammars using the `dgdb:` URI.

Case 2: Single-use GSL

Consider an application that deals with ambiguous responses from users by constructing a grammar with these ambiguous responses, playing a prompt to the caller such as “Did you mean ...?”, and performing recognition on that grammar.

Without just-in-time compilation, you would:

- Create a static grammar that includes a reference to a dynamic grammar and compile the static grammar
- At runtime, compile the GSL code dynamically and store in a dynamic grammar database
- Insert the dynamic grammar into a static grammar
- Recognize using that static grammar
- Delete the dynamic grammar

In contrast, using a just-in-time grammar, you would only use a code fragment like the following:

```
String gsl = "[Cisco Systems) (Sysco Foods)]";
jsc.playAndRecognize(gsl);
```

Case 3: A VoiceXML interpreter

A VoiceXML interpreter needs to recognize speech using all in-scope grammars in parallel. For instance, consider the following VoiceXML fragment (used for illustration purposes only):

http://host/vxml/main.vxml

```
<vxml>
  <link next="Navigation.vxml#GoForward">
    <grammar src="..\grammars\GoForward.gsl" />
  </link>
  <link next="Navigation.vxml#GoBack">
    <grammar type="application/x-nuance-gsl">
      (go back)
    </grammar>
  </link>
  <form id="mainmenu">
    <field name="mainmenu" slot="command">
      <grammar src="..\src\grams\roomorama2.gsl" />
    </field>
  </form>
</vxml>
```

Without just-in-time compilation, you would need to write fairly complex code to implement this logic. Recognizing with a just-in-time grammar is far simpler. You would write code to assemble these three grammars into one grammar and use it as follow:

```
String gsl = "[<http://host/grammars/GoForward.gsl>\n" +
              "(go back)\n" +
              "http://host/src/grams/roomorama2.gsl ]";
jsc.playAndRecognize(gsl);
```

Using the dynamic grammar gate technique

You can use dynamic grammars to dynamically enable or disable various branches in an otherwise static grammar, using what Nuance calls the *gate* technique. When you use the gate technique, you can explicitly “open” or “close” each branch of the grammar. You can also use gate *resistors* to dynamically change the probability associated with specific branches.

For example, in a personalized brokerage application you might only want to allow users to select a stock to sell if their portfolio actually includes that stock. When a given user wants to sell a stock, you insert a dynamic grammar into each of the placeholders that either opens the gate (allowing the grammar to be

recognized) or closes the gate (disabling the grammar), based on whether the current user owns each stock:

- To open a gate, refer to a dynamic grammar containing the following keyword:

```
special:passthrough
```

- To close a gate, refer to a dynamic grammar containing the following keyword:

```
special:roadblock
```

These expressions function as gates for the following reasons:

- An AND construction represents a set of expressions that must be matched. The keyword `special:passthrough` is an empty AND construction, which means that nothing needs to be matched. The gate is therefore “open”—the recognizer continues trying to match the other phrases in the larger grammar.
- An OR construction represents a list of alternatives, one of which must be matched. The keyword `special:roadblock` is an empty OR construction. It provides no alternatives and therefore a match is impossible. The gate is “closed”—the recognizer stops trying to match the utterance against the phrases in this grammar.

Note: In a Nuance GSL grammar, you can also use the `NULL` and `VOID` special keywords, as defined by W3C grammars: `NULL` and `passthrough` are synonyms; `VOID` and `roadblock` are synonyms.

The syntax of the special reference is shown in the following example:

```
;GSL2.0
TEST_SENTENCE:public
[
    (the quick <special:passthrough> brown fox)
    (the sentence <special:roadblock> that cannot be spoken)
]
```

In this example, the first sentence—“the quick brown fox”—can be recognized while the second sentence—“the sentence that cannot be spoken”—will never be matched.

To use dynamic grammar gates in an application, you can create a static grammar with dynamic grammar placeholders. At runtime, you can redefine the placeholders to the “open” dynamic grammar or “closed” dynamic grammar, as appropriate for the current user.

You can also use a *resistor* instead of an open or closed gate grammar. This lets you change the probability that any given branch will be taken instead of simply disabling the branch. To do this, insert a dynamic grammar record containing the following keyword:

```
<special:resistor?weight=XYZ>
```

where XYZ is the probability for the branch. For example:

```
<special:resistor?weight=0.5>
```

Probabilities that are less than 1.0 diminish the chance that the branch will be parsed. The closer the probability gets to 0, the less likely the branch is to be recognized. Specifying a probability of 0 is the same as closing a gate, and specifying a probability of 1 is the same as opening a gate. You can also assign probabilities that are greater than 0, for example:

```
<special:resistor?weight=1.5>
```

These cause the recognizer to favor that branch rather than avoid it. You might use this, for example, in a stock quote system where you want to boost the likelihood of recognizing a stock in the current user's portfolio without disallowing other stocks.

GSL syntax for enrollment

Enrollment is the feature that lets users add spoken phrases to a dynamic grammar. During enrollment, the recognition system listens to the user speak a phrase, generates pronunciation for that phrase by performing phonetic recognition of the utterance, and returns the pronunciation, along with a phrase identifier, to the application.

To perform voice enrollment, the recognition package of your application must include the special grammar *EnrollmentGrammar*, defined in the Nuance grammar file *enrollment.grammar*. This grammar enables phonetic recognition, and allows the system to generate pronunciations for spoken phrases. Add this as a subgrammar in a top-level grammar, for example:

```
; enrollment.grammar defines "EnrollmentGrammar"  
#include "enrollment.grammar"  
.PersonalPayeeList EnrollmentGrammar
```

When you perform enrollment, you use this as the recognition grammar. This grammar can also include any words or subgrammars needed for your application. For example, your application might allow the user to say things like "help" and "cancel" during enrollment, in which case your enrollment grammar should include those phrases. You compile this grammar just like any

other grammar. The resulting recognition package can be used for both standard (non-enrollment) recognition tasks, or to generate the pronunciations for utterances.

Note that the enrollment grammar is what you specify only during the enrollment process. To use this grammar for recognition, you still have to insert it into a static grammar, at a location defined by the syntax:

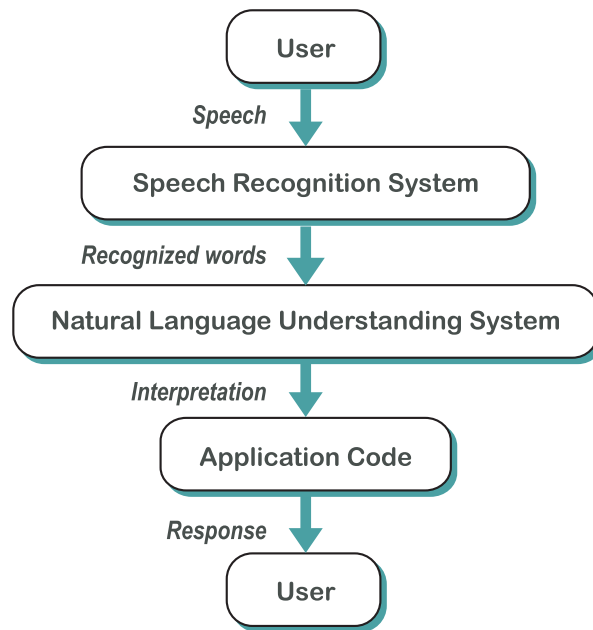
```
GrammarName:dynamic []
```

For more information on enrollment see the *Nuance Application Developer's Guide*.

Natural language understanding

4

The Nuance System lets you define natural language understanding in your recognition packages. A natural language understanding system takes a sentence (typically a recognized utterance) as input and returns an *interpretation*—a representation of the meaning of the sentence. The application code can then take action based on the user's request:



The natural language system simplifies the work your application needs to do to correctly respond to user utterances. Often, there are many ways a user can express the same meaning. For example, the following phrases:

“Withdraw fifteen hundred dollars from savings”
“Take fifteen hundred out of savings”
“Give me one thousand five hundred dollars from my savings account”

map to the same semantic units:

action = withdraw
amount = 1500.00
account = savings

The Nuance natural language understanding engine returns a simple, structured interpretation, based on a predefined set of *slots* with associated *values*. Applications can then access this interpretation directly without needing to parse the recognized text. You define the natural language commands for an application right in the grammar file, making it easy to update your definitions and recompile the recognition package.

Defining natural language interpretations

To define natural language understanding for a recognition package, you:

- Define a fixed set of slots that correspond to the types of information supplied in utterances in the application’s domain
- Determine how each phrase in the grammar causes slots to be filled with specific values

In an automated banking application, for example, slots might include *command-type*, *amount*, *source-account*, and *destination-account*. Each slot has an allowed set of values—the *source-account* slot could be filled with the strings “checking,” “savings,” “money market,” while the *amount* slot is filled with a numeric value, and so on.

The utterance “Transfer five hundred dollars from savings to checking” would generate an interpretation structure with the following values:

command-type = transfer
source-account = savings
destination-account = checking
amount = 500

Natural language commands

You add natural language understanding capability to a grammar by defining the set of supported natural language slots, and then including special natural language *commands* in the grammar file that map values to those slots. A natural

language command is specified within curly braces ({ and }) and attaches to the grammar construct immediately preceding it. For example, in the next grammar, the command attaches to the word “from”:

```
Phrase ( from {...} checking )
```

while in the next one, the command attaches to the entire OR construction, [checking savings]:

```
Phrase ( from [ checking savings ] {...} )
```

There are two kinds of natural language commands:

- A *slot-filling* command that indicates a slot is to be filled with a particular value
- A *return* command that lets you associate a particular value with a grammar without actually filling a slot

Slot-filling commands

A slot-filling command has the form:

<Slot Value>

The `Phrase1` grammar below uses slot-filling commands to specify the string with which the *source-account* slot is filled:

```
Phrase1 (from [
    (?my checking ?account) {<source-account checking>}
    (?my savings ?account)  {<source-account savings>}
] )
```

Note: Notice that although “checking” can be expressed in four possible ways (“*checking*”, “*checking account*”, “*my checking account*”, and “*my checking*”), the *source-account* slot is always filled with the value “checking.”

In addition to string values, a slot can also be filled with an integer value, as in the next example:

```
Phrase2 ( transfer
    [ fifty    {<transfer-amt 50>}
      sixty    {<transfer-amt 60>}
    ]
    dollars )
```

If the natural language system can treat a value as an integer, it does so. You can force a value to be treated as a string by enclosing it in double quotes. You might want to do this in digit grammars, where treating the value as an integer causes leading zeros to be omitted. So instead of creating a digit grammar with constructs like:

```
( zero zero zero ) {<digits 000>}
( zero zero one ) {<digits 001>}
```

you should create a grammar like the following:

```
Digits [ ( zero zero zero ) {<digits "000">}
         ( zero zero one ) {<digits "001">}
         ( zero zero two ) {<digits "002">}
         . . .
       ]
```

You can also fill multiple slots within a single natural language command. The `Phrase3` grammar below fills two slots:

```
Phrase3 ( from savings to checking )
        {<source-account savings> <destination-account checking>}
```

Return commands and variables

Return commands let you associate values with a grammar, without filling any slots. The commands in the *Account* grammar below cause one of the string values “checking” or “savings” to be returned:

```
Account ( ?my
         [checking {return(checking)}]
         [savings {return(savings)}]
         ?account )
```

Return values do not appear in the interpretation structures produced by the natural language understanding system—you must use the return value to fill a slot. You do this by using *variables*. In your grammar specification, you assign the return value of a grammar to a variable, and then reference that variable to fill a slot in a slot-filling command.

Assigning variables

To assign the return value of a grammar to a variable, follow the grammar name with a colon and the variable name, as in:

```
( from Account:acct )
```

where *Account* is the grammar described above and *acct* is the variable. In this case, the variable *acct* is assigned one of the values “checking” or “savings.” You can then reference the value of a variable in a slot-filling command using the “\$” character:

```
Phrase4 ( from Account:acct ) {<source-account $acct>}
```

The variable *acct* is filled with the return value of the *Account* grammar, and that value is then referenced to fill the slot *source-account*.

Caution: If you omit the \$ before the variable name, the slot is filled with the string “acct” rather than the value of the variable.

Variables must be set before they can be referenced. In the following grammar the *source-account* slot is *not* filled:

```
Phrase5 ( from {<source-account $source>} Account:source )
```

Return values are most useful when a certain type of value is used to fill multiple slots. In the `Phrase6` grammar below the variable `acct` is used to specify how both the *source-account* and *destination-account* slots get filled:

```
Phrase6 [ ( from Account:acct ) {<source-account $acct>}  
          ( to Account:acct ) {<destination-account $acct>}  
          ]
```

Variable names can only contain alphanumeric characters and the following special characters: dash (-), underscore (_), single quote ('), and the "at" sign (@).

When returning a string value, enclosing the character sequence between double quotes is optional. Thus both expressions below are valid and equivalent:

```
{ return(cat) }  
{ return("cat") }
```

However, the use of double quotes is encouraged to clearly differentiate the expressions:

```
{ return(1) }  
{ return("1") }
```

where the first one returns an integer and the second a string.

The *\$return* variable

All rule references are automatically assigned to the magic variable *\$return*, unless an explicit variable assignment like the following is present in the file:

```
SomeGrammarName:aVariableName
```

\$return is a read-only variable that refers to the return value of the last rule reference in a given rule expansion.

Here are a couple of fragment specifications, in GSL and GrXML, illustrating the use of the *\$return* magic variable.

GSL

```
Main      (Command Number:n shares)  
           {<command $return> <num $n>}  
  
Command [Buy Sell] {return($return)}  
  
Buy [buy purchase] {return(buy)}  
Sell [sell]         {return(sell)}  
Number ...
```

GrXML	<pre> <rule id="Command"> <one-of tag="return(\$return)"> <item> <ruleref uri="#Buy"/> </item> <item> <ruleref uri="#Sell"/> </item> </one-of> </rule> </pre>
The assign command	<p>To copy the value of a <i>\$return</i> variable to a specific variable, Nuance introduced the assign command.</p> <p>This is a general purpose command that takes a destination variable as its first argument and any valid r-value—such as, a <i>\$variable</i> or a command—as its second argument, assigning the computed value of the second to the first.</p> <p>Here are a couple of fragment specifications, in GSL and GrXML, illustrating the use of the assign directive.</p>
GSL	<pre> Main (First:f Second:s) {assign(v add(\$f \$s)) return(mul(\$v 2))} </pre>
GrXML	<pre> ... <rule id="Main"> <item tag="&lt;command \$v&gt; &lt;num \$n&gt;"> <ruleref uri="#Command" tag="assign(v \$return)"/> <ruleref uri="#Number" tag="assign(n \$return)"/> </item> please </rule> ... </pre>
The <i>string</i> variable	<p>The variable mechanism also provides a way to use the actual string of spoken words to fill a slot instead of explicitly specifying the string value to return. You can do this using the special variable <i>string</i>. When you use the <i>string</i> variable, the slot is filled with the portion of the input utterance that matched that grammar. Consider the following grammar:</p> <pre> Day1 [sunday monday tuesday wednesday thursday friday saturday] {return(\$string)} </pre> <p>The special variable <i>string</i> allows you to write cleaner specifications than specifying each return value, as in:</p> <pre> Day2 [sunday {return(sunday)} monday {return(monday)} . . .] </pre>

How slots are filled

At runtime, the natural language system generates interpretations by matching the input utterance with a phrase (defined by a grammar) and executing any natural language commands attached to that grammar (subsequently referred to as the *matching* grammar).

Whenever a matching grammar has a construct with a natural language command attached to it, that command is executed and slot/value pairs are added to the interpretation. Commands attached to a grammar construct that is not matched by the utterance are not executed and, therefore, no slot/value pair is added to the interpretation.

For example, if the utterance “from savings” is processed against the following grammar:

```
Phrase [ ( from Account:acct ) {<source-account $acct>}
        ( to Account:acct ) {<destination-account $acct>}
        ]
```

the *source-account* slot is filled, but the *destination-account* slot is not.

When defining slot-filling commands, keep in mind the following points:

- *A variable must be assigned before it is referenced*

If a slot-filling command is executed but contains a variable that is not assigned, the slot is not filled. So in the next grammar example, the *source-account* slot is not filled when the input utterance is “to checking”:

```
Banking [(from Account:source)
         (to Account:dest)
        ] { <source-account $source>
          <destination-account $dest>}
```

- *Do not fill a slot with more than one value*

No interpretation is produced if a slot is filled with more than one value in the matching grammar. The runtime system generates a warning when it finds a construct in a matching grammar that violates this principle.

- *Commands have precedence over the unary operators*

This rule resolves how constructs preceded by a unary operator (*?*, ***, or *+*) and followed by a command are parsed.

Consider the following grammar:

```
ImperativeGo ( ?please {<polite yes>} go )
```

If “please” is *not* present in the utterance, then the slot *polite* is not filled. This is because the command is attached to the construction `please` and not the construction `?please`. If the command were attached to the construction

?please, the slot would be filled whether the input is “go” or “please go,” as in:

```
ImperativeGo ( (?please) {<polite yes>} go )
```

The slot definitions file

The slot definitions file is a plain text file that lists all the slot names used in your grammars, one per line. This file must contain at least one slot name.

Using the command-line utilities, you need to manually create the slot definitions file and name it *name.slot_definitions*, where *name* is the name of the recognition package that you are defining. It must reside in the same directory as your grammar files.

For the *Phrase* grammar in the preceding section, the slot definitions file would contain the following list:

```
source-account  
destination-account
```

As in a grammar file, any line in a slot definitions file that begins with a semicolon is considered a comment and therefore ignored by the grammar compiler.

Compiling a grammar with natural language support

The grammar compilation program *nuance-compile*, described in Chapter 6, automatically compiles natural language capabilities into the recognition package when the directory contains a *slot_definitions* file.

Nuance also provides the tool *nl-compile* that compiles only enough of the recognition package to support interpretation rather than providing full support for speech recognition. You can use this program to quickly compile grammars for natural language testing.

To use *nl-compile*, just specify the package name, for example:

```
> nl-compile %NUANCE%\sample-packages\banking1
```

Ambiguous grammars

A grammar is ambiguous if a sequence of words can produce multiple interpretations. For example, the following grammar:

```
.Command (call Name:nm) {<command call> <name $nm>}
```



```

Name [ [john (john smith)]      {return(john_smith)}
      [mary (mary jones)]      {return(mary_jones)}
      [john (john brown)]      {return(john_brown)}
      . . .
    ]

```

is ambiguous because the word sequence “call john” produces two interpretations:

```

{<command call> <name john_smith>}
{<command call> <name john_brown>}

```

When ambiguous sentences occur, the natural language system returns multiple interpretations, sorted by probability (if specified). See “Grammar probabilities” on page 20 for information on specifying probabilities in a grammar.

In some situations you may not be able to avoid ambiguity in your grammars—your application will have to resolve the ambiguity by asking the caller which interpretation was intended, by using:

- Program logic
- Subdialogs

To test your grammar for ambiguity from the command line, use the program *generate*, described in “Command-line tools for grammar testing” on page 114.

Advanced features

You can create many speech recognition applications that incorporate natural language understanding using only the features described in “Defining natural language interpretations” on page 52. However, some applications require more advanced features, such as:

- Filling a slot with a function of multiple values
- Filling a slot with a complex data type

Functions

In the grammars described so far, all slot and return values have been constants—either strings or integers. The natural language specification system also lets you fill a slot or specify a return value with a *non-constant* function. GSL supports a standard set of functions for manipulating integers and strings, and you can define your own functions.

Standard functions

GSL supports standard functions for filling a slot or specifying a return value, based on simple integer arithmetic or string concatenation. The following grammar demonstrates how you can specify a return value that is the sum of two other return values:

```
NonZeroDigit [ one      {return(1)}
               two      {return(2)}
               three    {return(3)}
               four     {return(4)}
               five     {return(5)}
               six      {return(6)}
               seven    {return(7)}
               eight    {return(8)}
               nine     {return(9)}
             ]

TwentyToNinety [ twenty   {return(20)}
                 thirty   {return(30)}
                 forty    {return(40)}
                 fifty    {return(50)}
                 sixty    {return(60)}
                 seventy  {return(70)}
                 eighty   {return(80)}
                 ninety   {return(90)}
               ]

TwoDigit [ ( TwentyToNinety:num1 NonZeroDigit:num2 )
           {return(add($num1 $num2))}
           TwentyToNinety:num1
           {return($num1)}
         ]
```

This grammar creates return values for numbers between 20 and 99 by adding the values returned by two subgrammars, *NonZeroDigit* and *TwentyToNinety*. For example, for the utterance “fifty nine,” the *TwentyToNinety* grammar returns 50, and the *NonZeroDigit* grammar returns 9. These are assigned to variables, which the grammar *TwoDigit* then adds to return the value 59.

The function mechanism can also handle variables that have not been set, as in the following alternative specification of the *TwoDigit* grammar:

```
TwoDigit2 ( TwentyToNinety:num1 ?NonZeroDigit:num2 )
           {return(add($num1 $num2))}
```

When the *TwoDigit* grammar processes an utterance such as “thirty,” the `add` function substitutes 0 for the variable *num2* and correctly returns the value 30.

The following table lists the available standard functions for integer and string manipulation, and describes how each function treats a non-set variable:

Table 4: standard integer and string functions

Function	Description	Unset variable
add	Returns the sum of two integers.	0
sub	Returns the result of subtracting the second integer from the first.	0
mul	Returns the product of two integers.	1
div	Returns the truncated integer result of dividing the first integer by the second (e.g., <code>div(9 5)</code> evaluates to 1).	0 if first argument; 1 if second argument
neg	Returns the negative counterpart of a positive integer argument, or the positive counterpart of a negative integer argument.	0
strcat	Returns the concatenation of two strings. This is a binary operator, but it accepts nested calls as in <code>strcat(\$a1 strcat(\$a2 \$a3))</code> , which effectively concatenates three strings.	"" (empty sting)

You can create more complex specifications by composing standard functions. The *ThreeDigit* grammar below builds on the previous *TwoDigit* grammar to cover three-digit numbers spoken in the form “two twenty three”:

```
ThreeDigit ( NonZeroDigit:num1 TwoDigit:num2 )
            {return(add(mul($num1 100) $num2))}
```

Note: See `%NUANCE%\data\lang\English.America\grammars\number.grammar` for an example of a grammar file that makes extensive use of the standard functions.

User-defined functions

The grammar syntax also supports the creation of your own functions for use in generating slot values. You define your functions in a file called *name.functions* (where *name* is the name of the package you are defining) that you include in the directory with your other grammar files, and it is processed during compilation by *nuance-compile* or *nl-compile*.

To define a function, specify the possible arguments it can take and the values it yields on those arguments. The function below, for instance, maps a day of the week to the following day of the week:

```
next_day ( <sunday    monday>
           <monday    tuesday>
           <tuesday    wednesday>
```

```

        <wednesday    thursday>
        <thursday     friday>
        <friday       saturday>
        <saturday     sunday>
    )

```

The `next_day` function takes a single string argument—you can also define functions that take more arguments. The following function sample maps a month and year into the number of days in that month in that year:

```

days_in_month ( <january 1999 31>
                 <january 2000 31>
                 <february 1999 28>
                 <february 2000 29>
                 etc.
    )

```

Each triplet, delimited by angle brackets in this function, specifies one possible group of arguments to the function and the result of the function when applied to those arguments. The function result is always the last value in the set, regardless of the number of arguments.

You use user-defined functions in the same way as standard functions. You might use the `next_day` function in a grammar as follows:

```

DaySpec ( the day after DayOfWeek:dow ) {return(next_day($dow))}

```

Complex values

Strings and integers are the two types of simple values supported by the natural language system. The system also supports two types of *complex* values:

- Structures, which contain a set of slot/value pairs
- Lists, which let you fill a slot with a set of values

A value used in a structure or list may be either of simple or complex type, and the values in a list or structure don't need to be of same kind; thus a list could consist of an integer, a string, and another list.

Structures

A structure lets you create values with multiple name/value pairs. Suppose you want to create a date grammar that simultaneously returns values representing the month, day, and year. This is most naturally represented as a set of slot/value pairs whose elements are:

```

<month may>, <day 15>, <year 1995>

```

The following simple date grammar illustrates how to return a structure:

```

Month [ january february march april
        may june july august september
        october november december ] {return($string)}
Day [ first      {return(1)}
     second     {return(2)}
     etc.
]
Year [ ( nineteen ninety five ) {return(1995)}
      etc.
]
Date ( Month:m Day:d ?Year:y )
      {return([<month $m> <day $d> <year $y>])}

```

Each slot within a structure is referred to as a *feature*, and has a value that is a string, integer, or another structure.

Referencing feature
values

You can create specifications that reference values of features within a structure through variable expressions.

Assume that a variable, *d*, is set to a structure that has *month*, *day*, and *year* features, and that you want to include the value of the *month* feature of *d* in a slot called *depart-month*. The following grammar illustrates how to do so:

```

Phrase (departing on Date:d) {<depart-month $d.month>}

```

The expression `$d.month` evaluates to the value of the *month* feature of the structure to which the variable *d* refers. (If *d* has no *month* feature, then the expression `$d.month` has no value, and the slot *depart-month* does not get filled.)

The *DateConstraint* grammar below illustrates a realistic use of the structure returned by the previous *Date* grammar:

```

DateConstraint [
    (departing on Date:dep-date)
    {<depart-month $dep-date.month>
    <depart-day $dep-date.day>
    <depart-year $dep-date.year>}
    (returning on Date:ret-date)
    {<return-month $ret-date.month>
    <return-day $ret-date.day>
    <return-year $ret-date.year>}
]

```

Although there is a single *Date* grammar, the structure returned is used to fill the departing and returning slots. This kind of specification would not be possible without using structures.

You can also fill slots with entire structures, rather than components of those structures. For example:

```
DateConstraint [
    (departing on Date:date) {<depart-date $date>}
    (returning on Date:date) {<return-date $date>}
]
```

The application handling the interpretation can access the value of each feature within a structure using the natural language API functions. Consult the online documentation for information on natural language APIs.

Nested structures

You can create specifications that include nested structures. For example, the structure:

```
[<time 1100>
  <date [<month may> <day 16> <year 1995>]>
]
```

includes a *date* structure as a feature value of a larger structure. This grammar illustrates how you can access the *month* feature of the *date* feature of that structure:

```
Phrase ( departing on TimeDate:time-date )
    {<depart-month $time-date.date.month>}
```

The expression `$time-date.date.month` evaluates to the value of the *month* feature of the *date* feature of the structure referred to by *time-date*. If the utterance is “on May 16th, 1999 at 6 AM” this feature has the value “may.” If *time-date* has no *date* feature, or if its *date* feature has no *month* feature, then the slot *depart-month* is not filled.

Several grammars in the Grammar Library make extensive use of structures. See the file *date.grammar*, under the *Date* folder, or the file *time.grammar* under the *Time* folder.

Note: See the grammar file *date.grammar* in `%NUANCE%\data\lang\English.America` for the complete contents of a grammar using nested structures.

Lists

Using lists, you can fill a slot or return value with a sequence of values. Here is a simple example of a grammar that uses a list:

```
DigitString ( one two three ) {<digit_slot (1 2 3)>}
```

When this grammar matches the input string “one two three,” it sets the slot *digit_slot* to a list containing the integers 1, 2, and 3. In general, you can fill a value with a list of any sequence of values by enclosing the sequence between parentheses.

You can construct lists explicitly as well as using more complex mechanisms such as functions and commands, described in the following sections.

The natural language system includes a number of built-in functions you can use to construct list values. The following example illustrates a grammar that matches any three-digit string:

```
Digit [ one {return(1)}
      two {return(2)}
      etc.
]
TwoDigit ( Digit:d1 Digit:d2 ) {return(($d1 $d2))}
ThreeDigit ( Digit:d TwoDigit:list )
           {return(insert-begin($list $d))}
```

The `insert-begin` function creates a new list by adding an item to the beginning of an existing list. If the utterance is now “one four five,” the *ThreeDigit* grammar returns a value by inserting “1” before the two-digit list “4 5” returned by the *TwoDigit* grammar.

The following table lists all the functions provided for working with lists:

Table 5: Built-in list functions

Function	Description
<code>insert-begin</code>	Returns a list with an item inserted at the beginning
<code>insert-end</code>	Returns a list with an item inserted at the end
<code>concat</code>	Returns the concatenation of two lists
<code>first</code>	Returns the first item in a list
<code>last</code>	Returns the last item in a list
<code>rest</code>	Returns a list with the first item removed

As with the integer and string functions, you can apply the list functions `insert-begin`, `insert-end`, and `concat` to arguments that are unset variables. In each case, the unset variable is treated as if it were an empty list.

Caution: It is an error to apply *first*, *last*, or *rest* to an unset variable, a non-list argument, or an empty list.

Several of the list functions can also be used as commands when constructing list values. Commands differ from functions in that:

- Commands appear within curly braces and perform an action. For example, the following command fills a slot:

```
{<month may>}
```

- Functions merely specify a value and must be used within a command to have any effect. For instance, the following function specifies a list value:

```
concat($list1 $list2)
```

To use the value returned by this function, it must be used within a command:

```
{return (concat($list1 $list2))}
```

List commands differ from `return` and slot-filling commands in that they actually change the value of a variable. The grammar below uses a list command to create a list value that matches a digit string of arbitrary length:

```
DigitString +( Digit:d {insert-end(list $d)} )
{return($list)}
```

As the grammar continues to match additional digits, it updates the value of the variable *d* until the full list is created and returned, representing the spoken string of digits in the appropriate order.

The *insert-end* command differs from the `insert-end` function in that the command *changes* the value of the list passed to it by inserting the value of the second argument at the end of the list. Like the `insert-end` function, the command treats an unset variable as an empty list.

This table lists the supported list commands:

Table 6: List commands

Command	Description
<i>insert-begin</i>	Modifies a list variable by inserting an item at the beginning
<i>insert-end</i>	Modifies a list variable by inserting an item at the end
<i>concat</i>	Modifies a list variable by concatenating it with the contents of another list

These commands modify the value of the first argument. So the command

```
insert-begin(list1 2)
```

modifies the value of the variable *list1*, adding the value 2 to the beginning of the list.

Note: Note that the first argument is a variable, *not* a reference to a variable. So *insert-end(list 2)* executes the command, whereas *insert-end(\$list 2)* is an instance of the function.

Semantic uniqueing

In general, before a set of answers is returned during N-best processing they undergo a process known as “semantic uniqueing.” This means that if multiple answers have the same natural language interpretation, only one of those answers is returned. If the recognition engine identifies two word strings with the same interpretation, then only the highest-scoring result is returned in the N-best list. As a consequence, you are guaranteed that all the answers in your N-best list differ semantically.

Semantic uniqueing is performed unless you have turned off interpretation by setting `rec.Interpret` to `FALSE`

There is one scenario you should be aware of. If you have no natural language within a particular top-level grammar, presumably you do not want semantic uniqueing to be performed with that grammar. However, if another top-level grammar in your package does include natural language definitions, semantic uniqueing is performed for all grammars (assuming `rec.Interpret` is `TRUE`). A future release may include the ability to explicitly disable semantic uniqueing.

Robust natural language engine

Nuance 8.0 offers a robust NL parsing capability that lets you write slot-filling grammars for the meaningful phrases *only*. The NL interpretation engine spots these meaningful phrases in the text output of the recognizer and fills the appropriate slots. This technology is best used with SLM grammars. For more information about robust NL interpretation, see “Robust natural language interpretation” on page 86.

Say Anything: Statistical language models and robust interpretation

5

The Say Anything™ feature, which includes Nuance's statistical language models (SLM) and robust natural language interpretation (robust NL) technologies, allows users to speak freely to an application and have their sentences interpreted by the NL engine without having to write complex grammar rules covering the entire sentence.

This chapter describes the SLM technology and explains how the robust interpretation of spoken sentences is performed. It then describes how to create an SLM grammar and how to use robust NL interpretation.

Overview

This section describes the motivation behind the Say Anything feature and explains how it can improve the accuracy of an application without the need to write complex GSL grammar rules.

The challenge

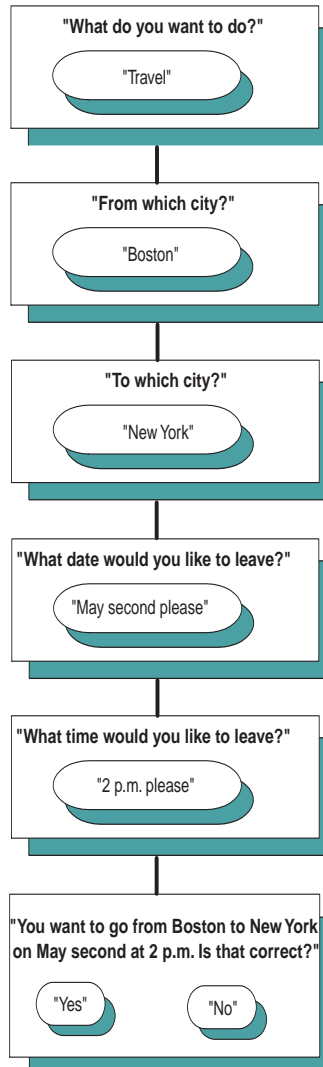
Consider an application that requires an open-ended prompt such as “What can I help you with?”. Not only can the user’s response be highly variable and hard to predict, but it can also contain *dysfluencies*—things such as re-starts, filled pauses (“um” and “uh”), and ungrammatical sentences. In addition, the grammar for this application may need to fill several NL slots from only one utterance.

The challenge is to write a GSL grammar that lets callers say any arbitrary phrase within the domain of the task, fills many slots, and still achieves high accuracy.

Coming up with such a set of grammar rules can be tedious. In most cases, the out-of-grammar rate obtained with hand-crafted GSL rules is very high and any attempt to add more GSL rules often leads to poor in-grammar recognition accuracy.

A simple solution

One attempt to solve this problem is to design a dialog in the form of nested sets of menus. For example, consider an application for airline ticket reservations. A dialog that used nested sets of menus would look like the following:



While this approach works well for certain applications (for example, a simple StockQuote application), it may not be suited for applications that must request a lot of information from callers. Users may find it difficult and frustrating to navigate through such complex dialogs.

An alternative approach is to construct a grammar such as:

```
.Sentence +Vocab
```

where *Vocab* is a subgrammar consisting of the list of all relevant words in the task domain:

```
Vocab [word1 word2 ... wordN]
```

For example, in a flight reservation application, the subgrammar would resemble:

```
Vocab [a Boston direct flight (san francisco) from i is next (new  
york) ninth on sixth the there to travel uh um want]
```

Suppose that the number of words in the vocabulary is *N*. Since every word in this vocabulary can follow every other word with the same likelihood, such an open-ended grammar usually performs poorly when *N* is large enough to cover most relevant words in the domain, typically several thousand.

The performance of the above grammar can be improved by assigning probabilities to the words in the vocabulary, based on the likelihood of their occurrence in a typical sentence. These probabilities can be evaluated by collecting a set of training examples and estimating the frequencies of words. In fact, Nuance provides a tool, *compute-grammar-probs*, that performs this task automatically.

However, a list of words with probabilities—commonly referred to as a *unigram*—is an overly simplistic model for many reasons, mostly because the probability of a word is *independent* of its position in a sentence. In this model, for example, a sentence such as *I want to travel to Boston next Monday* is considered as likely as any permutation of its words, such as *To want to Boston I travel next Monday*.

A better solution: The SLM and the robust interpretation approach

A language model that assigns probabilities to sequences of words is called a statistical language model (SLM). The unigram model described above is an example of a very simple SLM. Using *n-gram* SLMs provides a more interesting solution.

An *n-gram* SLM is one in which the probability of a word depends on the previous *N-1* words. *N* is called the *order* of the model. The unigram model

described above is an example of an n-gram model of order 1. In a unigram, the probability of a word is not affected by the context preceding that word. A 2-gram SLM—often called a *bigram*—is a model in which the probability of a word changes according to the word that precedes it. Hence, a bigram with a vocabulary of size V contains V^2 probabilities. Similarly, an SLM of order 3—often called a *trigram*—assigns V^3 probabilities.

Unlike a GSL grammar, an SLM grammar is not manually written but *trained* from a set of examples that models the user's speech. To train an SLM grammar, you pass a set of examples (and optionally a domain-specific vocabulary) to a Nuance utility, which estimates the model probabilities.

Since the probability assigned to a phrase depends on the context and is driven by data, an SLM provides a grammar in which more plausible phrases are given higher probabilities, a feature that the unigram model described above supports in a limited way only.

SLMs are useful for recognizing free-style speech, especially when the out-of-grammar rate is high. SLMs are not meant to replace GSL grammars, which are quite suitable when the application's prompts are sufficient to restrict the user's response. Since SLMs need a large set of examples to train, a data collection system or a pilot based on a GSL grammar may need to be developed to gather training examples.

Creating an SLM grammar

To create an SLM grammar, you:

- Create the training set
- Create a vocabulary file (optional)
- Determine the order of the model
- Train the SLM grammar
- Use the grammar in your application
- Measure the perplexity of a model

These steps are described below. The first sections describe the procedure in detail and provide conceptual information you need to perform the procedure. For a summary of the procedure, see “Summary of the procedure for creating an SLM grammar” on page 85.

Create the training set

The first step in creating an SLM grammar is to create the *training set*. A training set is a text file consisting of transcriptions of sample sentences, one per line, that users can say. Just as with any GSL grammar, the words cannot contain any uppercase letters, since *nuance-compile* assumes that uppercase words are rule names. The following example shows an excerpt of a training set for a flight reservation application:

```
um friday june the ninth uh from new york to austin
i want a flight from san francisco to los angeles on july sixth
is there a direct flight from boston to kalamazoo
i want to travel next monday
```

Your training set should not include punctuation and special characters. Also, all words and abbreviations should be written out the way they will be spoken. For example, *January 23* should be written *january twenty third* and *St. Patrick St.* should be *Saint Patrick Street*.

There is a strong correlation between the quality of an SLM and the quality of the training set. The best training sets are those collected from actual users of a system similar to what is to be deployed. You need several thousand examples at a minimum. For example, for vocabulary sizes up to 2,500 words, about 20,000 training examples are adequate to properly train an SLM.

Note that you may also want to create another file of sample sentences to use in your test set, to measure the quality of your model (as described in “Measure the perplexity of a model” on page 84).

Overview of class-based SLMs

The phrases in a training set can contain grammar rules; each grammar rule is called a *class*. An SLM that contains grammar rules is called a *class-based* SLM. For each class, you either write a GSL grammar rule that generates the phrases in the class or train a separate SLM. This procedure is described in the next section.

Class-based language models are useful when you have a limited amount of training data. In that case, certain important words or phrases may be rarely seen and therefore have poorly estimated statistics. A class lets you group similar words or phrases. For example, the class *City* in the above training set will probably be seen in your training data more often than individual city names, so its statistics will be better estimated and thus yield a better language model.

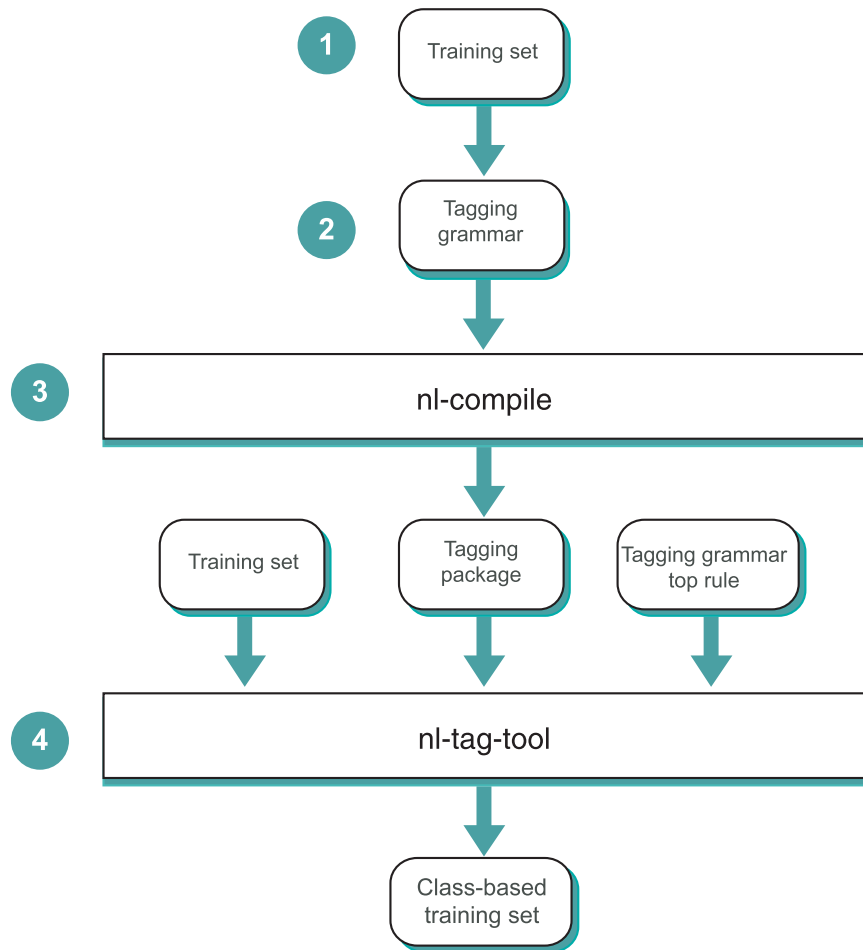
For example, consider the training set described above. You could generalize these examples by replacing the city names by the class name *City*. The *City* grammar rule would need to be defined as a separate GSL or SLM.

Class-based SLMs also let you leverage the existing standard grammars developed by Nuance for complex concepts such as time, date, and money amount. Furthermore, class-based language models let you dynamically modify the SLM vocabulary. For example, if the *City* class is a dynamic grammar, your application can dynamically change the entries in the *City* subgrammar at runtime, without retraining or recompiling the SLM.

Creating a class-based training set

To create a class-based SLM, you use the *nl-tag-tool* utility, which replaces phrases in a training set by the names of grammar rules. This process is known as *tagging*.

The *nl-tag-tool* utility works as follows. First, you look at your training set to determine the appropriate classes (step 1). You then write a grammar for these rules, called a *tagging grammar* (step 2). You compile this grammar into a tagging package (step 3). You then call *nl-tag-tool* with this tagging package, the training set, and the top rule of the tagging grammar. This utility converts the training set into a class-based one by processing the sentences and tagging them according to the grammar definitions. The following figure illustrates this process.



The *nl-tag-tool* utility takes sentences from standard input—you can enter sentences one at a time and see the resulting interpretation(s). You can also use *nl-tag-tool* in batch mode, as shown above, by creating a file with a list of sentences and using input redirection.

Note: *nl-tag-tool* also lets you filter which rules in a grammar should be tagged. See the *Nuance API Reference* for more information about this option and for a detailed description of the *nl-tag-tool* utility.

To create a class-based SLM:

- 1 Look at your training set to determine the appropriate class rules.

For example, in the training set described above, you could create the following classes: *City*, *DayOfWeek*, and *Month*.

- 2 Write a GSL grammar that defines each of these rules. This grammar, called the *tagging grammar*, will be used to relabel the training set.

For example, the following grammar, called *flightinfo.tagging.grammar*, defines the *City*, *DayOfWeek*, and *Month* rules:

```
.Tagging [ City DayOfWeek Month ]

City [ austin~0.15 boston~0.2 kalamazoo~0.05 (los angeles)~0.2
      (new york)~0.2 (san francisco)~0.2]

DayOfWeek [ sunday monday tuesday wednesday thursday friday
            saturday ]

Month [ january
        february
        march
        april
        may
        june
        july
        august
        september
        october
        november
        december
      ]
```

Note: If you have statistically valid information about the probabilities of the words in your application, Nuance recommends that you use these probabilities. For details on this topic, see “Probabilities in GSL” on page 20.

- 3 Create a tagging package using *nl-compile*.

For example, the following command compiles the *flightinfo.tagging.grammar* grammar:

```
> nuance-compile flightinfo.tagging.grammar English.America
```

Use a simple master package like *English.America* to produce the tagging package. This package will not be used for recognition.

- 4 Run *nl-tag-tool* to tag the training set, as follows:

```
> nl-tag-tool -package flightinfo.tagging -grammar .Tagging
-no_output < training.txt > training.tagged.txt
```

This command takes the current training set (for example, *training.txt*) and relabels it according to the package and grammar specified (for example, the *flightinfo.tagging* package and the *Tagging* grammar). The new training set is saved in file *training.relabeled.txt*.

For example, the relabeled training set would look as follows:

```
um DayOfWeek Month the ninth uh from City to City
i want a flight from City to City on Month sixth
is there a direct flight from City to City
i want to travel next DayOfWeek
```

Create a vocabulary file (optional)

You can also create a file that contains the *vocabulary* of your grammar—all the words to be covered by the SLM. This option is useful to constrain the vocabulary if the examples include typos, noise markers, or data from other applications containing non-relevant words that you do not want to include in your vocabulary.

This list of words is kept in a text file, one word per line. You need to split each utterance and phrase into unique words. For example, an utterance like “New York” needs to be split over two lines. Again, all the vocabulary words cannot contain any uppercase letters, since *nuance-compile* assumes that uppercase words are rule names.

The following is an excerpt of a vocabulary file for the flight reservation application example described above:

```
a
direct
flight
francisco
from
i
is
next
new
ninth
on
san
sixth
the
there
to
travel
uh
um
```

want
york

If you are using a class-based SLM, you need to determine the vocabulary from the relabeled training set. You also need to include the rule names. For example:

City
DayOfWeek
Month
a
direct
flight
from
i
is
next
ninth
on
sixth
the
there
to
travel
uh
um
want

Do not include words that are meaningless for the application and that occur once in the relabeled training set (for example, vulgar expressions).

While Nuance recommends that you use a vocabulary file, it is not required. If you don't specify a file, all unique words in the training set are included in the vocabulary.

Open and closed
vocabulary SLMs

There are two types of vocabularies: an open vocabulary SLM and a closed vocabulary SLM:

- An *open* vocabulary SLM is an SLM that contains a special class, called unknown (UNK), that lets you add new words to an SLM without retraining the model.
- A *closed* vocabulary SLM is an SLM to which you cannot add new words.

By default, an open vocabulary SLM is trained.

The UNK class has the following properties:

- If you have a vocabulary file, all the words in the training set that are not included in the vocabulary are tagged as belonging to the UNK class before the training process begins.

- If you do not have a vocabulary file or if all the words in the training set are in the vocabulary file, the UNK class is still generated but is assigned a small probability weight.

You need to define the UNK class in the GSL grammar that will include the SLM grammar. You can define it as an empty subgrammar or turn it into a background noise model for robustness, using items such as @reject@. You can also add new words to the UNK class. Nuance recommends that you use probabilities for all items defined or added to the UNK class.

Note that using the UNK class to add new words to an SLM should be a temporary solution only. Ideally, you should get new training samples that include the new words and then repeat the SLM training process to obtain proper statistics on the new words.

You can direct *train-slm* to train a closed vocabulary SLM—that is, with no UNK class—by passing the command-line option *-no-unknown*.

Determine the order of the model

The next step is to determine the order of the model. Is a trigram SLM better than a bigram for your application? It is correct to assume that the larger the N, the more powerful the n-gram model is since a larger context is used to assign a probability to a word. However, the number of probabilities needed in the model grows as a power of N and are therefore more difficult to estimate at large N values. Also, by increasing N when the training samples are limited, the model may experience over-training, that is, the model has memorized the training set and hence has lost its ability to model sentences not covered in the training set.

The optimum value of N is usually empirically determined by training a number of different n-gram SLMs and measuring their performance by running *batchrec* experiments on a set of example sentences (at least 500 sentences are needed to get a reasonable estimate of the performance). Your *batchrec* test set must not include sample sentences that were used to train your SLM. For more information on running *batchrec* to verify the performance of recognition packages, see Chapter 8, “Testing recognition performance.”

You can also use perplexity measurements, as described in “Measure the perplexity of a model” on page 84, to determine the value of N. However, *batchrec* is the preferred method.

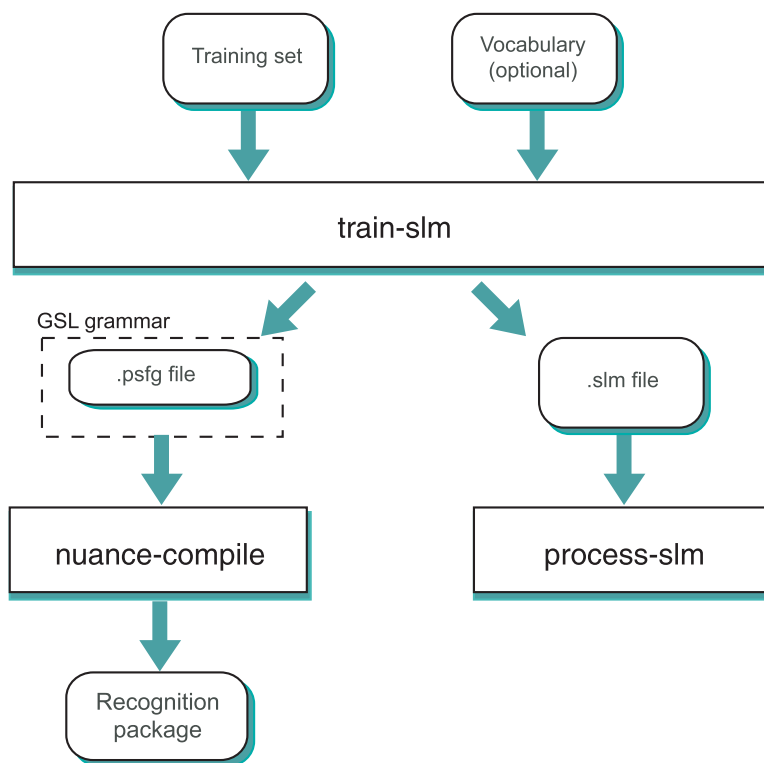
Train the SLM grammar

The next step is to train the SLM grammar. To train an SLM, you use the *train-slm* tool, which works as follows. It takes as input a training set and

optionally a vocabulary file, and it generates two output files that share the same base name but have different extensions:

- A *.psfg* file, which is a format that can be processed by the Nuance speech recognition engine. This is the file that you will be using in your Nuance application.
- An *.slm* file, which contains the basic probability values for the SLM. This file can be used for perplexity measurement, as described in “Measure the perplexity of a model” on page 84. Advanced users can also use this file for further manipulation. See Appendix B, “Advanced SLM features” for more information.

The following figure summarizes the training process.



Procedure

To train an SLM, enter the following command:

```
train-slm -corpus training-set-filename  
          -slm output-slm-basename  
          [-vocab vocab-filename] [-n N] [-no-unknown]
```

Where:

- *training-set-filename* is the name of the file containing the training set. See “Create the training set” on page 73 for more information on creating the training set.
- *output-slm-basename* is the base name of the output *.pfsg* and *.slm* files created by the trainer. The files are saved in the current directory.
- *vocab-filename* is the name of the optional vocabulary file. If you don’t specify this option, all unique words in the *training-set-filename* are included in the vocabulary. See “Create a vocabulary file (optional)” on page 77 for more information on creating the vocabulary.
- *N* is the order of the SLM to be generated. The default value is 3, a trigram SLM.
- `-no-unknown` trains a closed vocabulary SLM.

For example, to train a class-based SLM grammar for the flight reservation application with an open vocabulary SLM, enter the following command:

```
> train-slm -corpus training.relabeled.txt -vocab vocab -slm
flightinfo
```

The *train-slm* tool takes the class-based *training.relabeled.txt* training set and the *vocab* vocabulary file and generates the following two files:

- *flightinfo.slm*
- *flightinfo.pfsg*

The SLM will be a trigram, open vocabulary SLM.

Use the grammar in your application

You are now ready to use the SLM grammar in your application. To use an SLM grammar, you include the *.pfsg* file in a GSL grammar file using the following GSL rule:

Note: The `:slm` syntax replaces the Nuance 7.0 `:lm` syntax.

```
SLMSubGrammar:slm "output-slm-basename.pfsg"
```

Where:

- *SLMSubGrammar* is a subgrammar created from the *.pfsg* file. This grammar is treated like any other GSL subgrammar and can be included in other GSL grammars.
- *output-slm-basename.pfsg* is a *.pfsg* file created by *train-slm*.

Note: You cannot specify the *.slm* file instead of the *.pfsg* file.

For example, to use the SLM grammar created for the flight reservation application, you would create the following grammar, called *flightinfo.grammar*:

```
.Top SLMRule
SLMRule:slm "flightinfo.pfsg"
#include "flightinfo.tagging.grammar"
UNK [ @reject@ ]
```

You then compile the GSL grammar using *nuance-compile*, as follows:

```
> nuance-compile Your_GSL_Grammar Your_Master_Package -auto_pron
-dont_flatten
```

For example:

```
> nuance-compile flightinfo.grammar English.America.1 -auto_pron
-dont_flatten
```

Compilation of large vocabulary SLMs—especially class-based language models—can be slow and require large amounts of memory. Therefore, when compiling a grammar containing statistical models, Nuance recommends that you:

- Use the *nuance-compile* option *-dont_flatten*.
- If you find that compiling your SLM package takes too long or takes too much memory, use the *nuance-compile* option *-dont_optimize_graph* to disable the compilation optimization pass.

Once the speech recognition system starts processing a sentence using the *SLMRule* grammar, it can handle any phrase within the vocabulary of the model.

Using a class-based SLM

If you use a class-based SLM, you need to define the classes in the GSL grammar. For example, the sample grammar above includes the *flightinfo.tagging.grammar* grammar, which defines the *City*, *DayOfWeek*, and *Month*.

Using an open vocabulary SLM

If you are using an open vocabulary SLM, you need to define the UNK class in the GSL grammar. You can define it as an empty subgrammar or turn it into a background noise model for robustness—using items such as @reject@, as shown in the example above. You can also add new words to the UNK class. Remember to use probabilities for all items defined or added to the UNK class.

Final notes

Keep in mind that:

- When statically compiling a grammar containing an SLM rule, *nuance-compile* first looks in the current working directory for the .pfsg file. If it does not find the file, it then looks in directory %NUANCE%\data\lang\language\grammars.

- An SLM grammar must be a subgrammar in another top-level grammar; it cannot be a top-level grammar itself. However, the grammar can be as simple as:

```
.Top SLMRule
SLMRule:slm "flightinfo.pfsg"
```

Set appropriate parameters

Some recognition parameters significantly affect the recognition performance in SLM-based applications and should be set appropriately:

- 1 Set `rec.GrammarWeight` to a value of 9 or higher.

This parameter controls the weight of the probabilities in the grammar or statistical model, versus the acoustic knowledge.

Since all word strings in an SLM are valid utterances, it is important to help the recognizer prune unlikely word sequences based on the probabilities. The default value of 5 for this parameter is too low for SLM grammars and often the optimum value is around 9 or higher. Lower values are used when the SLM is not well trained.

- 2 Set `rec.Pruning` appropriately.

This parameter is highly correlated with the `rec.GrammarWeight` parameter. As `rec.GrammarWeight` increases, the recognizer runs faster and, typically, the pruning value needs to be increased to avoid extensive pruning.

- 3 Set `rec.pass1.gp.WTW` to 0 or -50.

This parameter is used to penalize word transition. It penalizes longer sentences in word loop type grammars—such as `+[word1 word2 ...]`.

In SLM-based grammars, the model probabilities already perform this task. Therefore, typically, this parameter should be set to 0 or -50 instead of its default value of -200.

- 4 Set `rec.PPR` to `TRUE`, except for small vocabulary cases, such as name spelling.
- 5 Turn N-best processing off (set `rec.DoNBest=FALSE`) or constrain the size of the N-best list (using `rec.NumNBest`) as much as possible, as N-best processing can be slow with large vocabulary SLMs.

Note: To get an N-best list with an SLM grammar, you must set `rec.Interpret=FALSE`.

The best way to determine the exact values of the parameters is to run multiple *batchrec* experiments with different settings to determine which values give the best performance. Keep in mind that large vocabulary SLMs typically run slower with extended master packages (packages with the suffix *.2*) than with standard master packages (packages with the suffix *.1*).

Measure the perplexity of a model

Perplexity is a measure of the quality of a language model. You can use this measure to tune the recognition accuracy of a model. The lower the perplexity, the better the model. In general, perplexity is correlated with recognition accuracy and a relative perplexity improvement (decrease) of 10% or more is indicative of better recognition accuracy.

Perplexity is interesting as a tuning tool because it is much faster to obtain than recognition accuracy, which requires running *batchrec*. To measure perplexity, you need a test set, which is a set of transcribed sentences from the application. This set must contain sentences that were not used to train the SLM.

To measure the perplexity of an SLM, use the *process-slm* tool, as follows:

```
process-slm
    slm_basename1 slm_basename2 ... slm_basenameN
    [-ppl-corpus test-set-filename]
```

Where:

- *slm_basename1 slm_basename2 ... slm_basenameN* are the filenames of the SLMs to measure. These are the *.slm* files obtained from *train_slm* without the *.slm* extension.
- *test-set-filename* is the name of the test set.

For example, suppose that two models were trained, a bigram model and a trigram model. You want to determine if the speed hit incurred with the trigram model is worth the performance gain. You can measure the perplexity of both models (*2gram.slm* and *3gram.slm*) on a test set called *test.txt* using the *process-slm* tool, as follows:

```
> process-slm 2gram 3gram -ppl-corpus test.txt
```

The output of the tool looks like:

```
Loading SLM file 2gram.slm
Done loading SLM file
Computing perplexity of corpus test.txt
8306 observation(s) logprob = -10285.300236 ppl = 17.310023
Perplexity = 17.310022
Loading SLM file 3gram.slm
```

```
Done loading SLM file
Computing perplexity of corpus test.txt
8306 observation(s) logprob = -10249.900334 ppl = 17.140981
Perplexity = 17.140982
```

In this case, the perplexity improvement going from a bigram to a trigram model is minimal: from 17.31 to 17.14. Since this is smaller than 10%, the trigram will probably not yield a significant recognition accuracy improvement.

Summary of the procedure for creating an SLM grammar

In summary, to create an SLM grammar:

- 1 Create a training set to train the language model. This text file, containing one phrase per line, should be representative of what a user of the target application would say. To create a class-based SLM:
 - a Look at your training set to determine the appropriate class rules.
 - b Write a GSL grammar that defines each of these rules. This grammar, called the *tagging grammar*, will be used to relabel the training set.
 - c Create a tagging package using *nl-compile*. For example:

```
> nuance-compile flightinfo.tagging.grammar English.America
```
 - d Run the *nl-tag-tool* utility to relabel the training set. For example:

```
% cat training.txt | nl-tag-tool -package flightinfo.tagging
-grammar .Tagging -no_output > training.relabeled.txt
```
- 2 Create a vocabulary for the model (optional). This vocabulary is used by *train-slm* to constrain the words incorporated in the model. The vocabulary should be a subset of the words used in the training set. If one is not supplied, *train-slm* will use all distinct words in the training set as the vocabulary.
- 3 Decide whether to use an open or a closed vocabulary.
- 4 Determine the order of the model.
- 5 Run *train-slm* to produce the statistical model in *.pfs* and *.slm* formats. For example:

```
> train-slm -corpus training.relabeled.txt -vocab vocab -slm
flightinfo
```
- 6 Include the produced *.pfs* file in a GSL grammar file.
- 7 Compile the GSL grammar using *nuance-compile*. For example:

```
> nuance-compile flightinfo.grammar English.America.1  
-auto_pron -dont_flatten
```

8 Set appropriate parameters.

9 Measure the perplexity of the SLM using *process-slm*. For example:

```
> process-slm 2gram 3gram -ppl-corpus test.txt
```

You may need to repeat the preceding steps several times, until your speed and accuracy requirements are met.

SLM and dynamic grammars

SLM grammars can be used with dynamic grammars just like conventional GSL grammars. The only restriction is that the GSL string to be compiled and used in a dynamic grammar cannot start with the *.pfsg* file name.

For example, the following GSL string is not valid for dynamic compilation:

```
"appslm.pfsg"
```

but the following format allows it to be compiled dynamically:

```
SlmRule  
SlmRule "appslm.pfsg"
```

For information about dynamic grammars, see “Dynamic grammars: Just-in-time grammars and external rule references” on page 23.

SLMs and just-in-time grammars

You can use an SLM in a just-in-time grammar. To do so, you include the *.pfsg* grammar in a static grammar, and then compile that grammar using the *nuance-compile-ngo* utility.

Robust natural language interpretation

SLMs are useful for recognizing free-style speech. But to understand the meaning of the spoken phrase, you still need to write grammars that fill slots with the appropriate values.

With the conventional NL parsing, writing these grammar rules can be a tedious task and defeats the advantages of using an SLM in the first place. To address this problem, Nuance offers a robust NL parsing capability that lets you write slot-filling grammars for the meaningful phrases *only*. The NL interpretation

engine spots these meaningful phrases in the text output of the recognizer and fills the appropriate slots.

Coupling the robust NL interpretation with the SLMs gives users more flexibility in what they can say and improves the performance of systems with high out-of-grammar rates. Specifically, this technology allows the NL engine to:

- Interpret more spontaneous speech effects such as hesitations, dysfluencies, and out-of-grammar sentences
- Maintain a high level of accuracy while using SLMs
- Handle flexible responses typical of mixed-initiative systems

Full and robust interpretation

The conventional NL interpretation is obtained as follows:

- Audio data is collected by an application
- The audio data is passed to the recognition engine
- The recognition engine generates text out of the audio data
- The text is passed to the NL engine
- The NL engine produces an interpretation result

In this conventional operating mode, called the *full* mode, the recognition engine and the NL engine are driven by the same GSL, which defines the valid phrases and how the slots are filled. However, you can use two different grammars: one driving the recognition phase and another driving the interpretation phase.

For example, the recognition can use an SLM grammar, allowing the application to recognize a large range of user's speech. The text output by the recognition engine is then processed by the NL engine running a GSL grammar that parses certain phrases only—the meaningful ones—and fills the appropriate slots.

While the conventional NL parser requires a *full parse* of the speaker's sentence by a top-level grammar rule—all the words in the sentence must be matched by a single grammar rule—the robust parser eliminates this requirement. If a full parse is not found, the robust NL parser attempts to fill as many slots as it can from partial parses using subsentence phrase fragments.

For example, consider the grammar for the flight reservation application above and the following user's sentence:

"I'd like to um I want to go to boston tomorrow."

The speech recognition engine, driven by the SLM, recognizes the sentence and sends the result to the NL engine, which tries to interpret the text.

In full mode, the text is not parsed by the NL engine since the sentence is not completely parsed by the grammar, and therefore no slots are filled. The sentence will be rejected. However, in robust mode, the *destination* and *day* slots will be filled by “boston” and “tomorrow” respectively.

Using robust interpretation

To use robust interpretation, you need to:

- Write the recognition and interpretation grammars
- Enable the robust interpretation mode

These tasks are described below.

Writing the recognition and interpretation grammars

The first step is to define both the recognition grammar and the interpretation grammar. To assign the interpretation grammar to the recognition grammar, you link them using the character = (equal) as follows:

```
Recognition_Grammar = Interpretation_Grammar
```

For example, let's extend the SLM-based travel application trained earlier with robust interpretation rules:

```
SLMRule:slm "flightinfo.pfsg" = InterpretationSubGrammar
```

```
InterpretationSubGrammar
[City:s {<destination $s>}
DayOfWeek [ sunday monday tuesday wednesday
            thursday friday saturday ] {<day $string>}
Month [ january
        february
        march
        april
        may
        june
        july
        august
        september
        october
        november
        december] {<month $string>}]
```

This grammar specifies that:

- The SLM in the file *flightinfo.pfsg* is used by the speech recognition engine to generate a text result.
- The subgrammar *InterpretationSubGrammar* is used by the NL engine to fill the *destination*, *day*, and *month* slots based on the recognition output.

Note: The use of the = (equal) character in GSL grammars is not restricted to SLMs. The following example is also valid:

```
SubGrammar +[word1 word2 ... wordN] = [today {<day today>}
                                         tomorrow {<day tomorrow>}
                                         ...]
```

Guidelines for writing grammar rules for robust interpretation

To take advantage of the robust parser mode, you must follow some guidelines when writing grammar rules.

The robust parser returns slots filled by grammar rules that parse phrase fragments. Therefore, to take advantage of robust parsing, Nuance recommends that you write grammar rules that fill slots with minimally valid phrase fragments.

The following example illustrates this recommendation. Consider the following GSL grammar:

```
.Sentence Date
Date:slm "date.pfsg" = (Month ?the DayOrd)
Month [january february march ... december] {<month $string>}
DayOrd [first {<ord 1>} second {<ord 2>}]
```

and the sentence:

“january uh first”

Here the sentence contains the filled pause “uh,” so in full operating mode the *Date* rule will not match that sentence, because the whole sentence was not matched.

However, the robust parser is able to match *january* using the *Month* rule—and therefore fills the slot *month* with the value *january*—and to match “first” using the *DayOrd* rule—and therefore fills the slot *ord* with the value “1.” This shows how robust parsing allows an application to recover from a filled pause and still return the correct NL interpretation.

Now consider what happens if the grammar is written this way:

```
Date:slm "date.pfsg" = (Month:m ?the DayOrd:d) {<month $m><ord $d>}
Month [january february march ... december] {return($string)}
DayOrd [first {return(1)} second {return(2)}]
```

In this case, with the same sentence, no NL interpretations are returned because the *Date* rule does not match the sentence and that is the *only* rule that fills slots.

By writing the grammar in this way, you are stating that the *Date* rule is the minimally valid—that is, shortest—phrase fragment that is acceptable for filling the slots. To be valid, the whole sentence must be matched. So, the rule of thumb for using the robust parsing is to fill at the smallest possible subgrammar.

Note: As the application developer, you must make sure that the words used by the interpretation grammar are included in the vocabulary of the SLM (recognition) grammar. The *nuance-compile* tool does not check for this type of error, which can cause a significant impact on the performance of the application. For example, consider the grammar `InterpretationSubGrammar` on page 88. You must make sure that words in the *City*, *DayOfWeek*, and *Month* subgrammars are included in the *flightinfo.pfsg* grammars, otherwise some of the words that fill the slots may never appear in the recognition result, therefore, the associated slot would never get the correct value.

Enabling robust interpretation

The next step is to direct the NL engine to spot the key phrases anywhere in the sentence by setting its interpretation mode to `robust` with the parameter `rec.InterpretationEngine`. This parameter can be set at runtime. Its default value is `full`, which directs the NL engine to use the normal interpretation mode.

Typically, you write your application so that it sets the value of `rec.InterpretationEngine` to `robust` when robust interpretation is to be performed and resets it back to `full` to resume the normal operation mode.

If a complete parse of a sentence is found, the robust NL will find the same result as the conventional full NL, but it may require extra computations. Nuance recommends you switch the `rec.InterpretationEngine` to `full` when the robust functionality is no longer required.

Ambiguity and robust interpretation

If you allow your application to handle free-style speech, it is likely that a given sentence will lead to multiple interpretations.

Consider, for example, the following grammar:

```
.Date [ (Month DayOrd)
      (DayOrd ?of Month)
      ]
Month [ january~0.2 {<month jan>}
      february~0.1 {<month feb>}
      ]
DayOrd [ first {<day 1>}
```



```
(twenty first) {<day 21>}  
]
```

And consider the following sentence:

“uh january twenty first of february”

No grammar rule parses this ambiguous sentence completely, so the robust parser finds multiple partial interpretations. The following interpretations are produced (higher ranked first):

- {<day 21> <month feb>}—filled by the phrase “twenty first of february”
- {<day 21> <month jan>}—filled by the phrase “january twenty first”
- {<day 1> <month feb>}—filled by the phrase “first of february”
- {<day 1> <month jan>}—filled by the phrases “january” and “first”

The first thing to notice is that the robust parser never fills a slot more than once and always uses non-overlapping phrases. For example, the phrases “january twenty first” and “first of february” cannot be used in the same interpretation, because they share the word “first” and they would fill the *month* slot with the conflicting values of *jan* and *feb*.

Note: The number of the interpretations returned by the NL engine is controlled by the runtime settable parameter `rec.MaxNumInterpretations`. By default, up to 10 interpretations can be returned.

To understand the order of interpretations, we need to know that the robust NL engine ranks interpretations according to the following rules:

- 1 Maximize the number of words covered in a sentence.
- 2 Minimize the number of grammar rules used.
- 3 Maximize the probability of the phrase fragments (using grammar weights).

Therefore, interpretations are first ordered by the number of sentence words covered by a phrase, then by the number of rules used, and finally by the probability of the phrases used in the parse.

In the example above, the first interpretation is filled by the phrase “twenty first of february” since it covers four words. The second and third interpretations, corresponding to “january twenty first” and “first of february”, cover three words. Note that “january twenty first” is ranked before “first of february” since, while they both cover three words with the single grammar rule `.Date` the grammar weight on “january” (0.2) is larger than that for “february” (0.1).

Finally, the last interpretation corresponds to the disjoint phrases “january” and “first”. This interpretation has the lowest rank since it only covers two words and uses two grammar rules (*Month* and *DayOrd*) to do so.

Statistical NL technology

Although the combination of the SLM and robust NL technologies offer a very flexible tool to develop free-style dialog applications, you still need to write grammar rules for the robust NL engine to interpret the spoken utterance. Nuance has developed a statistical NL technology that can eliminate the need to write NL grammar rules for call routing applications where a caller is routed to the correct destination based on the spoken utterance.

In these applications, there is only a single slot to be filled and that slot usually corresponds to the topic or the correct destination. To take advantage of this technology, a large number of example utterances must be collected and then transcribed and tagged with the correct destination route. Nuance provides an offline training tool that operates on the training data. Through statistical methods, this tool discovers the correlations between routing destinations and key phrases. The output of the training tool is called a *router package*. To take advantage of the router package, the parameter `rec.RouterDirectory` is set to point to the router package when the server is started. At runtime, the application can choose to use the router to interpret the sentence by setting the value of the `rec.InterpretationEngine` to `router`.

If you are interested in implementing applications based on this statistical NL technology, please contact Nuance professional services.

Compiling grammars

6

This chapter describes how to compile a grammar. It first presents the different ways of compiling a grammar. It then describes the master packages used in the process and other issues related to the generation of a recognition package.

Compiling a grammar


Nuance offers five ways to compile and store a grammar:

- Static compilation (*nuance-compile*)
- Dynamic compilation to a database
- Dynamic compilation to a file (*nuance-compile-ngo*)
- Just-in-time compilation
- Implicit compilation

Static compilation

Static compilation is performed using the utility *nuance-compile*. This utility takes a grammar file, a master package, and other resources and produces a recognition package containing compiled grammars, contexts, and parameter settings. This package is loaded when a recognition server is started.

This process is shown below.



To statically compile a grammar, enter the following command:

```
> nuance-compile package_name master_package [OPTIONS]
```

where

- *package_name* specifies the path and root name for the package files to compile. It can be a relative or an absolute path. The *nuance-compile* tool compiles the file *package_name.grammar* and reads additional pronunciations from *package_name.dictionary* if that file exists. Nuance's natural language understanding capability requires the file *package_name.slot_definitions* and optionally other files. The package directory created by *nuance-compile* is named *package_name* unless the *-o* option is used.
- *master_package* is the name of a master package. Master packages are provided by Nuance and found in %NUANCE%\data\lang. See “Choosing a master package” on page 101 for complete information on master packages.
- [OPTIONS] is any of the options described in “Summary of compiler options” on page 94.

For example, the following command tells the compiler to look for the grammar file *myapp.grammar* in the directory *c:\home\MyApp\grammars* and compile it using the American English master package:

```
> nuance-compile c:\home\MyApp\grammars\myapp English.America
```

The grammar file passed to *nuance-compile* must contain at least one top-level grammar once all the *#include* directives have been expanded.

The package directory created by *nuance-compile* contains all the files that the Nuance System needs for speech recognition. The actual contents of this directory vary depending on the options that you have specified. The file *out.compile* is always created and contains the command line used to generate the package.

Listing available options

Executing *nuance-compile* without arguments returns information about its usage and default settings. Further usage information can be obtained by specifying the command with *-options*.

Specifying a name for the recognition package

By default, *nuance-compile* generates a package directory named *package_name* in the current directory. To specify a package name different from the default one, use the *nuance-compile* option *-o output_package_name*.

Summary of compiler options

The table below introduces some of the options to *nuance-compile*. For a complete list of options, see the *Nuance API Reference*.

Table 7: *nuance-compile* options

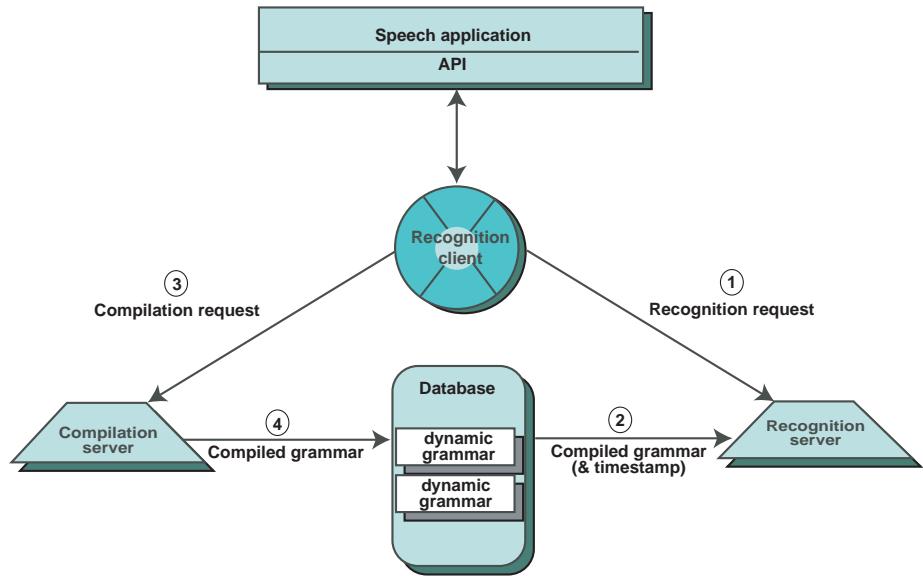
Option	Notes
<i>-merge_dictionary filename</i>	Merges standard dictionary entries with custom dictionary entries.
<i>-override_dictionary filename</i>	Replaces master dictionary entries with custom dictionary entries.
<i>-auto_pron</i>	Enables the generation of pronunciations for missing words. The generated pronunciations are written, by default, to file <i>package.autopron</i> .
<i>-nooverwrite</i>	Prevents deletion of previously compiled files.
<i>-write_auto_pron_output filename</i>	Writes the generated pronunciations to a specified file. It is meaningful only when used with the option <i>-auto_pron</i> .
<i>-enable_jit</i>	Enables the grammar for just-in-time compilation. If you do not specify this option, your application can perform recognition using the package's static grammars and contexts only.
<i>-optimize_graph</i>	Enables graph optimization pass (default). Leads to faster recognition but takes longer to compile.
<i>-dont_optimize_graph</i>	Disables graph optimizations. Speeds up compilation, but may slow down runtime performance.
<i>-dont_use_grammar_probs</i>	Tells the compiler to omit grammar probabilities from the compiled package. Useful to compare performance of package with and without probability specifications.
<i>-dont_flatten</i>	Expands each grammar only once during compilation. Produces smaller binary; limits the generation of compound-word pronunciations; must be used with recursive grammars.
<i>-o name</i>	Overwrites default package output name.
<i>-options</i>	Displays all compiler options available. Use with no other option.

Dynamic compilation in a database

Dynamic compilation in a database is initiated programmatically with a call to the Compilation Server method `NewDynamicGrammar`. The grammar passed to this function is compiled into a database that the recognition server accesses as needed.

For more details on dynamic grammars, see “Creating a dynamic grammar in a database” on page 23.

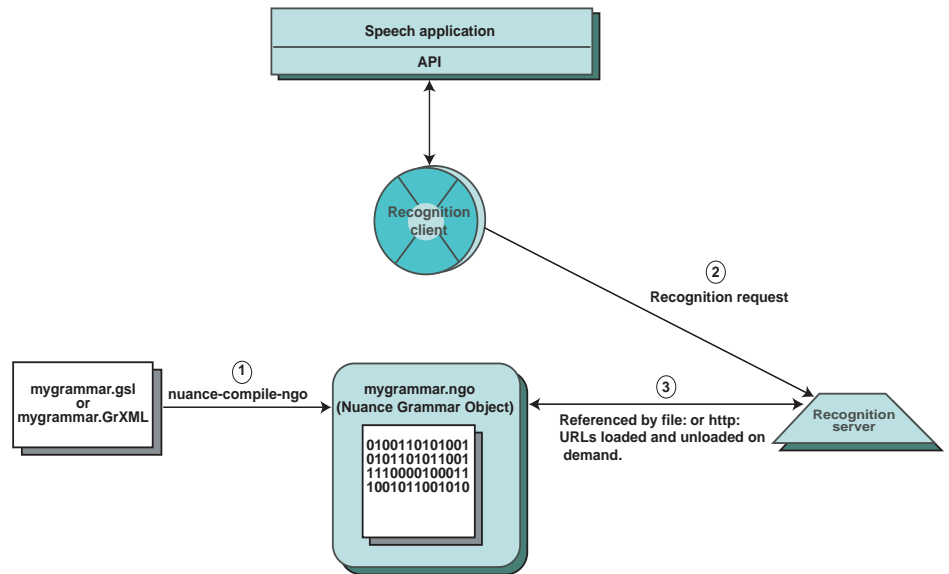
This process is shown below.



Dynamic compilation to a file

You can also compile a dynamic grammar with the utility *nuance-compile-ngo*. This utility can compile a GSL, GrXML (W3C format), or SLM grammar file into a *Nuance Grammar Object*—a binary grammar file with the *.ngo* extension. References to NGO grammars are done using the `file:` or `http:` protocols and are resolved by the recognition server, which loads and unloads these grammars on demand. Since an NGO file is precompiled, it is faster to load at runtime.

This process is shown below.



nuance-compile-ngo

To compile a grammar with *nuance-compile-ngo*, enter the following command:

```
> nuance-compile-ngo gram_file_name -package pack_name [OPTIONS]
```

Where:

- *gram_file_name* is the name of the file containing the grammar to compile. The extension of this file should match the type of grammar. For example, a GSL file should have the *.gsl* extension and a VoiceXML grammar file should have the *.grxml* extension. If another extension is specified, the header of the grammar is examined to determine its type. If that fails, the grammar is parsed as a GSL grammar. If that also fails, the grammar is not compiled.

gram_file_name is used as the base name of the output Nuance Grammar Object (*.ngo*) file created.

- *pack_name* is the name of just-in-time-enabled package that has been produced with *nuance-compile*, using the following options:
 - The master package that you want your dynamic grammars to use
 - The *-enable_jit* option

The grammar file that you pass to *nuance-compile* to produce the package *pack_name* is irrelevant.

- [OPTIONS] includes the following:

- `-do_crossword`, to generate compound word pronunciations for all words in the grammar.
- `-dont_add pauses`, to specify not to add pauses after words in the compile grammar. By default, pauses are added.
- `-dont_flatten`, to direct the compiler not to flatten the grammar during compilation. By default, grammars are flattened.
- `-dont_optimize_graph`, to direct the compiler not to perform any optimizations. By default, no optimization is performed.
- `-dont_use_grammar_probs`, to disregard any grammar probabilities specified in the grammar file. By default, probabilities are used.

Note that when referring to an NGO grammar in a JIT grammar, the NGO grammar and the JIT grammar must have been compiled using the same master package. Otherwise, the reference to the NGO grammar will be treated as a `fail: reference`.

Compilation of just-in-time grammars

The compilation of a just-in-time grammar is initiated when you pass GSL or GrXML content directly to the recognizer at runtime through an API call, as shown in the following example:

```
String gsl = "[hello (hi there)]";
jsc.playAndRecognize(gsl);
```

where `jsc` is a *`nuance.core.sc.NuanceSpeechChannel`*.

The recognition server then:

- Resolves all references found in the just-in-time grammar, according to the rules and conventions described in “Assembling grammars” on page 100
- Compiles the resulting grammar
- Uses the compiled binary for recognition and discards it

This section describes the compilation options available for just-in-time grammars.

To compile a recognition package to use with a just-in-time grammar, run *`nuance-compile`* with the *`-enable_jit`* option, as follows:

```
> nuance-compile my.grammar English.America.1 -enable_jit
```

In this example, the *`English.America.1`* package can be used to perform recognition of a just-in-time grammar. If you do not specify the *`-enable_jit`*

option, your application can only perform recognition using the package's static grammars and contexts.

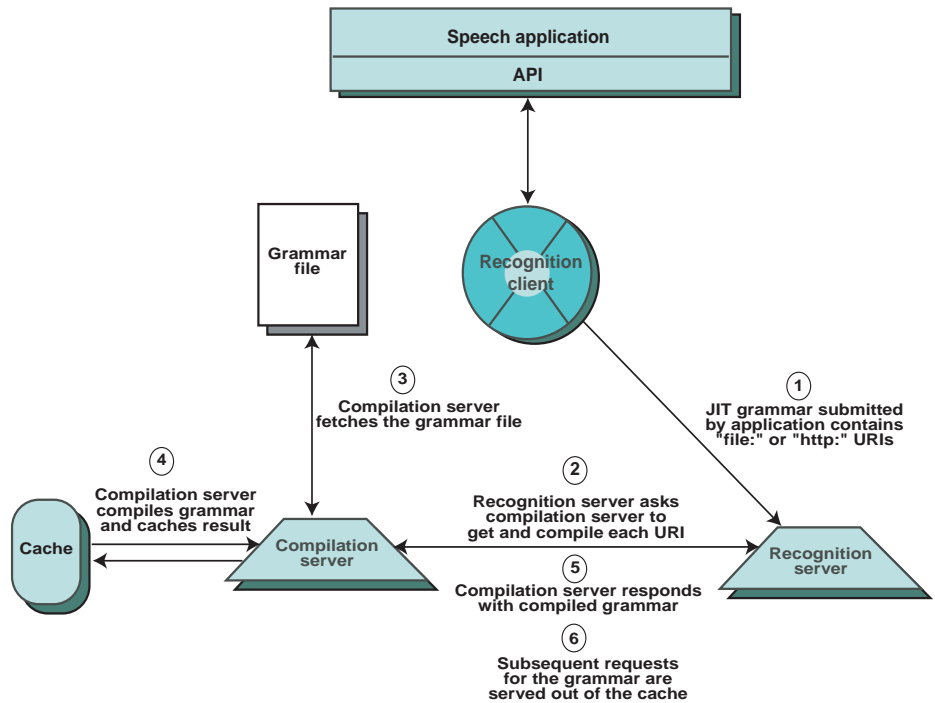
Implicit compilation

Implicit compilation is performed when a just-in-time grammar contains a `file:` or `http:` URI. Implicit compilation works as follows:

- 1** A just-in-time grammar submitted by the application contains a grammar reference in the form of a `file:` or `http:` URI.
- 2** The recognition server requests the compilation server to fetch and compile each of the references.
- 3** The compilation server gets the referenced grammars, compiles them, and caches the compilation result.
- 4** This result is passed to the recognition server, completing the request for compilation. Subsequent requests for grammars already compiled are served out of the cache.

The Recognition Server assembles the grammars that are used during recognition according to the rules and conventions described in “Assembling grammars” on page 100.

This process is below.



Assembling grammars

The recognition server resolves all references in a grammar specification. This process is called *assembling* a grammar.

The recognition server parses a grammar specified by an application and prepares to recognize with that grammar according to the following rules.

If the grammar has a name and it corresponds with a static grammar or context, the recognition server uses that grammar for recognition.

If the grammar is a just-in-time grammar, then the process is as follows:

- The recognition server requests the loading of all the components specified in references, until all references are resolved.
- The assembled grammar is compiled, used for recognition, and discarded after recognition ends, but the individual components are cached for later use.
- URI references specified with the `file:`, `http:`, and `dgdb:` protocols are always tested before each recognition.

Choosing a master package

The Nuance System provides multiple sets of acoustic-phonetic Hidden Markov Models also called *master packages*. These packages have been optimized for telephone-quality audio—that is, channels with 8-kHz/8-bit mulaw or alaw encoding, significant noise, narrow bandwidth (300 to 3300 Hz), and a variable linear frequency response. However, these packages also work well with most other microphones and audio channels.

Nuance provides master packages for multiple natural languages and packages that are optimized for use in several different types of environments. Master packages other than American English are *not* included in the “Custom” Nuance install option. These packages need to be specifically downloaded from the Nuance Developer Network website.

Master package naming convention

The default name format for a given master package is:

language[.dialect][.package_type]

where the square brackets indicate an optional string. Examples of this format are:

- *English.UK*
- *German*
- *French.Canada.1*

Note: If a master package is for a language for which only a single dialect is supported or for which the dialect name is equivalent to the language name, the dialect modifier is omitted. For example, the master package supporting Canadian French speakers is *French.Canada*, while the master package supporting European French speakers is simply *French*.

The specific package name format is:

language[.dialect].package_type.version.release

All languages have master packages with a type number equal to 1, and this version is called the *compact* package. A language may have master packages with types 1, 2 and/or 3:

- Packages with type number 2 are called *extended* packages. Extended packages perform additional recognition processing that may improve accuracy but requires additional memory.

- Packages with type number 3 are called *enhanced wireless* packages. They are designed for use in all environments including hands-free environments, on cellular phones, and in other environments with high levels of background noises. Enhanced wireless packages are capable of robust handling of any type of incoming speech, including normal land lines, cellular phones, and hands-free equipment. They require additional memory over the compact and extended packages.

The *version* number is the version number of the package; this number increases when there are significant changes to the package, for example when new acoustic models are available. The *release* number is the release number of the package; this number increases when there are minor changes to the package, for example minor changes to the dictionary.

Typically, you should leave out the *version* and *release* numbers when referencing a master package—the compiler defaults to using the latest version and release. If you need to use a specific version and release of a master package, you can specify them explicitly.

To determine how default names are mapped to specific master package names, use the utility *nuance-master-packages*.

Language-specific packages

A master package is language specific. The Nuance System provides multiple packages for some languages and dialects, such as American English. The package you use has an impact on your application's performance.

Language-specific packages provided with the current release include the following:

- Cantonese
- Czech
- Danish
- Dutch
- American English
- Australian-New Zealand English
- British English
- Singapore English
- South African English
- European French

- Canadian French
- German
- German (Switzerland)
- Greek
- Hebrew
- Italian
- Japanese
- Korean
- Mandarin (mainland China)
- Mandarin (Taiwan)
- Norwegian
- Brazilian Portuguese
- American Spanish
- European Spanish
- Swedish
- Turkish

Nuance releases new language-specific master packages as they become available. For a complete list of the master packages and languages supported, visit the NDN website at *extranet.nuance.com*.

See the Languages Documentation area on the NDN website for details on the language modules shipped with your version of the Nuance software, including the phoneme set, grammar writing conventions, and locale requirement.

Grammars for each language must be written in a specific locale, defined in the Languages Documentation. The following languages must be run under a specific locale:

- Hebrew: In Windows, Hebrew_Israel.1255; in Solaris, “he”
- Korean: In Windows, Korean or Korean_Korea.949; in Solaris, “ko”

Recommended usage

The command-line program *nuance-master-packages* lists the master packages installed on your system and the default package names mappings.

To determine which master package to use, keep in mind the following guidelines:

- When present, enhanced wireless master packages are generally the default package for a language. These master packages provide robust, accurate recognition for all applications. For example, *English.America* and *English.America.3* both map to the latest version of the *English.America.3* master package.
- If your system has memory constraints, try using the compact package, for example *English.America.1*, which uses less memory than *English.America.3*.

Mapping special characters

You may specify how certain characters in the English character set map to characters in the target language character set. Specifically, the characters that can be mapped are (space-separated listing):

() < > = + - * / , . ! ? # ^ ; : { } \$ " ' & @ | _ %

This specification is written in a plain text file, one line per mapping, according to the following syntax:

- The first (ASCII) character in a line is one of the listed above.
- The remaining non-blank characters, possibly multiple-byte characters, define what the first character is mapped to. The map may be one-to-many, in which case the list of equivalents, in the target language, are space-separated.
- If the first and the second characters are the same, the line is considered a comment, unless it is a special compiler directive (see paragraph at the end of this section).

This file should be installed in the directory *%NUANCE%\data\locale*. The name of this file follows the syntax that specifies the locale in your particular system. For example, the name of the French locale is *French-France@1252* on Windows and *fr_FR@ISO-8859-1* on Unix systems. Accordingly, the mapping file can be named:

- *French*, *French-France*, or *French-France@1252*, on the Windows platform.
- *fr*, *fr_FR*, or *fr_FR@1252*, on Unix systems.

For example, the locale equivalent mapping file *fr_FR*, may contain the following lines:

```
"«»
;; make «guillemots» equivalent to double quotes
```

The special comment-like line:

```
;;fast
```

is not treated as a comment but directs the compiler to speed up the processing of characters *only* when the locale name contains the suffix ISO-8859-1 (so-called Latin-1 languages) on Unix systems or the suffix 1252 on Windows platforms. Specifying such a directive in a mapping file on systems where the locale deviates from that suffix requirement leads to compilation errors.

Support for multiple languages

The Nuance System supports multilingual recognition. This feature lets you build applications that allow speakers to:

- Select their language of choice at the first prompt
- Use the correct pronunciation of foreign words in multi-lingual regions, for example, using the French pronunciation of French names within an English utterance
- Pronounce foreign words in international applications, for example, letting English or German-speaking callers use a Japanese dialing application and have names correctly recognized with the caller's native accent

To implement an application that uses multiple languages, you must use a master package that includes acoustic models for all the languages you want to support. Nuance can generate special mixed-language master packages on request. If you are interested in this feature, contact Nuance Technical Support with details about the languages you want to use, the languages you want automatic pronunciation for, the Nuance version and the operating system you are using, and the type of mixed-language support your application requires.

Grammar

Note that languages can be combined in the same grammar without special markings. For example:

```
Welcome ( *@hes@ [ hello bonjour ] )
```

Adding missing pronunciations

The *nuance-compile* program finds pronunciations for each word in your grammar using a set of standard dictionary files. However, you might need to create grammars containing words that do not have pronunciations in the standard dictionary—for example, unusual names or domain-specific names.

If *nuance-compile* finds any words in a grammar file that have no pronunciations in the standard dictionary, the program displays an error and creates a file in the current directory called *package_name.missing*. This file lists each word in the grammar with no known pronunciation.

You can provide pronunciations for these words in one of two ways:

- 1 By creating a custom dictionary for the package.
- 2 By using the *nuance-compile* automatic pronunciation generator option, *-auto_pron*.

Creating a dictionary file

To create a customized dictionary for a recognition package, create a file in the package directory called *package_name.dictionary*. The easiest way is to rename or copy the *package_name.missing* file generated by *nuance-compile* to *package_name.dictionary* and place it in the corresponding package directory. Then add a pronunciation for each word using the phonetic symbols listed in “The Nuance phoneme set” on page 133. Each line in the file should contain one pronunciation—consisting of the word followed by its phonetic sequence, for example:

```
telegraph t E l * g r a f
```

You can use the Nuance utility *pronounce* to see pronunciations for similar words. For example, if you want to create a pronunciation for the name “Vaughan” you could get the pronunciation for “gone” by running:

```
> pronounce English.America.2 gone
```

The program outputs the pronunciation:

```
gone g O n
```

From this you can determine that the pronunciation for “Vaughan” is “v O n”.

To provide multiple pronunciations for a word, include each pronunciation on a different line. For example:

```
process p r A s E s  
process p r o s E s
```

Multilingual dictionaries

A multilingual dictionary contains pronunciations in all the languages indicated and includes entries from the unilingual dictionaries. Words are not marked by language, while phonemes are tagged with a language code.

To create an application-specific dictionary, list the words and pronunciations as specified above, but tag the pronunciations with a language. Use phonemes from the languages you expect the callers accents to be in. For example, a person speaking French with a native American English accent requires that the French words in your application be pronounced using American English phonemes. The following dictionary entries illustrate how “hello” and “bonjour” can be handled in both the American English and French accents:


```
hello h.EA E.EA l.EA o.EA
hello E.FR l.FR O.FR
bonjour b.EA *.EA n.EA Z.EA U.EA r.EA
bonjour b.FR o~.FR Z.FR u.FR r.FR
```

Automatic pronunciation is supported in one language only.

Merging and overriding dictionaries

You can specify your custom dictionary to interact with the master dictionary in one of two ways:

- *Override*, where word pronunciations in your dictionary replace those of the standard dictionary
- *Merge*, where alternate word pronunciations are added to the standard dictionary

Command line directives

The *nuance-compile* option *-override_dictionary* lets you specify word pronunciations that replace any pronunciations existing in the standard Nuance dictionary. The following command line demonstrates how to use this feature:

```
> nuance-compile myGrammar English.America.2
      -override_dictionary c:\tmp\myDictionary.dictionary
```

The *nuance-compile* option *-merge_dictionary* allows you to add alternate pronunciations for words that already have a pronunciation in the standard Nuance dictionary. The following command line demonstrates how to use this feature:

```
> nuance-compile myGrammar English.America.2
      -merge_dictionary myAltProns.dictionary
```

The file specification for both of these options may be an absolute or a relative file path name.

Note: Detecting words that aren't in the master dictionary or in your custom dictionary helps you find possibly misspelled words in your grammars.

Automatic pronunciation generator

The automatic pronunciation generator is a compiler feature that tries to create pronunciations for missing words. The pronunciations that the compiler generates, if any, are written into a file called *package_name.autopron*, and automatically included in your compiled package. Then, if you want to improve your pronunciations, you can examine the generated pronunciations, accept or improve them, and include them in your dictionary. (If you do not include them, they are overwritten next time you compile your package.)

To use this feature, recompile your package using the option `-auto_pron` on the *nuance-compile* command line as illustrated in the command:

```
> nuance-compile myGrammar English.America.2 -auto_pron
```

You control the name of the output file where the missing pronunciations are written by using the option `-write_auto_pron_output`.

The following command line illustrates the use of this option. Here the compiler writes the automatically generated pronunciations to the file *myMissingWords* in the current directory:

```
> nuance-compile myGrammar English.America.2
    -auto_pron
    -write_auto_pron_output myMissingWords
```

The file specification for the `-write_auto_pron_output` option can be an absolute or relative file path name. If no pronunciations are generated automatically, any pre-existing *.autopron* file is removed.

Note: Using the option `-write_auto_pron_output` without the `-auto_pron` option has no effect.

Phrase pronunciations

You can often improve recognition accuracy by identifying compound words in your grammars. Also referred to as *crossword* or *multiword* pronunciations, compound-word pronunciations let you explicitly identify phrases that users tend to co-articulate. For example, users often say “give me” as “gimme,” or “I want to” as “I wanna.”

The grammar compiler *nuance-compile* and *nuance-compile-ngo* can use compound-word pronunciations for these types of phrases when they are specified in the dictionary.

Using *nuance-compile-ngo* is particularly efficient for dynamic grammars since it lets you generate crossword pronunciations for a specific subgrammar only, instead of generating them for the top-level grammar in a package and its entire contents.

For example, consider an application that uses a large top-level grammar, with a lot of fillers and alternate paths, but you only need to generate crossword pronunciations for one of the subgrammars. You can compile this subgrammar as a separate grammar document with *nuance-compile-ngo* and refer to the resulting NGO at runtime. Generating crossword pronunciations using *nuance-compile-ngo* provides more control and flexibility over the grammars that are compiled.

Note: If you want to use multiword pronunciations in your package, you must not use the option *-dont_flatten*.

For example, with the following grammar and dictionary specifications, multiple pronunciations are included for the phrase “what is”:

```
; grammar specification
Balance (what is my ?account balance

; dictionary specification
what      w ^ t
is         I z
(what is) w ^ z
(what is) w ^ t z
```

The last two lines in the previous dictionary specification provide compound-word pronunciations explicitly in a dictionary file. The format to use is similar to the single word pronunciation format, except that the compound word should be enclosed in parentheses:

```
(<compound_word_sequence>) <phoneme_sequence>
```

You can also use compound-word modeling to improve recognition accuracy for sequences of very short words (where each word contains one or two phonemes). Using compound words in these situations allows the Nuance System to assign the most detailed context-dependent acoustic models possible. This usually yields better performance than using a sequence of individual words, where the many word boundaries result in the assignment of less-detailed models. Some examples of this type of pronunciation are:

```
(i b m)          a j b i E m
(a t and t)      e t i * n t i
```

See “Tuning recognition performance” on page 111 for a description of the trade-off between accuracy and recognition performance, when a grammar is compiled with the crossword option.

Note: Previous versions of the Nuance system supported the ability to explicitly specify compound-word pronunciations by using “new” words, like “what_is”, in your grammar and explicitly adding pronunciations for the word “what_is” in the dictionary. This mechanism is still supported. The underscore mechanism, used to create a compound word out of a sequence of several words, does not permit pauses between the individual words in the sequence. Therefore, you should use the parenthesis mechanism when you want *both* compound *and* isolated pronunciations of compound words to be supported.

Optimizing recognition packages

This section presents ways in which you can optimize your recognition packages, including:

- Minimizing the size of your compiled package
- Speeding up the compilation process
- Improving recognition

Creating unflattened grammars

Large grammars—especially those that contain large subgrammars frequently used by other grammars—create packages that consume a large amount of memory. This is because, by default, grammars are fully expanded so that a subgrammar referenced multiple times in a grammar is included multiple times in the binary representation.

Packages compiled with the *-dont_flatten* option have each subgrammar included only once in the binary representation no matter how many times that subgrammar is referenced. This mechanism can cause a significant reduction in the size of the binary files for a package—however, it can also slightly slow recognition (on the order of 5%) for some applications.

Note: You must use the *-dont_flatten* option if any of the grammars in your application use recursion. On the other hand, this option disables the use of multiword pronunciations.

Speeding up compilation

By default, *nuance-compile* performs extensive graph optimization prior to the actual compilation. This results in an efficient, faster runtime package. However, it takes more time to complete the compilation process. To speed up the compilation, you can turn the optimization passes off by using the option *-dont_optimize_graph*.

Filler grammars

GSL gives you the ability to specify a subgrammar as *filler*. When you use filler grammars in your package, the recognition engine scores the confidence of a phrase recognition using only non-filler words.

This can sometimes give you more accurate results by preventing correct recognition of filler phrases from boosting the score of a recognition result into an acceptable range. For example, if the recognizer hypothesizes that the

out-of-grammar phrase “I want to leave next month” is actually the in-grammar phrase “I want to travel by car,” that result might be scored highly enough to be accepted because of the match of the “I want to” portion. If this portion is disregarded, the score of the remaining words should be low enough so that the utterance is correctly rejected.

To create a filler grammar, use the GSL keyword `filler`. For example, the following creates a filler grammar defining the phrases “I want to” and “I wish to:”

```
IWantTo:filler      [(i want to) (i wish to)]
```

You could use this grammar as follows:

```
.Main (?IWantTo travel by car)
```

You have access to filler/non-filler scoring differences at runtime, so that you can compare recognition performance. See the online documentation on the functions `RecResultOverallConfidenceWithoutFiller` and `RecResultOverallConfidence`.

Tuning recognition performance

The accuracy and speed of the Nuance recognizer are determined by several factors:

- The input speech: noise level, distortions, speech clarity, and so on
The cleaner the speech sample, the faster and more accurate the recognition.
- The dialog design
Recognition is hurt when the speaker is unsure of what to say. A good dialog design dispels any doubts on the part of the speaker as to what can or should be said in a dialog.
- The complexity of the grammar
A large or complex grammar usually results in a slower recognition system. Recognition may be less accurate because of the larger number of possible word sequences, or may be more accurate if the grammar better reflects actual input speech.
- The confusability of the grammar
Grammars that depend on the ability to distinguish between words that sound similar, such as “John Smith” and “Jon Smits,” typically result in higher error rates.

- The complexity of the acoustic models

More complex models are more accurate but can result in slower recognition.

- The use of the crossword option

Using the *nuance-compile* and *nuance-compile-ngo* option *-do_crossword* to compile a grammar will improve accuracy. However, it may result in slower recognition performance. Most gains are accomplished when used for small grammars with short words and many word boundaries, such as grammars for digits, alpha-digits and currency.

Nuance recommends that you turn the crossword option on for grammars where recognition performance is not an issue. This recommendation is true regardless of the language, but specially true for those languages with many short words, like Cantonese. The best way to determine if this is an appropriate option for your grammar is to try it out and compare results.

- The amount of search performed

The recognition system can be configured to perform a wider or narrower search. Less searching speeds up the system, but might increase the error rate if good theories are missed.

Sound grammar and application design are essential to good recognition performance. The performance of the Nuance recognizer can be further optimized by varying system configurations and by using more specialized grammar specification techniques. These include:

- Experimenting with different master packages
- Tuning parameter settings
- Creating more accurate pronunciations
- Using grammar probabilities

When developing and refining an application, Nuance recommends that you collect recordings of application usage and use them to measure the speed and accuracy of the recognizer offline under various configurations using the *batchrec* program. This tool is briefly described in “Command-line tools for grammar testing” on page 114.

Testing grammars

7

A grammar definition is “code” and, therefore, prone to have bugs just like any other piece of software or data. This chapter presents a methodology for testing a grammar that will help you find and fix bugs in your grammars.

The first section presents an overview of various grammar tests. The second section lists the command-line programs that Nuance provides to conduct basic testing. Typically, these programs are used in scripts that you create and run to test your grammars.

Grammar testing overview

A good plan for testing a grammar should include the following kind of tests:

- Coverage test

The goal of a coverage test is to verify that your grammar is able to parse a prescribed set of phrases. Usually you maintain a set of phrases—and, possibly, interpretations—that you require your grammar to parse after any change is introduced or any compilation parameter is modified.

- Interpretation test

This test verifies that slots are filled accurately and that all grammars return the correct values—that is, that your grammar delivers the expected natural language interpretation for a prescribed collection of phrases.

- Over-generation test

This test is intended to expose phrases that should *not* be parsed by your grammar. It helps you verify that your grammar will not accept unwanted sentences.

- Ambiguity test

This test exposes phrases parsed by your grammar that have multiple interpretations.

- **Pronunciation test**

This test is used to detect words with unknown pronunciations. It also exposes misspellings in your grammars.

Regression tests

Whenever you change a grammar, it is important to test the new version of the grammar to ensure that no errors have been introduced. If you run an old test set against a new grammar, you notice if any errors have been introduced to the previously working portions of the grammar. This is the basic paradigm of regression tests.

Whether you are testing a new grammar or doing a regression test, the steps are quite similar. The only difference is that in a regression test, you compare the results to a previous version, but when you test a new grammar, there is no previous version to compare against. The grammar tests described in previous sections may be used to test new grammars or as regression tests.

Command-line tools for grammar testing

The Nuance System toolkit includes several utility programs you can use to test your grammars. This section briefly describes the use and functionality of the following programs:

- *parse-tool*
- *generate*
- *nl-tool*
- *Xapp*
- *batchrec*

See the online documentation for complete reference information on these utilities.

parse-tool

This tool assists you in testing whether your grammar parses a given sentence correctly. It takes sentences from standard input and returns a value indicating whether the sentence was successfully parsed. Use the option

`-print_trees` to get a display of the grammar paths traversed to match the utterance. For example:

```
> parse-tool -package numbers -print-trees
```

produces the following output:

```
Your input ————— Ready
                                thirty two
                                Sentence: "thirty two"
                                .N0-99
                                DECADE
                                thirty
                                NZDIGIT
                                two
```

Note: *parse-tool* does not support dynamic grammars and just-in-time grammars.

generate

This tool outputs the sentences defined by paths in your grammars. The generation scheme could be exhaustive or random. This tool is typically used to test for overgeneration—that is, to detect nonsensical or inappropriate sentences that your grammar supports.

To run *generate*, specify the package and the name of the grammar (top-level or subgrammar) within the package that you want to analyze.

By default, *generate* continues randomly generating possible sentences until you stop the program. Use the *-num* option to specify the number of sentences to generate, as in the following command:

```
> generate -package numbers -grammar -N0-99 -num 5
```

Then the output might look like:

```
eighty one
seven one
fifty
fourteen
thirty two
```

If the program generates unwanted or unexpected sentences, your grammar is overgenerating.

To print out the interpretations for each generated sentence, specify the *-nl* option. For example:

```
> generate -package numbers -grammar -N0-99 -num 5 -nl
```

To test your grammars for ambiguity, use the option *-ambig*. This option causes *generate* to output only sentences defined by the grammar that have more than one interpretation. When using the *-ambig* option, the grammar you specify *must* be a top-level grammar, as in the following:

```
> generate -package myPackName -grammar .myTopLevelGrammar -ambig
```

If you are using the robust interpretation mode, note that *generate* only generates sentences using the interpretation portion of a grammar package, without using any grammars in the recognition portion of the package.

If you are using a grammar with the robust interpretation to parse just the meaningful phrase fragments, *generate* will only generate valid phrase fragments in that grammar, without any extra filler words that could be included in phrases parsed by an SML grammar.

Note: *generate* does not support dynamic grammars and just-in-time grammars.

nl-tool

This tool takes sentences from standard input and outputs the interpretations for them. It requires the specification of a package name and a top-level grammar within that package to use for recognition, as illustrated in the following command line:

```
> nl-tool -package myPackageName -grammar .Sentence
```

This tool also lets you perform batch mode analysis of interpretations using the arguments *-gen_ref* and *-compare_to*.

This utility works in both full and robust interpretation mode. Statically compiled packages can be used with *nl-tool* in both the full and robust parsing modes.

By default *nl-tool* operates in full parser mode. To enable the robust parsing mode, use the *-robust* command-line option.

Since the robust parser attempts to find the best interpretations from parts of a phrase, ambiguous interpretations will most likely be generated. For example, consider the grammar:

```
Month [january {<month jan>} february {<month feb>}]
```

and the utterance:

“january february”

This phrase would fail to be parsed in full parser mode. However, the robust parser generates two valid interpretations for it:

```
{<month jan>}  
{<month feb>}
```

See “Robust natural language interpretation” on page 86 for more information on robust parsing.

Note: The parameter `rec.MaxNumInterpretations`, used by *nl-tool*, can be used to limit the number of interpretations generated.

Xapp

Xapp is a graphical tool that performs recognition of audio utterances by supplying a transcription of each recognized phrase and the confidence score of the recognition.

This tool allows you to run a simple real-time test on a grammar by speaking utterances that you expect to be parsed by your grammar, and then checking the result.

Note: *Xapp* does not support dynamic grammars and just-in-time grammars.

To run *Xapp*:

- 1 Start a *recserver* process from a command-line window by typing a command like the following:

```
> recserver -package my-package lm.Addresses=my-server-machine
```

where *my-package* is a compiled package name, and *my-server-machine* is the name of the host machine where the Nuance Licence Manager (*nlm* process) is running.

The Nuance System also lets you load and unload recognition packages from a recognition server without restarting the server. For more details on this feature, see the *Nuance Application Developer's Guide*.

The recognition server is completely initialized when it displays the message:

```
Recserver ready to accept connections on port XXXX
```

- 2 In another window, invoke *Xapp* with the following command:

```
> Xapp -package my-package lm.Addresses=my-server-machine
```

where *my-package* and *my-server-machine* are identical to the strings used for the command in step 1. On the Windows platform, you can find the *Xapp* command in the Nuance start menu.

- 3 Select the grammar you want to test.

4 Cycle through the following steps until you are done:

- Click the Listen button.
- Say a phrase that you expect to be parsed by your grammar.
- Check out the result returned by *Xapp* in the “Recognized Speech” area of the window. No interpretation is displayed if your grammar does not have any natural language commands.

batchrec

An important tool that you can use to test your grammar is *batchrec*. This tool is rather sophisticated as it requires a set of pre-recorded audio data—usually gathered from field data or from a special data collection.

Basically, you run the spoken data through the *batchrec* and it evaluates recognition accuracy scores for each utterance in your data set. See Chapter 8 for more information on *batchrec*.

Testing recognition performance

8

This chapter describes how to test recognition performance of an application using the command-line program *batchrec*. The *batchrec* program performs recognition on a set of recorded audio files. If you also provide the correct results (either utterance transcriptions or natural language interpretations, or both) for the utterances in the audio files, *batchrec* scores the result for each file and prints cumulative accuracy statistics.

You should use *batchrec* to analyze and tune the performance of your live applications. Recognizing pre-recorded utterances instead of live speech lets you both process a lot of data quickly and process the same data multiple times. This enables you to:

- Establish a baseline test for recognizer accuracy
- Measure the speed of the recognizer
- Estimate performance on a live task, in advance
- Tune the recognizer configuration by varying parameter values while holding the speech data constant

You can record digital audio files for use with *batchrec* via a running application, by using the Nuance waveform editor *Xwavedit*, or by using third-party recording software. Nuance software records files in either the SPHERE-headered *.wav* format or the RIFF format commonly used on Windows. The *batchrec* tool can process either of these formats. You can also use the *wavconvert* tool to convert files between the *.wav*, RIFF, and *.au* formats.

Note: See the online documentation for information on the *wavconvert* and *Xwavedit* utilities.

Choosing *batchrec* test sets

The *batchrec* program provides you with a statistical measure of recognition accuracy for a set of audio files. To get useful results, you must test the right set of recorded utterances, and interpret the resulting statistics carefully.

The utterances in your *batchrec* test set should be as representative as possible of the utterances that your application will recognize when deployed.

Characteristics you should consider include:

- Gender and accent of speaker
- Content of speech
- Rate (speed) of speech
- Method of audio capture
- Ambient noise conditions

In addition, there is considerable variation among speakers, and even among different utterances by the same speaker. Therefore, for statistically reliable results you must provide a large sample set. To accurately measure speaker-independent recognition accuracy, your sample set should include at least 500 utterances spoken by at least 20 different people whose speech characteristics mirror the target speaker population.

Note: The recognition system uses several adaptive parameters to compensate for variation across audio channels and to estimate the background/channel noise in the signal. This means that each recognition result is somewhat dependent on the preceding utterances. At the start of the first utterance there is no information about the signal, so the system uses a set of default values. If the signal is very noisy or very clean these default values are not good, so the system makes an estimate of the signal noise and stores it for the next utterance to use. All subsequent utterances can improve this estimate, so the value for a particular utterance depends on all utterances before it. The effect on recognition is usually negligible, but if you change the order of your test utterances, the recognition score and perhaps the result may change. The recognizer always starts in the same state, so if you do not change the order of the utterances, your test set always produces the same results.

Using *batchrec*

To invoke *batchrec*, you must specify, at a minimum:

- The recognition package

- A file containing a list of the files to be processed

For example:

```
> batchrec -package package_name -testset testset_list
```

batchrec also supports a number of additional arguments, described in “Additional options” on page 124. The following section describes the format of the file list passed to *batchrec*.

Creating the testset file

The file you specify with the *-testset* option lists the digital audio files to be recognized, one audio file per line. Each audio file is processed as a single utterance for output and scoring purposes, regardless of its content.

The testset file also identifies the grammar to use to recognize each audio file. Many recognition packages contain several top-level grammars, to be used for different recognition tasks. Use the **Grammar* keyword to identify the grammar to use to recognize particular files. Typically, you create command blocks, listing the set of files to recognize with each grammar. For example:

```
*Grammar .YesNo
/home/usr1/waveforms/yes.wav
/home/usr1/waveforms/no.wav
*Grammar .StockQuote
/home/usr1/waveforms/stocks001.wav
/home/usr1/waveforms/stocks002.wav
/home/usr1/waveforms/stocks003.wav
```

With this test set, the files *yes.wav* and *no.wav* are recognized using the *.YesNo* grammar, while the files *stocks001.wav*, *stocks002.wav*, and *stocks003.wav* are recognized using the *.StockQuote* grammar.

Note: If the recognition package you specify on the *batchrec* command line contains only one top-level grammar, this grammar is used by default. However, it is good practice to always specify the grammar on the first line of your testset file. You can also set a default grammar for a *batchrec* run with the option *-grammar*.

Additional commands

In addition to the **Grammar* command, *batchrec* supports a number of additional commands you can include in your testset file:

```
*Echo text
```

Lets you specify text to be written out to the screen during processing.

```
*Exit
```

Lets you explicitly exit the testset at a given point.

`*SetParam param_name=param_value`

Sets the specified Nuance parameter to the given value. The parameter must be runtime-settable. For information about Nuance parameters see the online documentation or the *Nuance Application Developer's Guide*.

`*GetParam param_name`

Prints out the current value of the specified parameter.

`*NewAudioChannel`

Indicates that the following audio files were generated on a different audio channel than the previous file. Clears inserted dynamic grammars with `CALL` persistence.

You can also include commands for working with dynamic grammars and for performing speaker verification. Dynamic grammar commands are described in the next section. For information on *batchrec* speaker verification commands, see the *Nuance Verifier Developer's Guide*.

Dynamic grammar commands

You can use *batchrec* to evaluate the performance of applications that include dynamic grammar functionality. This includes both recognizing against dynamic grammars, and using the enrollment facility to add a new voice-trained phrase to a dynamic grammar (this requires that you provide a transcriptions file when you start *batchrec*). The related *batchrec* commands are described here. See “Dynamic grammars: Just-in-time grammars and external rule references” on page 23 for information on the Nuance System's dynamic grammar functionality. In the references below, the database handle (such as *db_handle*) must be a positive integer. Zero is not allowed as a database handle value.

Note: A compilation server is required for the `*NewDynamicGrammar` functions. Make sure that you start *batchrec* with the `-rcengine` option (see “`-rcengine`” on page 126), that you start a resource manager and a compilation server, and that you set the `rm.Addresses` parameter correctly. See the *Nuance Application Developer's Guide* for more information.

`*OpenDynamicGrammarDatabase db_handle DBDescriptor_arguments`

Creates a connection to the dynamic grammar database containing the dynamic grammars you want to use for testing. This must be a database where `-dbclass` is equal to `dgdb`. See the *Nuance Application Developer's Guide* for information on the arguments required for a given database provider.

`*CloseDynamicGrammarDatabase db_handle`

Closes the connection to the specified database.

`*NewDynamicGrammarEmpty db_handle db_key ok_to_overwrite`

Creates a new, empty record in a dynamic grammar database.

`*NewDynamicGrammarFromGSL db_handle db_key gsl_file
ok_to_overwrite`

Creates a new record in a dynamic grammar database containing the given grammar content. This can be a GSL extended right-hand side or a complete GSL or GrXML grammar document.

`*NewDynamicGrammarFromPhraseList db_handle db_key pl_file
ok_to_overwrite`

Creates a new record in a dynamic grammar database containing the specified set of phrases. The phrase list file (*pl_file*) contains one phrase per line, in the format:

phrase_id phrase_text nl_command probability

For example:

`user_112 "john doe" "{<user john_smith>}" 1`

Note that the natural language command *must* be enclosed in double quotes.

`*DeleteDynamicGrammar db_handle db_key`

Removes a record from a dynamic grammar database.

`*CopyDynamicGrammar source_db_handle source_db_key
target_db_handle target_db_key ok_to_overwrite`

Creates a new record in a dynamic grammar database, containing the contents of an existing record.

`*AddPhraseToDynamicGrammar db_handle db_key p_id p_text p_nl
*AddPhraseList db_handle db_key pl_file`

Adds a single phrase or a set of phrases to a dynamic grammar.

`*RemovePhrase db_handle db_key p_id`

Removes a phrase from a dynamic grammar.

`*QueryDynamicGrammarExists db_handle db_key`

Prints out whether the given database contains a dynamic grammar at a given key.

`*QueryDynamicGrammarContents db_handle db_key`

`*QueryDynamicGrammarContentsWithID db_handle db_key phrase_id`

Prints out the contents of a specified dynamic grammar database record, or of the phrases in that record with a given ID.

`*InsertDynamicGrammar db_handle db_key label {CALL | PERMANENT}`

Inserts a dynamic grammar from a database into a static grammar at the given label, with the specified persistence (CALL or PERMANENT). CALL persistence means that the insertion remains until next `NewAudioChannel` line.

`*EnrollNewPhrase grammar db_handle db_key {NEEDED | ALL}`

Begins a dynamic grammar enrollment session. The enrolled phrase is added to the specified dynamic grammar. Note that to be able to perform enrollment, you must provide either a transcription file and a natural language transcription file when you start *batchrec*, using the `-transcriptions` and `-nl_transcriptions` options. The format of these enrollment transcription files is the same as the format of the transcription files used for recognition, except for the use of noise markers, which are not allowed in enrollment transcription files. See “Additional options” on page 124 for a description.

For the final argument, specify `NEEDED` to have the new phrase added to the grammar as soon as enough consistent enrollments are found. Specify `ALL` if you want all enrollments to be used as listed in the testset file, before the phrase is added to the grammar.

See “Example: Using *batchrec* to perform enrollment” on page 130 for an example of performing enrollment with *batchrec*.

`*EnrollCommitPhrase`

Commits an updated dynamic grammar to the database.

`*EnrollAbortPhrase`

Ends an enrollment session without updating the dynamic grammar.

`*ModifyPhrase db_handle db_key old_phrase_id new_phrase_id
new_phrase_nl`

Modifies a voice-enrolled phrase in a dynamic grammar.

`*CompileAndInsertDynamicGrammar gsl_file label persistence {CALL |
PERMANENT}`

Compiles a dynamic grammar from a GSL file and inserts it into a static grammar at the given label, with the specified persistence (CALL or PERMANENT).

Additional options

batchrec supports a number of options:

`-transcriptions transcriptions_file`

If you supply a file containing a transcription for the utterance in each audio file in the test set, *batchrec* prints accuracy statistics, both per sentence and cumulative. Each line in the transcriptions file contains a file name followed by a transcription string. The file name can be as short as the base name of the audio file, or as long as the complete path of the audio file. Supply as much of the path as is necessary to uniquely identify all files. If all files in the test set are found in one directory and have unique base names, only the base name needs to be supplied. Everything after the file name is the transcription of the audio file. For example:

```
datadir1/file3.wav phoenix arizona
datadir1/file1.wav chicago illinois
datadir2/file3.wav boston massachusetts
```

Transcriptions can be listed in any order. Words in the transcription must be spelled exactly as they appear in the grammar, including letter case. If digits are spelled out in the grammar (“four” instead of “4”), they must be spelled out in the transcriptions file, too.

`-nl_transcriptions nl_transcriptions_file`

If your grammar includes natural language interpretations, you can supply a file with transcriptions of the correct natural language interpretation for each audio file, and *batchrec* prints accuracy statistics for the natural language interpretations returned by the recognizer. Each line in the file contains a file name and then a natural language interpretation. Everything after the file name is the natural language interpretation for that audio file. You can also generate the NL transcriptions file from a standard transcriptions file using the *generate-nlref* tool.

You can list natural language transcriptions in any order. A natural language interpretation is a set of named slots, with a value for each. Each slot/value pair is enclosed in angle brackets. Inside the angle brackets, the first word is the slot name and everything after the first space is the slot value—so slot names are limited to a single word, while slot values may be one or more words. If the slot name or value includes an angle bracket, the entire slot name or value must be enclosed in double quotes. Slot and value names must be spelled exactly as they appear in the grammar. For more information on how to specify a natural language component in your grammar and how to use the natural language recognition result, see “Defining natural language interpretations” on page 52. An example transcriptions file is:

```
datadir1/file3.wav <city phoenix> <state arizona>
datadir1/file1.wav <city chicago> <state illinois>
datadir2/file3.wav <city boston> <state massachusetts>
```

Note: A common scoring problem that arises with digit grammars is that the natural language system produces a string value, while the transcription gives an integer value. To force a sequence of digits to be treated as a string, enclose it in double quotes. The transcription file would look like this:

```
datadir1/file1.wav <digits "123456">
```

rather than like this:

```
datadir1/file1.wav <digits 123456>
```

If N-best recognition is turned on, *batchrec* also prints N-best accuracy statistics, both for recognition accuracy and natural language accuracy. The N-best accuracy statistics are just like the regular statistics except that they measure the accuracy you would achieve if the system could automatically choose the best item out of the N-best list. For more information on N-best recognition, see the *Nuance Application Developer's Guide*.

`-grammar grammar_name`

Sets a default grammar to be used if none is specified with a `*Grammar` command.

`-vrs`

This option causes *batchrec* to run in a client/server environment by causing the program to connect to a running recognition server or resource manager instead of performing recognition in the same process.

`-rcengine`

This option causes *batchrec* to run in a client/server environment, connecting to a recognition server process as a client using the RCEngine interface.

`-endpoint | -dont_endpoint`

The `-endpoint` and `-dont_endpoint` options specify whether or not endpointing is applied to the speech samples before recognition.

The default is `-dont_endpoint` because it is assumed that the recorded utterances have already been endpointed. If `-endpoint` is specified, then leading and trailing silence is stripped out of the speech sample before it is handed to the recognizer. Nuance recommends that you do not endpoint speech files that were recorded, and therefore already endpointed, by a Nuance application. Results differ, depending on how this option is set.

`-print_confidence_scores`

The recognition accuracy at various confidence thresholds is output for both transcriptions and natural language transcriptions. To obtain maximum output set the parameter `rec.ConfidenceRejectionThreshold` to 0 or -1 on *batchrec*'s command line.

`-print_word_confidence_scores`

This option causes *batchrec* to include the confidence score for each word in each utterance in its output.

`-debug_level 0-5`

`-debug_level_during_init 0-5`

You can use *-debug_level* and *-debug_level_during_init* to control how much information is printed about the processing. They take a value from 0 to 5; 0 suppresses most output, while 5 prints the largest amount of information. The default level for both of the settings is 2.

`-detailed_nl_scoring`

Keeps track of the insertion, deletion, and substitution errors at the semantic level, when your application uses the robust NL parsing. For details on robust NL parsing, see “Robust natural language interpretation” on page 86.

Setting Nuance parameters

You can also specify Nuance parameters on the *batchrec* command line. For example:

```
> batchrec -package package_directory
           -testset testset_list_file
           -transcriptions transcriptions_file
           -endpoint rec.DoNBest=TRUE
```

causes *batchrec* to perform N-best processing.

See the online documentation or the *Nuance Application Developer's Guide* for more information on Nuance parameters.

Output from *batchrec*

The following shows sample *batchrec* output for a single audio file, where *batchrec* was run with both a transcriptions and an NL transcriptions file:

```
File 72:      /home/tests/myapp/testset/numfood171.wav
Grammar:     .Sentence
Transcription: buy him seven fish now
NL Transcript: <get-how buy> <for him> <count seven> <food fish>
              <when now>

Result #0:   buy him seven fish now (conf: 66, NL conf: 63)
NL Res.#0:   <when "now"> <count "seven"> <for "him"> <get-how "buy">
              <food "fish">

Total Audio: 2.63 sec
Utt. Times:  0.37 secs 0.141xRT (0.37 usr 000 sys 0.141xcpuRT) 100%cpu
Avg. Time:   0.34 secs 0.132xRT (0.33 usr 0.00 sys 0.130xcpuRT) 99%cpu
```

```
Rec Errors:      0 ins, 0 del, 0 sub = 0.00% of 5 words.
Rec Total:       0 ins, 0 del, 2 sub = 0.56% of 360 words (2.78% of 72 files).
NL Status:       correct
NL Total:        0 rejects, 2 incorrect = 2.78% error on 72 files
```

This output includes the following information:

File 72: /home/tests/myapp/testset/numfood171.wav

Provides the name and location of the audio file being analyzed. In this case, this is the seventy-second file recognized by this *batchrec* run.

Grammar: .Sentence

Names the recognition grammar used to recognize this file.

Transcription: buy him seven fish now

Prints the transcription of the utterance in the audio file, as provided in the file specified with the *-transcriptions* option.

NL Transcript: <get-how buy> <for him> <count seven> <food fish>
<when now>

Prints the natural language transcription for the utterance, as provided in the file specified with the *-nl_transcriptions* option.

Result #0: buy him seven fish now (conf: 66, NL conf: 63)

Prints the recognition result returned by the recognizer and the confidence scores of that result.

NL Res.#0: <when "now"> <count "seven"> <for "him"> <get-how "buy">
<food "fish">

Prints the natural language interpretation generated from the recognition result.

Note: In the previous example, only one result is returned. If N-best processing is enabled (by setting the parameter *rec.DoNBest* to *TRUE*), multiple recognition and natural language results may be returned.

Total Audio: 2.63 sec

Provides the length of the recording.

Utt. Times: 0.37 secs 0.141xRT (0.37 usr 000 sys 0.141xcpuRT)
100%cpu

Provides the wall clock and CPU times used to process the file, followed by the percentage of CPU time dedicated to recognition processing during this period. This file, for example, took 0.37 seconds to process, and the recognition process got 100% of the CPU's time during processing.

```
Avg. Time: 0.34 secs 0.132xRT (0.33 usr 0.00 sys 0.130xcpuRT)
99%cpu
```

Provides the average wall clock and CPU times used to process files to this point, and the percentage of CPU time dedicated to recognition processing. The `xcpuRT` value is typically the most useful measure of speed.

```
Rec Errors: 0 ins, 0 del, 0 sub = 0.00% of 5 words
```

Prints error statistics for this file, including number of incorrect words inserted by the recognizer, number of words deleted, and number of words misrecognized. This line also includes the word count and the per-word error rate.

```
Rec Total: 0 ins, 0 del, 2 sub = 0.56% of 360 words (2.78% of 72 files)
```

Prints error statistics for all files processed so far, including per-word and per-file error rates. In this example, for the 72 files processed so far, there have been no insertions, no deletions, and two substitutions (misrecognitions). 0.56% of words have been misrecognized and 2.78% of the files have yielded a recognition error.

```
NL Status: correct
```

Indicates whether the natural language interpretation was correct, incorrect, or rejected.

```
NL Total: 0 rejects, 2 incorrect = 2.78% error on 72 files
```

Provides natural language error statistics for all files processed so far, including the number of rejects, the number of incorrect interpretations produced, and the per-file error rate. In this example, 2.78% of the files have yielded a recognition error. This error rate is the same as the recognition error rate for this example. This can occur when the grammar produces a one-to-one word to slot mapping.

Using batchrec for SLM grammars

The *batchrec* program determines the status of an utterance to be correct, incorrect, or reject by comparing NL interpretations to NL transcriptions.

Switching the recognition engine to the robust NL parsing mode may cause some NL results to be partially correct, in the sense that some slots contain the right values but others are either not filled or incorrectly filled.

Specify the *batchrec* command-line option *-detailed_nl_scoring* to keep track of the insertion, deletion, and substitution errors at the semantic level.

Example: Using *batchrec* to perform enrollment

To perform enrollment with *batchrec*, you need to create the following files:

- A testset file, which includes commands to enroll a set of utterances
- A transcription file, which lists all the utterances used for enrollment and associates a phrase ID with each utterance
- A natural language interpretation file, which specifies the natural language interpretation for each utterance

You then run *batchrec* with the following options:

```
> batchrec -package package_directory
           -testset testset_list_file
           -transcriptions transcriptions_file
           -nl_transcriptions nltranscriptions_file
```

For example, the following sample testset file opens a file system database and creates a record for a new speaker (*acct_1234*). It then performs enrollment of two difference phrases using the *.Enroll* enrollment grammar; three utterances are used to train each phrase. Finally, it closes the dynamic grammar database.

```
*OpenDynamicGrammarDatabase 1 -dbprovider fs -dbroot . -dbname
enroll_db -dbclass dgdb
*NewDynamicGrammarEmpty 1 acct_1234 1
*EnrollNewPhrase .Enroll 1 acct_1234 needed
utt01.wav
utt02.wav
utt03.wav
*EnrollCommitPhrase
*EnrollNewPhrase .Enroll 1 acct_1234 needed
utt04.wav
utt05.wav
utt06.wav
*EnrollCommitPhrase
*CloseDynamicGrammarDatabase 1
```

The phrase ID associated with each utterance is specified in the transcription file as follows:

```
utt01.wav word_1
utt02.wav word_1
utt03.wav word_1
utt04.wav word_2
utt05.wav word_2
utt06.wav word_2
```


The natural language interpretation for each phrase is specified in the natural language interpretation file as follows:

```
utt01.wav word_1 <name "word_1"> <phone_number "1234567">
utt02.wav word_1 <name "word_1"> <phone_number "1234567">
utt03.wav word_1 <name "word_1"> <phone_number "1234567">
utt04.wav word_2 <name "word_2"> <phone_number "7654321">
utt05.wav word_2 <name "word_2"> <phone_number "7654321">
utt06.wav word_2 <name "word_2"> <phone_number "7654321">
```


Creating application-specific dictionaries

9

The Nuance System represents each word in a recognition vocabulary as a sequence of phonetic symbols. These sequences provide the mapping between a word in the vocabulary and its pronunciation, and reside in dictionary files. The Nuance System dictionary provides phonetic pronunciations for more than 100,000 English words. Ideally, all the words you need for your application will be part of the Nuance dictionary file. However, some applications may require words or alternative pronunciations, such as proper names or special acronyms, that the Nuance System dictionary does not contain. In these situations, you can define additional word pronunciations in a separate, auxiliary dictionary file.

When you compile your grammars, the compilation process will tell you whether your grammars contain any words that are not part of the main Nuance System dictionary by generating an error message. To resolve this type of compilation error, you provide pronunciations for those words not listed in the Nuance dictionary, as described in Chapter 6, “Adding missing pronunciations” on page 105, either by using the automatic pronunciation generator or by providing a dictionary file with the missing pronunciations.

Note: Remember that the automatic pronunciation generator provides a fast, easy mechanism for generating pronunciations during application development. However, pronunciations that are created manually are typically more accurate, assuming they are carefully created and tested.

The Nuance phoneme set

The Nuance System uses a set of symbols to represent *phonemes*. Roughly speaking, there is one phoneme corresponding to each spoken sound in a given language. For spoken English, for example, there are more than 40 different phonemes. Because letters can be pronounced differently in different words and because different letters can represent the same sound, there is not a one-to-one

correspondence between letters and phonemes. For example, the [k] sound and corresponding /k/ phoneme occur for the *c* in *cat*, the *k* in *keep*, the *ck* in *tick*, the *ch* in *echo*, and the *q* in *Iraq*. Similarly, the letter *c* can be pronounced using the /s/ phoneme as well as the /k/ phoneme (as in *circle*). Therefore, a mapping between the pronunciations of a word (specified as a phoneme sequence) and the spelling is necessary input to the recognizer.

The Nuance phoneme set is specified using the Computer Phonetic Alphabet (CPA). The CPA provides a system for easily expressing the phonemes in notation defined by the IPA (International Phonetic Alphabet) using a standard keyboard. The following table lists the phonemes used by the Nuance System to express pronunciations for American English:

Table 8: Mappings for American English

Phonetic Category	Phoneme	Example	Phoneme	Example
Vowels	<i>i</i>	fleet	<i>u</i>	blue
	<i>I</i>	dimple	<i>U</i>	book
	<i>e</i>	date	<i>o</i>	show
	<i>E</i>	bet	<i>O</i>	caught ^a
	<i>a</i>	cat	<i>A</i>	father, cot ^a
	<i>aj</i>	side	<i>aw</i>	couch
	<i>Oj</i>	toy	<i>*r</i>	bird
	<i>^</i>	cut	<i>*</i>	alive
Stops	<i>p</i>	put	<i>b</i>	ball
	<i>t</i>	take	<i>d</i>	dice
	<i>k</i>	catch	<i>g</i>	gate
Flap	<i>!</i>	butter, hidden		
Nasals	<i>m</i>	mile	<i>g~</i>	running
	<i>n</i>	nap		
Fricatives	<i>f</i>	friend	<i>v</i>	voice
	<i>T</i>	path	<i>D</i>	them
	<i>s</i>	sit	<i>z</i>	zebra

Table 8: Mappings for American English

Phonetic Category	Phoneme	Example	Phoneme	Example
	S	shield	Z	vision
	<i>h</i>	have		
Affricates	<i>tS</i>	church	<i>dZ</i>	judge
Approximants	<i>j</i>	yes	<i>w</i>	win, which
	<i>r</i>	row	<i>l</i>	lame
Others (reserved for Compiler)	-	(silence)		
	<i>inh</i>	(inhale)	<i>exh</i>	(exhale)
	<i>clk</i>	(click)	<i>rej</i>	(other)

- a. Many American dialects do not distinguish between /A/ and /O/, and the distribution of the two sounds among specific words often varies. When adding words to your dictionary, look at similarly spelled words using the utility program pronounce to get an idea of which symbol is appropriate; perhaps providing a variant with each will be useful.

It is important that you be very careful when using the phoneme set, because the symbols representing phonemes do not consistently correspond to the letters in the spelling of a word. Because one phoneme can be represented by several different letters or letter combinations, those same letters can represent other phonemes in other words. The following table shows some typical mappings from English letters to phonemes.

Table 9: Mapping English-language letters

Letter	Phonemes	Example
<i>a</i>	<i>e</i>	able
	<i>A</i>	father
	<i>a</i>	apple, pan
	<i>O</i>	ball, caught, pawn
	<i>*</i>	alive, zebra (unstressed vowels)
<i>b</i>	<i>b</i>	ball
<i>c</i>	<i>k</i>	cat, echo

Table 9: Mapping English-language letters

Letter	Phonemes	Example
	<i>tS</i>	child
<i>d</i>	<i>d</i>	dogs
	<i>!</i>	bridle
	<i>t</i>	ask ed , kiss ed
<i>e</i>	<i>i</i>	even
	<i>E</i>	e cho
	<i>u</i>	new
	<i>*</i>	ag ent
<i>f</i>	<i>f</i>	f ather
<i>g</i>	<i>g</i>	g olf
	<i>dZ</i>	gi ant
<i>h</i>	<i>h</i>	h otel
<i>i</i>	<i>aj</i>	ivory, time
	<i>l</i>	I ndia
	<i>i</i>	Lisa, sh i eld
	<i>*</i>	sanity, aph i d
<i>j</i>	<i>dZ</i>	j ump
<i>k</i>	<i>k</i>	k ey
<i>l</i>	<i>l</i>	l aw
	<i>* l</i>	ab le
<i>m</i>	<i>m</i>	M ike, thumb
	<i>* m</i>	Ad am , wis d om
<i>n</i>	<i>n</i>	n ew, k not
	<i>* n</i>	wid e n, nation
<i>n+g</i>	<i>g~</i>	r ing
<i>o</i>	<i>o</i>	o pen, co a t, sh o w

Table 9: Mapping English-language letters

Letter	Phonemes	Example
	A	olive
	aw	cow
	O	bought
	u	shoe, loop
	U	book, would
	^	company
	*	carrot
	Oj	toy
<i>p</i>	<i>p</i>	pull
	<i>f</i>	aphid
<i>q</i>	<i>k w</i>	quick
	<i>k</i>	Iraq
<i>r</i>	<i>r</i>	ring
	<i>*r</i>	bird, rider, grammar
<i>s</i>	<i>s</i>	soup
	S	shell
	z	dogs, wisdom
	Z	vision
<i>t</i>	<i>t</i>	time
	!	butter, sanity
	D	them
	T	thin
	tS	question
	S	nation
<i>u</i>	^	under, putt
	U	put, would

Table 9: Mapping English-language letters

Letter	Phonemes	Example
	<i>u</i>	lunar, rude
	<i>j u</i>	usage, confuse
	<i>*</i>	focus
<i>v</i>	<i>v</i>	very
<i>w</i>	<i>w</i>	wisdom, which
<i>x</i>	<i>z</i>	xenophobia
	<i>k s</i>	axis
<i>y</i>	<i>j</i>	yes
	<i>i</i>	very
	<i>aj</i>	cry
<i>z</i>	<i>z</i>	zebra

As a general guideline for generating dictionary entries, ignore the spelling and focus on how the word sounds. Break it down into its component phonemes, sound by sound. Write the codes for those phonemes in the correct sequence. To verify, try to reproduce the sound of the word by reading the symbols in sequence. For example, in the word *backwards*, the sounds (represented by spelling, not phonemes) are as follows:

b, a, ck, w, ar, d, s

The letters are sufficient to tell you what the component sounds are, only because you can still see what the word is, and you know how to pronounce it. To tell the recognizer that the sounds are spelled b-a-c-k-w-a-r-d-s, you must be more precise about the exact sounds:

- In the Nuance phoneme set, the [b] sound is straightforward, and the phoneme symbol for it is /b/.
- The next sound is a little harder. The letter *a* can represent many different sounds, even in this one word. The English letter table shows that the second sound in *backwards* is the same as the first sound in *apple*, and is represented by the symbol/a/.
- The third sound is represented in the spelling by two letters, *ck*, but the symbol for this phoneme is just /k/.

- The fourth phoneme is represented by the symbol /w/.
- The fifth phoneme is tricky. It is a syllabic *r*, meaning that it has blended with the vowel to make up the root of the syllable, and is therefore one sound. The symbol for this phoneme is /*r/. (In a Boston accent, the *r* part of the phoneme is dropped, and the phoneme is /*/.)
- The symbol for the sixth phone is /d/, and the final phone, which is spelled with an *s*, is actually represented by the symbol /z/, which is what it sounds like.

The dictionary entry for this word, therefore, might look like this:

```
backwards  b a k w *r d z
backwards  b a k w * d z
```

where the second pronunciation indicates a Boston accent.

Note: Be especially careful with words borrowed from other languages, words that would be hard for a child to read, and all vowels. Also, the letter *s* can correspond to many sounds, such as the phoneme /z/ as it does in *busy* and *dogs*. Similarly, the letter *d* can correspond to the phoneme /t/, as in *asked*, *missed*, and *pushed*.

Multiple pronunciations

You can provide multiple pronunciations for any word. The number of pronunciations you want to provide depends on how many dialects you want your system to recognize. List multiple pronunciations on separate lines. For example, the following table lists some examples of words that could have multiple pronunciations:

Table 10: EXAMPLES OF WORDS WITH MULTIPLE PRONUNCIATIONS

Word	Phonemes
either	aj D *r
	i D *r
compass	k ^ m p * s
	k A m p * s
defense	d * f E n s
	d i f E n s
sandwich	s a n d w I tS

Table 10: EXAMPLES OF WORDS WITH MULTIPLE PRONUNCIATIONS

Word	Phonemes
	s a n w l tS

Sample dictionary file

The following lists a sample dictionary made up of the words in Table 9.

able	e b * l
adam	a ! * m
agent	e dZ * n t
alive	* l a j v
aphid	e f * d
apple	a p * l
asked	a s k t
ball	b O l
bird	b *r d
book	b U k
bought	b O t
butter	b ^ ! *r
cat	k a t
caught	k O t
child	tS a j l d
coat	k o t
compass	k ^ m p * s
confuse	k * n f j u z
cow	k a w
cry	k r a j
dogs	d A g z
dogs	d O g z
echo	E k o
even	i v * n
father	f A D *r
focus	f o k * s
golf	g A l f
grammar	g r a m *r
hotel	h o t E l
india	I n d i *
iraq	* r A k
iraq	I r A k
iraq	a j r A k
iraq	* r a k
iraq	I r a k
iraq	a j r a k
ivory	a j v *r i
ivory	a j v r i

jump	dZ ^ m p
key	k i
kissed	k I s t
knot	n A t
law	l O
lisa	l i s *
loop	l u p
lunar	l u n *r
mike	m a j k
nation	n e S * n
new	n u
new	n j u
olive	A l * v
open	o p * n
pan	p a n
pawn	p O n
proceed	p r * s i d
proceed	p r o s i d
pull	p U l
put	p U t
putt	p ^ t
question	k w E s tS * n
question	k w E S tS * n
quick	k w I k
rider	r a j ! *r
ring	r I g~
rude	r u d
sanity	s a n * ! i
shell	S E l
shield	S i l d
shoe	S u
show	S o
soup	s u p
them	D E m
thin	T I n
thumb	T ^ m
time	t a j m
under	^ n d *r
usage	j u s * dZ
usage	j u s I dZ
very	v E r i
vision	v I Z * n
which	w I tS
widen	w a j ! * n
wisdom	w I z d * m
would	w U d
xenophobia	z i n * f o b i *

yes	j E s
zebra	z i b r *

To create pronunciations, you can also use the command-line program *pronounce* to get pronunciations for words that have similar pronunciations. For example, to generate a pronunciation for the word “glyph,” you could base it on the pronunciation for the word “cliff.” To use *pronounce*, specify the master package you are using and a list of one or more words for which you would like to see pronunciations. See the online documentation for *pronounce* for more details.

Converting non-CPA dictionaries

If you have Nuance dictionary files you created with a version of the Nuance System older than 6.2, you need to convert these to CPA. The Nuance System includes the tool *arpabet-to-cpa* that can do this for you. See the online documentation for usage information.

Phoneme sets for other languages

In addition to American English, Nuance supports a variety of other languages and English dialects. To work with other languages, you need to use a specific Nuance master package to process that language. See the section “Choosing a master package” on page 101.

Tables for mapping the phonemes for all supported languages to their CPA representation are located at the documentation site on the Nuance Developer Network at *extranet.nuance.com*.

GSL reference



This appendix provides a formal definition of the syntax of the Grammar Specification Language (GSL), and a summary of all the package files used to compile a package.

The following conventions are used in the syntax description listed below:

Expression	Meaning
X	A nonterminal
“xyz”	Literal string xyz should be typed into the specification
X ::= Y	X consists of Y
X Y Z	X or Y or Z
X Y	X followed by Y, with white space allowed
X~Y	X followed by Y with no white space intervening

GSL syntax

The following rules formally define the GSL syntax. However, these rules may—in a few cases—generate some non-valid expressions, and therefore, their use should be complemented with the notes following the syntax rules.

```
GrammarFile ::= GrammarDefinitions
GrammarDefinitions ::=
    GrammarDefinition |
    GrammarDefinition GrammarDefinitions
GrammarDefinition ::=
    GrammarNameWithOptionalSuffixes GrammarConstruct |
    GrammarNameWithOptionalSuffixes
```

```

GrammarNameWithOptionalSuffixes ::=
    GrammarName~":"~Suffixes |
    GrammarName
Suffixes ::= Suffix | Suffix~","~Suffixes
Suffix ::= "dynaref" | "dynamic" | "filler"
GrammarConstruct ::= BareConstruct |
    BareConstruct Commands |
    "=" GrammarConstruct |
    GrammarName~":slm" "\"AnyString "\" |
    GrammarName~":slm" "\"AnyString "\" "=" GrammarConstruct
BareConstruct ::=
    BareConstructWithoutProb |
    BareConstructWithoutProb~"~"~Prob
BareConstructWithoutProb ::= Word |
    GrammarReference |
    "?" GrammarConstruct |
    "*" GrammarConstruct |
    "+" GrammarConstruct |
    "[" GrammarConstructs "]" |
    "(" GrammarConstructs ")"
GrammarConstructs ::=
    GrammarConstruct |
    GrammarConstruct GrammarConstructs
Word ::= LowerCaseLetter | Digit | WordChar~Word
GrammarReference ::= GrammarName |
    GrammarName~":"~Variable |
    Ext_Grammar_Ref
GrammarName ::= LimitedString
Commands ::= "{" CommandList "}"
CommandList ::= Command | Command CommandList
Command ::= "<" Slot Value ">" |
    "return (" Value ")" |
    "insert-begin (" Variable Value ")" |
    "insert-end (" Variable Value ")" |
    "concat (" Variable Value ")"
Slot ::= SlotString
SlotString ::= VOrFChar | VOrFChar~SlotString
Values ::= Value | Value Values
Value ::= Integer |
    LimitedString |
    "\""~AnyString~"\"" |
    Structure |
    "$"~VariableExpression |
    List |
    FunctionApplication
Structure ::= "[" FeatureValuePairs "]"

```

```

FeatureValuePairs ::=
    FeatureValuePair |
    FeatureValuePair FeatureValuePairs
FeatureValuePair ::= "<" Feature Value ">"
Feature ::= VariableOrFeatureString
Variable ::= VariableOrFeatureString
VariableExpression ::= Variable | Variable~"."~FeaturePath
FeaturePath ::= Feature | Feature~"."~FeaturePath
List ::= "()" | "(" Values ")"
FunctionApplication ::= Function~ "(" Values ")"
Function ::= LimitedString
WordChar ::= LowerCaseLetter | "-" | "_" | "." | "@" | "'" | Digit
LimitedString ::= LimitedChar | LimitedChar~LimitedString
LimitedChar ::= UpperCaseLetter |
    LowerCaseLetter |
    Digit |
    "-" | "_" | "@" | "'" | "."
Digit ::= 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9
VariableOrFeatureString ::=
    VOrFChar |
    VOrFChar~VariableOrFeatureString
VOrFChar ::= UpperCaseLetter |
    LowerCaseLetter |
    "-" | "_" | "@" | "'"
AnyString ::= Char | Char~AnyString

```

Notes

- A semicolon indicates a comment. The compiler ignores everything on the line, after a semicolon.
- Grammar names must contain at least one uppercase letter.
- White space is a sequence of at least one of the following characters: space, tab, new line, return, and page feed.
- Char signifies any character other than the double quote or white space.
- Prob denotes a *non negative* floating point number not greater than 1.0.
- The following strings are reserved for internal use and cannot be used to denote a word or grammar name (*n* designates any integer):

AND-*n*, OR-*n*, OP-*n*, KC-*n*, PC-*n*

External grammar reference syntax

The Grammar Specification Language has been extended to include the following grammar reference syntax rules.

```
Ext_Rule_ref ::= ExRef_Resolved | "<" ExRef_Resolved "%" ExRef_Special ">"
```

```
ExRef_Resolved ::= "<" ExRef_Simple
                  ["!" DynGrammLabel "=" ExRef_Simple
                  ["&" DynGrammLabel "=" ExRef_Simple]*] ">"

ExRef_Simple ::= "<ExRef_Type~:"~ExRef_Path[~#"~ExRef_StartRule]">"
ExRef_Type  ::= "http" | "file" | "dgdb" | "static" | "builtin"
ExRef_Special ::= "<special:"~Special_Path~">" | "<fail:"~ExRef_Path~">"
Special_path ::= "NULL" | "VOID" | "passthrough" | "roadblock" |
                  "resistor?weight=~Prob"
```

Some notes:

- The `Special_Path` specification `NULL` is equivalent to `passthrough`, and the `Special_Path` specification `VOID` is equivalent to `roadblock`.
- `ExRef_Path` is any string not containing the characters `#` or `>`.
- `ExRef_StartRule` is any valid public rule name.

`DynGrammLabel` is the name of a grammar annotated as `:dynamic []` in the associated simple reference.

Summary of package files

The table below lists all the files that *nuance-compile* takes into account when compiling a recognition package, with a brief description of the expected contents.

The common name of these files is the name of the package being compiled. For example, you compile the file *myapp.grammar* to generate a recognition package called *myapp*. All files should be stored in the same directory.

Table 1: Recognition package files

File extension	Required?	Contents
<i>.grammar</i>	Yes	GSL code
<i>.dictionary</i>	No	One word (or compound) followed by a sequence of phonemes per line
<i>.slot_definitions</i>	No	One slot name per line
<i>.functions</i>	No	User-defined functions

Advanced SLM features

B

This section provides advanced information on the Nuance SLM utilities. If you are new to SLM features, please follow the procedures described in Chapter 5, “Say Anything: Statistical language models and robust interpretation”.

For background information on SLMs, please see the following books:

- F. Jelinek, *Statistical Methods for Speech Recognition*, The MIT Press, Cambridge, 1997.
- S. Young, and G. Bloothoof, *Corpus-Based Methods in Language and Speech Processing*, Kluwer Academic Publishers, Norwell, 1997.

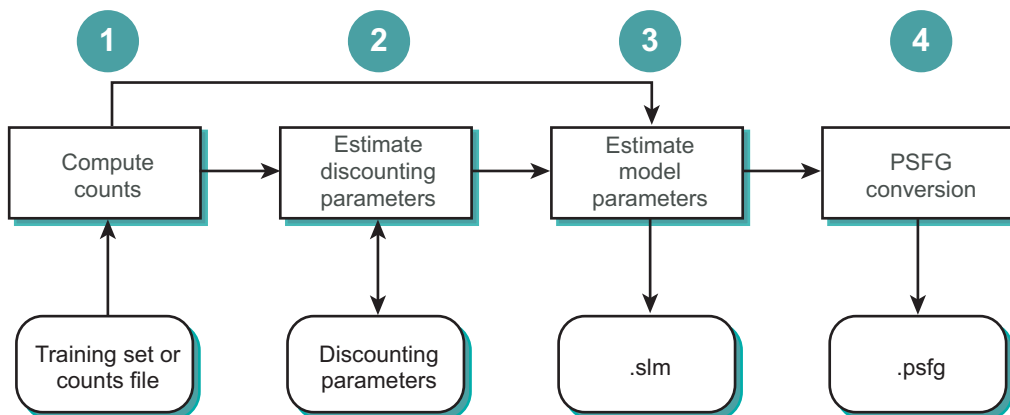
The *train-slm* tool

The *train-slm* tool, Nuance's SLM trainer, is a command-line tool that trains n-gram language models.

Training an n-gram model is done in four phases:

- 1 Load a training set, either from a text file (one sentence per line) or from a counts file. This phase computes the counts of all n-grams that occur in the corpus.
- 2 Estimate the discounting parameters that will be used to smooth the relative frequencies.
- 3 Estimate the model probabilities from the counts.
- 4 Convert the estimated model to a Probabilistic Finite State Grammar (PFSG), the structure used to compile recognition packages.

This process is illustrated in the figure below.



The trainer can perform the four steps in a single command-line invocation, as described in Chapter 5, or you can invoke the trainer to perform individual tasks. You can save the output of the intermediate steps to files and use them in future runs of the trainer. This capability lets you, for example, run Step 3 again with different parameters without having to repeat Step 1 and Step 2.

Once the model is trained, it can be saved in SLM format in a *.slm* file and in PFSG format in a *.psfg* file. The SLM format is a proprietary format for storing the language model. It is the input to *process-slm*, the utility that performs post-training operations such as perplexity measurements and conversion to PFSG. The PFSG format is the result of converting a language model to a Probabilistic Finite State Grammar, the data structure used by the Nuance System recognizer. This is the file read by *nuance-compile* to produce recognition packages.

Usage

The *train-slm* tool can be used with either a text corpus or a counts file as input:

```

train-slm
  -corpus input-corpus-filename | -counts input-counts-filename
  [DISCOUNTING_OPTIONS]
  [SLM_OPTIONS]
  [OPTIONS]

```

where [DISCOUNTING_OPTIONS] are:

```

[ -discounts input-discounts-filename |
  -discounting-type GOOD-TURING |
                        ABSOLUTE_DISCOUNTING |
                        KATZ_DISCOUNTING
  [-discounting-observ num-observ] ]

```

where [SLM_OPTIONS] are:

```
[ -slm output-slm-basename
  [-pfsg-type STANDARD [-keep-lowprobs] |
    FULL-EXPANSION |
    ADJUST-BOW ]
  [-no-pfsg]]
```

where [OPTIONS] are:

```
[-help]
[-no-unknown]
[-n N]
[-vocab vocab-filename]
[-verbose verbose_level]
[-write-counts output-counts-filename]
[-write-discounts output-discounts-filename]
```

Required input

The *train-slm* tool requires one of the following input files, specified with the *corpus* or *counts* option:

-corpus input-corpus-filename

(Step 1) Loads the n-gram counts from a text corpus file. The corpus files contains one sentence per line. See “Create the training set” on page 73 for more information.

-counts input-counts-file

(Step 1) Loads the n-gram counts from a counts file. Each line of the counts file contains an n-gram (a sequence of n words) followed by an integer (the count for that n-gram).

Discounting options

The *train-slm* tool can take the following optional discounting options:

```
[ -discounts input-discounts-filename |
  -discounting-type GOOD-TURING |
    ABSOLUTE_DISCOUNTING |
    KATZ_DISCOUNTING
  [-discounting-observ num-observ] ]
```

-discounts input-discounts-filename

(Step 2) Instead of estimating the discounting parameters from the loaded corpus, reads them from file *input-discounts-filename*. Typically, you produce this file by running *train-slm* with the *-write-discounts* option. You can then edit it manually to change the discounting parameters. See “Editing discounting parameters” on page 153 for more information.

-discounting-type discounting-method

(Step 2) Controls which type of discounting algorithm gets trained. This option cannot be used with the *-discounts* option: the discounting

parameters are either loaded from a file or estimated. Three discounting methods are supported: Each method uses discounting parameters that are estimated from the loaded corpus.

- **GOOD-TURING (default)**—The Good-Turing formula is used for counts in the range `[lower_bound, upper_bound]` where `lower_bound` and `upper_bound` are parameters estimated by *train-slm*. Under `lower_bound`, counts are discounted to 0. Above `upper_bound`, counts are multiplied by a discounting factor `alpha`, a number close to but smaller than 1.0.
- **ABSOLUTE_DISCOUNTING**—A count of one is reduced by a constant `b1`, estimated by *train-slm*. All other counts are reduced by the same constant, `b2`. The algorithm also uses lower and upper bounds for discounting. Below the `lower_bound`, a count is discounted to 0; above the `upper_bound`, no discounting occurs.
- **KATZ_DISCOUNTING**—This is a variant of GOOD-TURING in which the counts in `[lower_bound, upper_bound]` are discounted by using the Good-Turing formula and multiplied by a discounting factor `alpha` that is typically close to 1.0. Above `upper_bound`, the counts are unchanged.

`-discounting-observ num-observ`

(Step 2) Controls the number of observations used to estimate the discounting parameters. Only the first *num-observ* observations from the training corpus are used to estimate the parameters. Because the vocabulary must not be applied to the training corpus to train the discounting methods, Step 2 can take a lot of memory, even more than the model training itself. This option is useful to limit the amount of memory used by Step 2. The default behavior is to use the full training set. Note that this option cannot be used with the *-discounts* or *-counts* options.

SLM options

The *train-slm* tool can take the following SLM options:

```
[ -slm output-slm-basename
  [-pfsg-type STANDARD [-keep-lowprobs] |
                                FULL-EXPANSION |
                                ADJUST-BOW ]
  [-no-pfsg] ]
-slm output-slm-basename
```

(Step 3) Tells the trainer to train an SLM. The resulting model is saved to file *output-slm-basename.slm* in SLM format and to *output-slm-basename.pfsg* in PFSG format.

`-pfsg-type pfsg-type`

(Step 4) Controls the conversion algorithm used to produce the PFSG. Three methods are available:

- **STANDARD (default)**—Creates a standard PFSG where a backoff distribution is used for n-grams that were not observed in the training set. Low probability n-grams are pruned (see `-keep-lowprobs`).
- **FULL-EXPANSION**—The backoff distributions are not used and the low probability n-grams are NOT pruned. Each distribution has V transitions, where V is the vocabulary size.
- **ADJUST-BOW**—Similar to the **STANDARD** conversion but the backoff weight is decreased to ensure that no low probabilities remain. Thus, low probabilities are NOT pruned.

`-keep-lowprobs`

(Step 4) Only valid when `pfsg-type` is **STANDARD**. Instructs the trainer not to prune low probabilities when converting to PFSG. A low probability occurs when a distribution word's probability is smaller than the path through the backoff distribution. By default, the low probability n-grams are pruned.

`-no-pfsg`

(Step 4) Does not generate the PFSG. By default, a **STANDARD** PFSG is generated.

Additional options

The *train-slm* tool can take the following additional options:

```
[ -help ]
[ -no-unknown ]
[ -n N ]
[ -vocab vocab-filename ]
[ -verbose verbose_level ]
[ -write-counts output-counts-filename ]
[ -write-discounts output-discounts-filename ]
```

`-help`

Provides a detailed command-line usage.

`-no-unknown`

(Steps 1 and 3) Trains a closed language model, that is, one that does not assign probabilities to phrases containing out-of-vocabulary words. For an open language model, the PFSG uses the UNK grammar rule that must be defined by the user. Closed model training ignores the n-grams whose last word is out-of-vocabulary; also, n-grams whose history contains an out-of-vocabulary word are truncated to erase the word.

The default is to train an open language model and use the unknown class UNK to represent out-of-vocabulary words. See “Open and closed vocabulary SLMs” on page 78 for more information.

`-n N`

(Steps 1, 2, and 3) Trains an n-gram SLM of order *N*. This number determines the order of counts to compute and has implications as to the amount of memory used by *train-slm*. Default is 3 (trigram). See “Determine the order of the model” on page 79 for more information.

`-vocab input-vocab-filename`

(Steps 1 and 3) Uses the vocabulary file *input-vocab-filename* to train the n-gram SLM. This file contains the words that need to be covered by the produced SLM. See “Create a vocabulary file (optional)” on page 77 for more information.

A vocabulary file contains one word per line. Out-of-vocabulary words are processed depending on whether the option *-no-unknown* is specified or not. Using a vocabulary decreases the memory requirements of the training process.

`-verbose verbose-level`

(Steps 1, 2, and 3) Controls the amount of information output by the trainer; *verbose-level* is an integer between 0 and 4:

- 0—No verbose
- 1—Low verbose
- 2—Medium verbose
- 3—High verbose
- 4—Extreme verbose

`-write-counts output-counts-filename`

(Step 1) Saves the counts that were computed from a corpus to the file specified with *output-counts-filename*.

`-write-discounts output-discounts-filename`

(Step 2) Saves the discounting parameters to the file *output-discounts-filename*. The saved discounting parameters can be used in a future trainer run with the *-discounts* option.

Editing discounting parameters

You can edit the discounting parameters by creating a parameter file, editing it, and then using *train-slm*, as described below.

- 1 To output the discounting parameters to a file, enter:

```
> train-slm -corpus input-file -write-discounts  
output-discounts-filename other_appropriate_options
```

where *input-file* is the name of the training set, *output-discounts-filename* is the name of the file that will contain the discounting parameters, and *other_appropriate_options* is any of the options you want to specify (such as *-n*). For example:

```
> train-slm -corpus training.txt -write-discounts params.txt
```

This command outputs a text file that contains the discounting parameters according to the discounting type specified: GOOD-TURING (default), ABSOLUTE_DISCOUNTING, or KATZ_DISCOUNTING. See “Discounting options” on page 149 for more information about the discounting parameters.

- 2 Edit the file as required.
- 3 To train the SLM grammar using the updated parameters, enter:

```
> train-slm -corpus input-file -discounts  
input-discounts-filename other_appropriate_options
```

where *input-file* is the name of the training set, *input-discounts-filename* is the name of the discounting parameter file that you updated, and *other_appropriate_options* is any of the options you want to specify (such as *-n*). For example:

```
> train-slm -corpus training.txt -discounts params.txt
```

The *process-slm* tool

You use *process-slm* to perform post-training operations on SLMs. It takes a *.slm* file produced by *train-slm* or an ARPA file as input and can perform the following operations on it:

- Perplexity measurements
- Conversion to *.pfsg* format
- Model interpolation
- N-gram pruning

The resulting language model is then saved to a *.slm* or a *.pfsg* file.

Usage

The *process-slm* tool has the following command-line syntax:

```
process-slm
  slm_basename1 slm_basename2 ...
  [-ppl-corpus input-corpus-filename]
  [-interpolate interpolated-slm-basename -w wgt1 -w wgt2 ...]
  [-verbose verbose-level]
  [-prune-lowprobs]
  [-pfsg
    [-pfsg-type STANDARD [-keep-lowprobs] |
      FULL-EXPANSION |
      ADJUST-BOW ] ]
```

Required input

slm_basename1 slm_basename2 ...

Specifies the path of the models to operate on. You do not need to specify the *.slm* extension. You must specify at least two input models for the *-interpolate* option. For all other options, each model is processed one by one.

Options

-ppl-corpus input-corpus-filename

Computes the perplexity of each input model on the training corpus specified by *input-corpus-filename*. The text corpus contains one sentence per line. See “Create the training set” on page 73 for more information.

-interpolate interpolated-slm-basename

Performs the linear interpolation of the input models and saves the result to file *interpolated-slm-basename.slm*. The interpolation weights are specified by the *-w* option. There must be as many *-w* options as there are input models. The weights do not need to sum to one. If they don't, they are normalized.

-w wgtx

Specifies an interpolation weight (a float).

-verbose verbose-level

Controls the amount of information output by the trainer; *verbose-level* is an integer between 0 and 4:

- 0—No verbose
- 1—Low verbose
- 2—Medium verbose

- 3—High verbose
- 4—Extreme verbose

`-prune-lowprobs`

Prunes the low probability n-grams from each input model and saves it in SLM format to file *slm_basenameX.pruned.slm*. See the `-keep-lowprobs` option of *train-slm* for more information.

`-pfsg`

Converts each input SLM to a PFSG and saves it to file *slm_basenameX.pfsg*.

`-pfsg-type pfsg-type`

(Step 4) Controls the conversion algorithm used to produce the PFSG. Three methods are available:

- STANDARD (default)—Creates a standard PFSG where a backoff distribution is used for n-grams that were not observed in the training set. Low probability n-grams are pruned (see `-keep-lowprobs`).
- FULL-EXPANSION—The backoff distributions are not used and the low probability n-grams are NOT pruned. Each distribution has V transitions, where V is the vocabulary size.
- ADJUST-BOW—Similar to the STANDARD conversion but the backoff weight is decreased to ensure that no low probabilities remain. Thus, low probabilities are NOT pruned.

`-keep-lowprobs`

Only valid when *pfsg-type* is STANDARD. Instructs the trainer not to prune low probabilities when converting to PFSG. A low probability occurs when a distribution word's probability is smaller than the path through the backoff distribution. By default, the low probability n-grams are pruned.

W3C grammar support

C

The World Wide Web Consortium (W3C) has proposed a syntax for speech recognition grammar formats (SRGF). This proposal includes two syntax forms: Augmented BNF (ABNF) and XML (GrXML).

Nuance supports W3C grammars in the XML format. This section first describes how you can use a W3C grammar with the Nuance System. It then lists the Nuance extensions to the W3C format. Finally, it lists the W3C features not currently supported by Nuance.

Note: This chapter does not explain how to write a W3C grammar. For details on the W3C grammar syntax in GrXML form, see the specification published at www.w3.org/TR/2001/WD-speech-grammar-20010820/. Nuance is currently working from the draft dated **August 20, 2001**. Since the W3C specification will likely change, grammars developed in Nuance 8.0 may not be compatible with future releases.

Working with W3C grammars in the Nuance System

Most of the concepts in the Nuance Grammar Specification Language, introduced in this document, are also available in the W3C grammar specification. However, each language has some features and concepts that are slightly different from the other.

This section presents some aspects of W3C grammar features that you should be aware of when using this format with the Nuance System, such as:

- Supported encodings
- Differences in the notion of *token*
- GARBAGE rule handling

- DTMF mode

Supported encodings

Nuance supports the following encodings in a GrXML specification:

- UTF-8
- UTF-16
- ISO-8859-1
- US-ASCII

Support for other encodings may be added in future releases.

Note: You must make sure that the locale setting for your environment is appropriately set for the language of the grammar you're compiling. See the *Nuance System Installation Guide* for more information.

Tokens

Unlike the GSL specification syntax for a word, a token in a W3C grammar may contain whitespace—if properly quoted—and may contain uppercase characters.

To reconcile this distinction, any token in a W3C grammar that contains whitespace is further broken into individual words by the Nuance System. The Nuance multi-word feature—which allows for pronunciation of a group of words, such as “what is” and “new york”—works at dictionary lookup time and performs its search across *separate* words composing the multi-word. For more information, see “Phrase pronunciations” on page 108.

W3C grammars allow for capital letters in a token. However, for dictionary lookups and the auto-pronunciation to work, the Nuance implementation converts all letters to lowercase as part of the compilation process. This means that you cannot provide separate pronunciations for the same word written with different cases.

Handling of the GARBAGE rule

Nuance maps the special GARBAGE rule definition:

```
<ruleref special="GARBAGE"/>
```

to the following block:

```
<item repeat="0-1">  
  @reject@
```

</item>

This amounts to 0 or more occurrences of the @reject@ word being recognized wherever the GARBAGE ruleref is used.

Nuance recommends that you tune the value of the recognition parameter `rec.RejectWeight` when using the GARBAGE rule in your grammars.

Caution: To minimize inaccuracies that the GARBAGE rule processing may produce, use this special rule with care. Instead, try to come up with an explicit model of the expected words—or utterances—that you want your grammar to discard. For example, you could train these models using Say Anything grammars. See Chapter 5, “Say Anything: Statistical language models and robust interpretation,” for more information.

DTMF mode

The W3C specification indicates that a grammar can be in one of two modes: `voice` (the default mode) or `dtmf`, depending on the type of input that the speech engine should be detecting.

As with GSL, you can specify DTMF sequences in a W3C `voice` grammar, using the tokens `dtmf-1`, `dtmf-2`, and so on. These tokens are only valid in a `voice` grammar. In a `dtmf` grammar, an automatic translation of phone buttons to `dtmf` tokens takes place, according to the following map:

Phone button	dtmf token
0	dtmf-0
1	dtmf-1
2	dtmf-2
3	dtmf-3
4	dtmf-4
5	dtmf-5
6	dtmf-6
7	dtmf-7
8	dtmf-8
9	dtmf-9
star	dtmf-star

Phone button	dtmf token
*	dtmf-star
pound	dtmf-pound
#	dtmf-pound

Any token outside this range will cause a compilation error. Note that, even though a grammar may be in `dtmf` mode, it will still be processed in parallel with the recognition engine, natural language will still be interpreted by a recognition server, and high background noise may cause a rejection.

Note: When the Nuance System handles the telephony, DTMF-only recognition—that is, with no speech—is not supported.

Nuance extensions to W3C grammars

Nuance has extended the GrXML grammar specification to provide additional functionality available in GSL grammars.

The standard GrXML elements that were extended include:

- `<rule>`
- `<ruleref>`

The Nuance-specific GrXML attributes and elements are all in the namespace `http://voicexml.nuance.com/grammars`, as shown in the following DTD specification fragment:

```
<!ATTLIST grammar
    ...
    <!-- Nuance extensions -->
    xmlns:nuance CDATA #FIXED "http://voicexml.nuance.com/grammar"
    ...>
```

The rest of this document will use `nuance:` as an abbreviation for this namespace.

The `<rule>` element

The `<rule>` element was extended to include the `nuance:dynamic` attribute; this attribute lets you declare a grammar rule as dynamic.

The extension to the DTD that defines this attribute is (extension in bold face):

```
<!ELEMENT rule (%rule-expansion; | example)*>
```

```

<!ATTLIST rule
    ...
    nuance:dynamic (true | false) "false"
    ...>

```

Here is an example of how to declare a rule to be dynamic:

```

<grammar xmlns="http://www.w3.org/2001/06/grammar"
    xmlns:nuance="http://voicexml.nuance.com/grammar"
    xml:lang="en" version="1.0">
    <rule id="personal" nuance:dynamic="true">
        ...
    </rule>
</grammar>

```

To see more examples of dynamic rules, see “Sample grammars” on page 164.

The <ruleref> element

The <ruleref> element was extended to:

- Allow for the redefinition of external rules from the reference.
- Specify a backoff reference in case the referenced rule cannot be resolved—for example, because of an invalid URL or a timeout.

Note: The just-in-time grammar specification provides two operators to handle these situations: the operator `!`, to specify the redefinition of external rules, and the operator `%`, to specify backoff references. See “External rule references and redefinitions” on page 25.

The Nuance extension to <ruleref> allows it to contain, in addition to the W3C-specified elements:

- Zero or more <nuance:redef> elements
- An optional <nuance:backoff> element

Note that only an *external* reference <ruleref>, specified with the attribute `uri`, may contain redefinitions and backoffs.

The portion of the expanded DTD that accounts for these elements is:

```

<!ELEMENT nuance:redef EMPTY>
<!ATTLIST nuance:redef
    name CDATA #REQUIRED
    uri CDATA #IMPLIED
    special CDATA #IMPLIED>

<!ELEMENT nuance:backoff EMPTY>

```

```

<!ATTLIST nuance:backoff
    uri CDATA #IMPLIED
    special CDATA #IMPLIED>

<!ELEMENT ruleref (nuance:redef*, nuance:backoff?)>

```

For the `<nuance:redef>` element, the `name` attribute specifies the rule in the reference that is being redefined, and the `uri` or `special` attribute specifies the new rule definition.

The optional `<nuance:backoff>` element is used if the original reference cannot be resolved. Here is a GrXML fragment illustrating the use of these elements:

```

<grammar xmlns="http://www.w3.org/2001/06/grammar"
    xmlns:nuance="http://voicexml.nuance.com/grammar"
    xml:lang="en" version="1.0">
...
    <ruleref uri="http://faraway-server.com/cgi/grammar.pl">
        <nuance:redef name="personal"
            uri="http://myserver/grammars/addrbook.grxml"/>
        <nuance:redef name="gate" special="NULL"/>
        <nuance:backoff special="VOID"/>
    </ruleref>
...
</grammar>

```

To see more examples of how to use the `<ruleref>` element, see “Sample grammars” on page 164.

Unsupported features

This section discusses several W3C grammar features not currently supported by the Nuance System.

Note: Support for some of these features may be added in the near future, as part of option packs or service packs, as appropriate.

ABNF

ABNF, a compact GSL-like optional format, is not supported in this release.

Elements and attributes

The following are currently ignored by the Nuance System:

- `<meta>` elements of type `http-equiv`

- `<lexicon>` elements
- `type` attribute of the `<ruleref>` and `<alias>` elements

The Nuance System generates warnings when these elements are used.

W3C semantic interpretation tags

The exact format for a semantic interpretation tag—natural language interpretation, in GSL terms—in a W3C grammar is not yet specified.

To append a semantic capability to your W3C grammars:

- Use the Nuance natural language interpretation syntax.
- Set the optional grammar element attribute `tag-format` to Nuance to specify the semantic format that your grammar is using.
- Do not enclose the semantic tag in curly brackets. In a GrXML document, this semantic tag is already delimited by the tag element, as in the following fragment:

```
<tag><![CDATA[<command $n> <num $n>]]></tag>
```

NL features

Some NL operations available in GSL have no counterpart in W3C grammars. Most notably, there is no equivalent to the following way of referencing a grammar in GSL and assigning its returned value to a variable:

```
SomeGrammarName:aVariableName
```

Grammar syntax and semantic tags are completely disjoint in W3C grammars. Thus, the Nuance NL interpretation syntax was expanded to account for this fact by providing the *\$return* variable and the *assign* command. See “Assigning variables” on page 54 for more information.

Language specification

W3C grammars may contain language specifications at the grammar, expansion, and individual token level. The Nuance compiler ignores these specifications and tries to generate pronunciations for missing words as best as possible, according to the current master package.

Use the `Nuance-Package-Name` and `Nuance-Config-Name` keys of the `RECOGNIZE` header to select the appropriate master package and compilation server for your language. See “Specifying just-in-time grammar options: Using the keyword `RECOGNIZE`” on page 40 for more information.

Using W3C grammars

Once your W3C grammar is written, you can use it just as you use any other dynamic grammar, including:

- Through an external rule reference
- In a just-in-time request
- By compiling the grammar to an NGO file
- Using `newDynamicGrammar` functions

Note: The *nuance-compile* utility does not handle the static compilation of a W3C grammar.

See Chapter 3, “Dynamic grammars: Just-in-time grammars and external rule references,” for more information on external rule references and just-in-time requests. See Chapter 6, “Compiling grammars,” for more information on NGO files. See the *Nuance Application Developer’s Guide* for more information about the `newDynamicGrammar` functions.

Sample grammars

The following samples illustrate how to use W3C grammars.

Sample dialer application grammar

This grammar allows the caller to dial people by name or by number, in response to a question such as “who would you like to call?”

Entries dialed by name may either be from a pre-compiled static corporate grammar or a dynamically specified personal grammar. The top-level grammar ‘dial’ will fill either the number or the name slots. If the name slot is filled, it will fill the source slot to specify whether or not the corporate or personal grammar was used.

```
<?xml version="1.0"?>
<grammar xmlns="http://www.w3.org/2001/06/grammar"
  xmlns:nuance="http://voicexml.nuance.com/grammar"
  xml:lang="en" version="1.0" mode="voice"
  tag-format="Nuance" root="dial">

  <!-- Alias a standard telephone number grammar that returns -->
  <!-- the value of the telephone number recognized. -->
  <alias name="number" uri="telephonenumber.gsl" />

  <!-- Primary rule -->
  <rule id="dial" scope="public">
    <!-- Low-prob the optional pre_hesitation grammar -->
```

```

<item repeat="0-1" repeat-prob="0.01">
  <ruleref uri="#pre_hesitation"/>
</item>

<!-- Low-prob the optional pre_filler grammar -->
<item repeat="0-1" repeat-prob="0.01">
  <ruleref uri="#pre_filler"/>
</item>

<!-- The real contents of the grammar -->
<ruleref uri="#core"/>

<!-- Low-prob the optional post_filler grammar -->
<item repeat="0-1" repeat-prob="0.00001">
  <ruleref uri="#post_filler"/>
</item>
</rule>

<!-- Main contents of the grammar -->
<rule id="core">
  <one-of>
    <item>
      <ruleref alias="number#MAIN"/>
      <tag><![CDATA[<number $return>]]></tag>
    </item>
    <item>
      <!-- if the package into which this is inserted does -->
      <!-- not have a static grammar called CORPORATE, -->
      <!-- backoff to a roadblock -->
      <ruleref uri="static:CORPORATE">
        <nuance:backoff special="VOID"/>
      </ruleref>
      <tag><![CDATA[<source corporate> <name $return>]]></tag>
    </item>
    <item>
      <ruleref uri="#personal"/>
      <tag><![CDATA[<source personal> <name $return>]]></tag>
    </item>
  </one-of>
</rule>

<!-- Dynamic placeholder for the personal grammar -->
<rule id="personal" nuance:dynamic="true">
  <ruleref special="VOID"/>
</rule>

<!-- Fillers -->
<rule id="pre_hesitation">

```

```

    <one-of>
      <item>uh</item>
      <item>um</item>
      <item>hm</item>
    </one-of>
  </rule>

  <rule id="pre_filler">
    <one-of>
      <item>
        <one-of>
          <item>i would like</item>
          <item>i'd like</item>
          <item>i want</item>
          <item>on</item>
        </one-of>
        to
        <one-of>
          <item>call</item>
          <item>dial</item>
        </one-of>
      </item>
      <item>
        <item repeat="0-1">
          please
        </item>
        call
      </item>
      <item>
        <item repeat="0-1">
          please
        </item>
        dial
      </item>
    </one-of>
  </rule>

  <rule id="post_filler">
    <one-of>
      <item>please</item>
    </one-of>
  </rule>

</grammar>

```

You could then use this grammar in a just-in-time grammar and redefine the personal rule, as follows:

```

<?xml encoding="1.0"?>
<grammar xmlns="http://www.w3.org/2001/06/grammar"
          xmlns:nuance="http://voicexml.nuance.com/grammar"
          version="1.0" mode="voice" xml:lang="en-US"
          tag-format="Nuance" root="main">

  <rule id="main" scope="public">
    <ruleref uri="dialer.grxml">
      <nuance:redef name="personal" uri="#names"/>
    </ruleref>
  </rule>

  <rule id="names" scope="public">
    <one-of>
      <item>mom</item>
      <item>dad</item>
    </one-of>
    <tag>return($string)</tag>
  </rule>

</grammar>

```

Sample retail application grammar

This example shows a sample retail application grammar for purchasing items in a hockey equipment store:

```

<?xml version="1.0" ?>
<grammar xmlns="http://www.w3.org/2001/06/grammar"
          xmlns:nuance="http://voicexml.nuance.com/grammar"
          xml:lang="en-US" tag-format="Nuance" root="main">

  <rule id="main" scope="public">
    <item>
      <item repeat="0-1">
        <ruleref uri="#number"/>
        <tag>assign(n $return)</tag>
      </item>
      <ruleref uri="#items"/>
      <tag>assign(i $return)</tag>
    </item>
    <tag><![CDATA[<number $n> <item $i>]]></tag>
  </rule>

  <rule id="items" scope="public">
    <one-of>
      <item>
        <one-of>
          <item>helmet</item>
        </one-of>
        <item>helmets</item>
      </item>
    </one-of>
  </rule>

```

```

        <one-of>
            <item>stick</item>
        <item>sticks</item>
        </one-of>
    </item>
    <item>
        <one-of>
            <item>puck</item>
        <item>pucks</item>
        </one-of>
    </item>
    <item>
        <item repeat="0-1">
            <one-of>
                <item>pair of</item>
                <item>pairs of</item>
            </one-of>
        </item>
        <item repeat="0-1">hockey</item>
            <one-of>
                <item>pant</item>
                <item>pants</item>
            </one-of>
        </item>
    <item>
        <item repeat="0-1">
            <one-of>
                <item>pair of</item>
                <item>pairs of</item>
            </one-of>
        </item>
        socks
    </item>
    <item>
        <item repeat="0-1">
            <one-of>
                <item>pair of</item>
                <item>pairs of</item>
            </one-of>
        </item>
        skates
    </item>
    <item>
        <one-of>
            <item>jersey</item>
        <item>jerseys</item>
        </one-of>
    </item>
</one-of>
<tag>return($string)</tag>
</rule>

<!-- Simple number grammar -->

```

```

<rule id="number">
  <one-of>
    <item>one   <tag>return(1)</tag></item>
    <item>two   <tag>return(2)</tag></item>
    <item>three <tag>return(3)</tag></item>
    <item>four  <tag>return(4)</tag></item>
    <item>five  <tag>return(5)</tag></item>
    <item>six   <tag>return(6)</tag></item>
    <item>seven <tag>return(7)</tag></item>
    <item>eight <tag>return(8)</tag></item>
    <item>nine  <tag>return(9)</tag></item>
  </one-of>
</rule>

<!-- Fillers -->
<rule id="pre_filler">
  <one-of>
    <item>i would like</item>
    <item>i'd like</item>
    <item>i want</item>
    <item>one</item>
  </one-of>
  to
  <one-of>
    <item>buy</item>
    <item>get</item>
    <item>purchase</item>
  </one-of>
</rule>

<rule id="post_filler">
  <one-of>
    <item>please</item>
  </one-of>
</rule>

</grammar>

```

Sample YesNo grammar

The following code shows two sample grammars, *YesNo* and *StrictYesNo*.

- *YesNo* handles loose expressions, for example, “Yes,” “that’s correct,” “thank you.”
- *StrictYesNo* is limited to one-word responses.

```

<?xml version="1.0" ?>
<grammar xmlns="http://www.w3.org/2001/06/grammar"
  xmlns:nuance="http://voicexml.nuance.com/grammar"
  version="1.0" xml:lang="en" root="YesNo">

  <!-- Public grammars -->

  <!-- Loose YesNo grammar -->

```

```

<rule id="YesNo" scope="public">
  <item repeat="0-1" repeat-prob="0.01">
    <ruleref uri="#YesNo_PreHesitation"/>
  </item>
  <ruleref uri="#YesNo_CORE"/>
</rule>

<!-- Strict YesNo grammar -->
<rule id="StrictYesNo" scope="public">
  <ruleref uri="#StrictYesNo_CORE"/>
</rule>

<!-- Helper subgrammars -->

<!-- Prehesitation -->
<rule id="YesNo_PreHesitation">
  <one-of>
    <item>um</item>
    <item>uh</item>
    <item>hm</item>
  </one-of>
</rule>

<!-- Core grammars -->
<rule id="YesNo_CORE">
  <one-of>
    <item>
      <ruleref uri="#YesLoose"/>
      <tag><![CDATA[<YesNo yes>]]></tag>
    </item>
    <item>
      <ruleref uri="#NoLoose"/>
      <tag><![CDATA[<YesNo no>]]></tag>
    </item>
  </one-of>
</rule>

<rule id="StrictYesNo_CORE">
  <one-of>
    <item>
      <ruleref uri="#YesStrict"/>
      <tag><![CDATA[<YesNo yes>]]></tag>
    </item>
    <item>
      <ruleref uri="#NoStrict"/>
      <tag><![CDATA[<YesNo no>]]></tag>
    </item>
  </one-of>
</rule>

<!-- Subgrammar used in YesNo -->
<rule id="YesLoose">
  <one-of>

```



```

<item>yes please</item>
<item>
  yes
  <item repeat="0-1">
    it
    <item repeat="0-1">
      <one-of>
        <item>sure</item>
        <item>certainly</item>
      </one-of>
    </item>
    is
  </item>
</item>
<item>
  it
  <item repeat="0-1">
    <one-of>
      <item>sure</item>
      <item>certainly</item>
    </one-of>
  </item>
  is
</item>
<item>yup</item>
<item>yeah</item>
<item>okay</item>
<item>sure</item>
<item>you got it</item>
<item>
  <item repeat="0-1">
    <item repeat="0-1">yes</item>
    <one-of>
      <item>that's</item>
      <item>it's</item>
      <item>that is</item>
      <item>it is</item>
    </one-of>
  </item>
  <one-of>
    <item>right</item>
    <item>correct</item>
  </one-of>
</item>
<item>
  <item repeat="0-1">
    <one-of>
      <item>yes</item>
      <item>yeah</item>
    </one-of>
  </item>
  i
  <one-of>

```

```

        <item>would</item>
        <item>do</item>
    </one-of>
</item>
</one-of>
<item repeat="0-1">
    <one-of>
        <item>thanks</item>
        <item>thank you</item>
    </one-of>
</item>
</rule>

<!-- Subgrammar used in StrictYesNo -->
<rule id="YesStrict">
    <one-of>
        <item>yes</item>
        <item>yup</item>
        <item>yeah</item>
        <item>right</item>
        <item>correct</item>
    </one-of>
</rule>

<!-- Subgrammar used in YesNo -->
<rule id="NoLoose">
    <one-of>
        <item>nope</item>
        <item>absolutely not</item>
        <item>
            no
            <item repeat="0-1">way</item>
        </item>
        <item>
            <item repeat="0-1">no</item>
            <one-of>
                <item>it isn't</item>
                <item>it's not</item>
                <item>it is not</item>
            </one-of>
        </item>
        <item>
            <item repeat="0-1">
                <item repeat="0-1">no</item>
                <one-of>
                    <item>that's</item>
                    <item>it's</item>
                    <item>that is</item>
                    <item>it is</item>
                </one-of>
            </item>
            <one-of>
                <item>wrong</item>

```

```

        <item>
            not
            <one-of>
                <item>correct</item>
                <item>right</item>
            </one-of>
        </item>
        <item>incorrect</item>
    </one-of>
</item>
<item>
    <item repeat="0-1">no</item>
    i
    <one-of>
        <item>would not</item>
        <item>wouldn't</item>
        <item>do not</item>
        <item>don't</item>
    </one-of>
</item>
</one-of>
<item repeat="0-1">
    <one-of>
        <item>thanks</item>
        <item>thank you</item>
    </one-of>
</item>
</rule>

<!-- Subgrammar used in StrictYesNo -->
<rule id="NoStrict">
    <one-of>
        <item>no</item>
        <item>nope</item>
        <item>wrong</item>
        <item>incorrect</item>
    </one-of>
</rule>

</grammar>

```


Index

SYMBOLS

\$return variable 55

A

ABNF 162
acoustic models 10
 naming conventions 101
 usage 103
add operator 61
adding
 NL commands 9
 probabilities 20
AddPhraseList 123
AddPhraseToDynamicGrammar 123
ambiguity
 robust interpretation 90
 testing 113
anticipating responses 6
arpabet-to-cpa 142
assign command 56
audio file 121
-auto_pron option to *nuance-compile* 108

B

backoff, in grammars 32-33
batchrec 119-129
 commands 122
 dynamic grammar commands 122
 Nuance parameters 127
 optional arguments 124
 output 127
 recording 119
 testset example 121
 testset file 121
 transcriptions 122
 usage 120

batchrec commands

 *Echo 121
 *Exit 121
 *GetParam 122
 *NewAudioChannel 122
 *SetParam 122

builtin grammars 35

C

character mapping file 104
CITY grammar 8
class-based SLMs 73-77, 82
closed vocabulary SLMs 78
CloseDynamicGrammarDatabase 122
commands, natural language 52
 list 65
 return 54
 slot-filling 53
-compare option to *nl-tool* 116
compilation servers
 configuring 36
 multiple servers 41
 with just-in-time grammars 41
CompileAndInsertDynamicGrammar 124
compiling 98
 just-in-time grammars 44, 98
 recognition packages. *See*
 nuance-compile,
 nuance-compile-ngo
compound-word modeling 108-109
Computer Phonetic Alphabet. *See* CPA
config.EGRCacheMB 36, 37
config.EGRCacheMinFreshSecs 28, 37
config.EGRProxy 36, 37
config.EGRTimeoutMS 29, 37
configuring for just-in-time
 grammars 35-38
Content-Base key 35, 40, 43

- conventions
 - grammar writing 103
 - language-specific 22
 - naming acoustic models 101
 - syntax 143
- CopyDynamicGrammar 123
- coverage test 113
- CPA 134-139
 - converting to 142
 - for American English 134
- crosswords 108, 112
 - modeling 108-109

D

- DATE grammar 8
- DBDescriptorToQueryString 27
- debug_level* option to *batchrec* 127
- defining slots 4
- DeleteDynamicGrammar 123
- dgdb references 25, 27, 38
- dialog
 - design 3
 - directed 5
 - mixed-initiative 5
- dictionaries
 - creating 133-142
 - example 109, 140
 - merge 107
 - override 107
 - sample 140
 - standard 105
- directed dialogs 5
- div operator 61
- do_crossword* option to
 - nuance-compile-ngo* 112
- documentation (Nuance)
 - accessing xi
 - description x
- dont_endpoint* option to *batchrec* 126
- dont_flatten* option to *nuance-compile* 28

- dynamic grammars 23-49
 - and SLMs 86
 - commands 122
 - enrollment 48
 - gates 46
 - in databases 23-25
- dynamic keyword 23

E

- Echo. *See batchrec* commands
- egr.Backoff 32, 34, 37
- egr.builtin.context 35, 37
- egr.dgdb.context 38
- egr.file.context 35, 37
- enable_jit* option 44
- endpoint* option to *batchrec* 126
- EnrollAbortPhrase 124
- EnrollCommitPhrase 124
- enrollment 48
- EnrollNewPhrase 124
- Exit. *See batchrec* commands
- external rule references 25-35
 - backoff 32-33
 - built-in grammars 35
 - defining rules as public 30
 - dgdb 25, 27, 38
 - examples 30
 - fail 26, 33-34
 - file 25, 26-27, 37
 - http 25, 26-27
 - local 25, 43-44
 - parameters 36
 - redefining 31-32, 43
 - redefining to a local rule 43-44
 - relative paths 34-35
 - relative URIs 43
 - root rules 29-31
 - special 25, 32-33, 46-48
 - static 25, 27-28
 - web server configuration 36

F

- fail references 26, 33-34
- failed references 33-34

- features 63
- fetching grammars 28-29
- file references 25, 26-27, 37
- filling multiple slots 54
- flight info example 4, 5
- functions 59-62
 - standard 60
 - user-defined 61

G

- gate technique 46
- gen_ref* option to *nl-tool* 116
- generate* 115-116
- generate-nlref* 125
- GetParam. *See batchrec* commands
- getting help, Nuance xiii
- grammar development
 - designing responses 2
 - guidelines 1
 - predicting responses 2
 - tasks 3
- grammar files 13
 - comments 15
 - compound words 108-109
 - enrollment support 48
 - grammar naming conventions 14
 - including other files 18
 - lists 64
 - reserved strings 16
 - sample 17
- Grammar keyword 121
- grammar* option to *batchrec* 121, 126
- Grammar Specification Language. *See* GSL
- grammar writing recommendations 22

- grammars
 - backoff 32-33
 - conventions 22
 - core 7
 - creating 13
 - default location 18
 - definition 13
 - description 13, 14
 - design principles 1
 - example of naturally phrased numbers 17
 - fetching 28-29
 - filler 7, 8
 - flattening 110
 - hierarchies 16
 - just-in-time 38-46
 - compilation options 98
 - compilation servers 41
 - configuration 35-38
 - multi-package support 42
 - options 40
 - with W3C 163
 - writing 38
 - name 13, 14
 - operators 15
 - probabilities 20
 - prompt coordination 3
 - recursive 19
 - rule 7
 - subgrammars 7, 16
 - top-level 16
 - writing 1
 - See also* dynamic grammars 23
- GSL
 - code example 9
 - enrollment support 48
 - identifying text strings 21
 - probabilities 20
 - reserved strings 145
 - syntax 143
- GSL operators
 - concatenation 15
 - disjunction 15
 - kleene closure 15
 - optional 15
 - positive closure 15

H

help, Nuance xiii
http references 25, 26-27

I

include directive (*#include*) 18
InsertDynamicGrammar 124
integer functions, in natural language
 slots 61
International Phonetic Alphabet 134
interpretation test 113

J

just-in-time grammars 38-46, 98
 and SLMs 86
 compilation options 98
 compilation servers 41
 compiling 44
 configuration 35-38
 multi-package support 42
 options 40
 with W3C 163
 writing 38

L

labels, maximum number in package 24
language-specific conventions 22
lexicon element 163
lists, as interpretation values 64
 commands 65
 functions 65
local references 25, 43-44

M

master packages 101
-merge_dictionary option to
 nuance-compile 107
meta element 162
MIME types 36
misspellings 107

mixed-initiative dialogs 5
ModifyPhrase 124
mul operator 61
multibyte character maps 104
multi-package support 42
multiple language support 105
multiple pronunciations 139
multiword 108
 -dont_flatten option 109
modeling 108-109

N

namespace 160
natural language interpretation 51-66
 commands 9, 52
 compiling packages 58
 complex slot values 62
 features 63
 lists 64
 return commands 54
 robust 86-92
 slot commands 53
 variables 54
N-best processing 83, 126
neg operator 61
nested structures 64
NewAudioChannel. *See batchrec*
 commands
NewDynamicGrammarEmpty 123
NewDynamicGrammarFromGSL 123
NewDynamicGrammarFromPhrase
 List 123
NL commands and unary operators 57
NL transcriptions file 125
-nl_transcriptions option to batchrec 125
nl-compile 58
nl-tag-tool 74
nl-tool 116
Nuance namespace 160
Nuance phoneme set 134
Nuance System, getting help xiii
nuance-compile 10, 28, 58, 73, 93, 164
nuance-compile-ngo 26, 45, 93, 97
Nuance-Config-Name key 40, 41

nuance-master-packages 103
Nuance-Package-Name key 40, 42
-num option to *generate* 115

O

open vocabulary SLMs 78, 82
OpenDynamicGrammarDatabase 122
operators
 add 61
 div 61
 mul 61
 neg 61
 strcat 61
 sub 61
order of the model, determining 79
over-generation test 113
-override_dictionary option to
 nuance-compile 107

P

package files summary 146
package name convention 101
package optimization 110-111
 graph algorithms 110
 unflattened grammars 110
parameters
 batchrec 127
 external rule references 36
 SLMs 83-84
parse-tool 114
passthrough keyword 32, 47
perplexity of a model, measuring 84-85
pfsg files in SLM grammars 81
phonemes 133-139
 American English 134
 languages other than English 142
 mappings to letters 135
phrase pronunciation 108
principles of grammar development 1
-print_confidence_scores option to
 batchrec 126
-print_trees option to *parse-tool* 115

-print_word_confidence_scores option to
 batchrec 127
probabilities and
 expressions 20
 kleene closure 21
 optional operator 21
 positive closure 21
process-slm 84
prompt design 5
pronounce 106
pronunciations 105-109
 compound-word 108-109
 crossword 108
 missing words 107
 multiple 139
 multiword 108
 testing 114
proxy settings for HTTP grammars 37
public keyword 30
public rules 30

Q

QueryDynamicGrammarContents 123
QueryDynamicGrammarContentsWith
 ID 123
QueryDynamicGrammarExists 123

R

-rcengine option to *batchrec* 126
rec.DoNBest 83
rec.GrammarWeight 83
rec.Interpret 67, 83
rec.pass1.gp.WTW 83
rec.PPR 83
rec.Pruning 83
recognition
 package 10, 58
 servers 117
 tuning, with *batchrec* 119
RECOGNIZE header 35, 40
recursive grammars 19
redefining external rules 31-32, 43-44
references. *See* external rule references

- referencing variables 57
- referring to specific rules 29-31
- regression test 114
- `RemovePhrase` 123
- `resistor` keyword 32, 48
- `return` command 54
- `roadblock` keyword 32, 47
- robust natural language
 - interpretation 86-92
- root rules 29
- rule element in W3C grammars 160
- `ruleref` element in W3C grammars 161

S

- Say Anything grammars 69-92
- semantic interpretation tags 163
- semantic uniqueing 67
- servers 117
- `SetParam`. *See* *batchrec* commands
- `slm` keyword 82
- SLMs 8, 69-92, 147-155
 - class-based 73-77
 - measuring perplexity 84-85
 - parameters 83-84
 - process-slm* 84
 - process-slm* tool 153
 - training 79-81
 - train-slm* 80
 - train-slm* tool 147
- slots
 - complex values 62
 - definition 4
 - definitions file 58
 - features 63
 - filling 57
 - name 125
 - nested structures 64
 - slot-filling commands 53-57
 - structures 62
 - values 4, 57
- special character mappings 104
- `special` references 25, 32-33, 46-48
 - `passthrough` 32, 47
 - `resistor` 32, 48
 - `roadblock` 32, 47

- specific rules, referring to 29-31
- standard dictionary 105
- `static` references 25, 27-28
- statistical NL technology 92
- `strcat` operator 61
- StrictYesNo* grammar 169
- string functions, in natural language
 - slots 61
- string variable* 56
- structures 63
- sub operator 61
- subgrammar 16
- summary of package files 146

T

- `tag-format` attribute 163
- tagging 74
 - grammar 74
- technical support xiii
- testing
 - ambiguity 113
 - command-line programs 114
 - coverage 113
 - interpretation 113
 - over-generation 113
 - pronunciation 114
 - regression tests 114
- testset* option to *batchrec* 121
- training
 - SLM grammars 73, 79-81
- train-slm* 80
- transcriptions
 - batchrec* 122
 - example 125
 - file 125
- transcriptions* option to *batchrec* 125

U

- unary operators and NL commands 57
- uniqueing, semantic 67
- Unix syntax conventions xii
- URLs. *See* external rule references
- user-defined functions 61

V

variables 57
 assigning 54
 string 56
 valid characters 55
vocabulary files in SLM grammars 77-79
-vrs option to *batchrec* 126

W

W3C 157-173
 ABNF 162
 extensions 160-162
 rule element 160
 ruleref element 161
 namespace 160
 NL features 163
 sample grammars 164-173
 dialer application 164
 retail application 167
 StrictYesNo 169
 YesNo 169
 semantic interpretation tags 163

wavconvert 119
web server configuration 36
Windows syntax notations xii
word names 14
World Wide Web Consortium. *See*
 W3C 157
-write_auto_pron_output option to
 nuance-compile 108

X

Xapp 117-118
Xwavedit 119

Y

YesNo grammar 121

